

Transformers

CPSC483: Deep Learning on Graph-Structured Data

Rex Ying

Questions

- How to summarize what it means for an ML system to be **trusted**?
- Name one of the four major characteristics of a **trustworthy** ML system according to the book's opinion.

Explain what does it mean and why it matters

- Have you noticed any news, articles, policies, events that have implications in trustworthy deep learning in recent years?

Readings

- Readings are updated on the website (syllabus page)
- **Readings:**
 - [\[2401.05561\] TrustLLM: Trustworthiness in Large Language Models](#)
- This lecture is not explicitly tested
 - But in future lectures we will assume knowledge of this when developing trustworthy components on top of Transformers


Outline of Today's Lecture

1. Self-Attention and Transformers

2. Transformers for (Large) Language Models (LLMs)

3. Transformers in Other Modalities


Sequence Learning



$$h_i = f_W(x_i, h_{i-1}), y_i = f_Y(h_i)$$

What are the issues and challenges of RNNs?

Sequence Learning




$$h_i = f_W(x_i, h_{i-1}), y_i = f_Y(h_i)$$

Problems of RNNs

- Sequential computation prevents parallelization
- Capacity of handling long sequences
- Mainly focusing on modeling recurrence
 - does not capture other correlations (hierarchical, long-range, polysemy....) well

Sequence Learning




$$h_i = f_W(x_i, h_{i-1}), y_i = f_Y(h_i)$$

Problems of RNNs

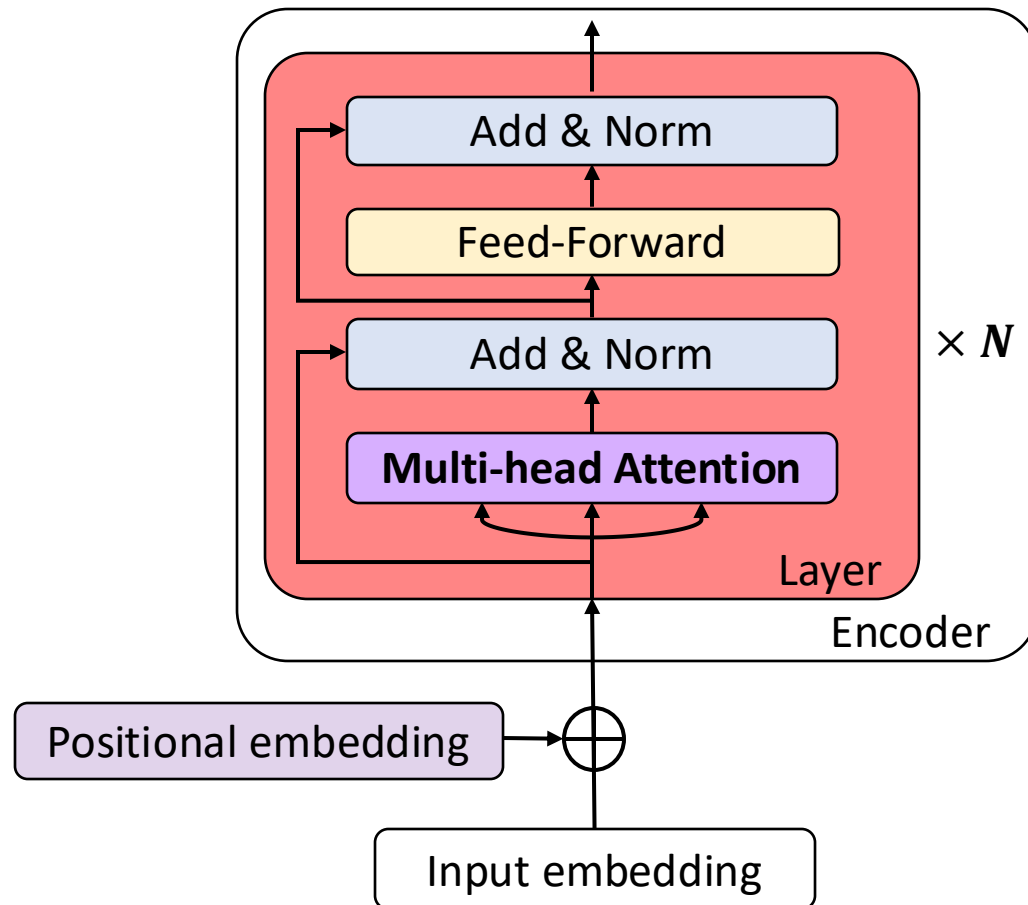
1. parallelization — — — — — ➔
2. long sequences — — — — — ➔
3. only recurrence — — — — — ➔

Solutions by Transformers

1. **Parallel input:** Input All tokens at the same time
2. **Self-Attention:** Enable attention in long-range
3. **Positional Embeddings p_i :** Model all possible correlations




Transformers — Overview



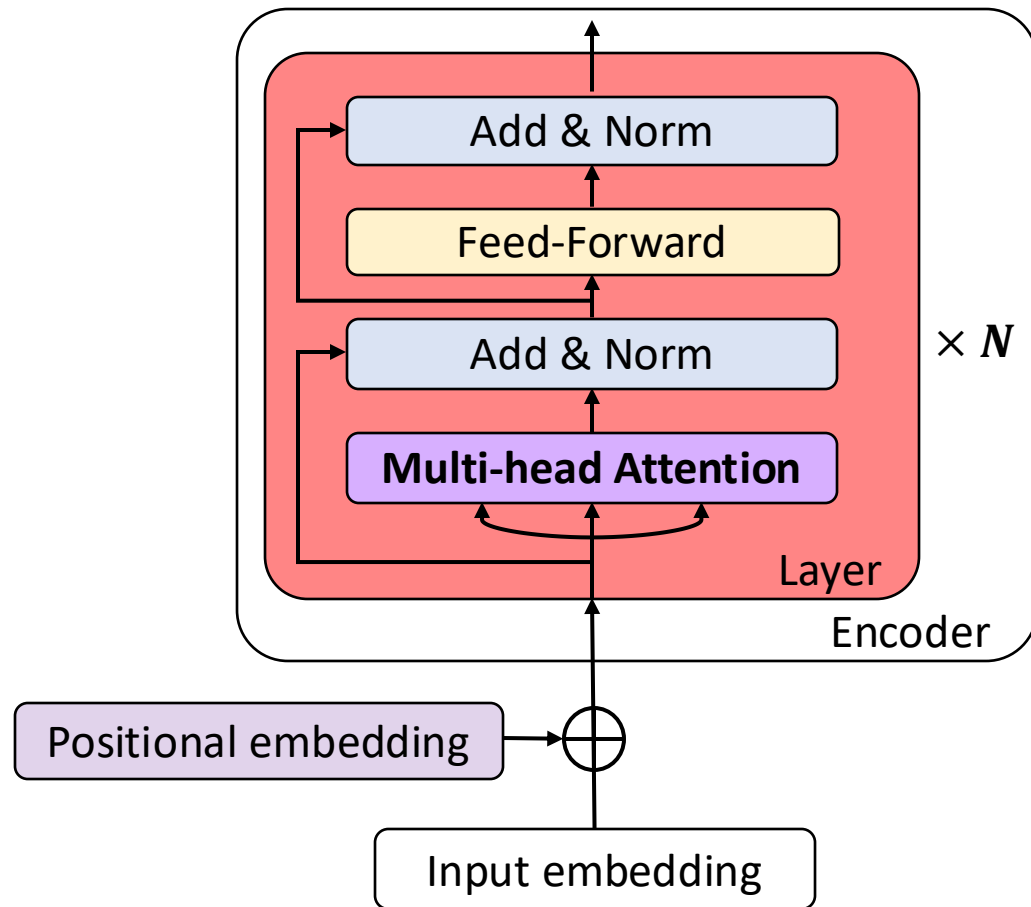
- **Original paper:** Attention is all you need [Vaswani et al., 2017].
- **Key component:** Multi-head self-attention
- **Other components** of a transformer layer: layer normalization, skip connection, position-wise feed-forward layer (FFN, or MLP)
- **Model usage:** Pre-training followed by fine-tuning. The transferred model can be:
 - **Encoder-only** (e.g BERT)
 - **Encoder-Decoder** (e.g [BART](#))
 - **Decoder-only** (e.g GPT)
 - We will show an example later

Transformers — Overview



- **Model usage:** Pre-training followed by fine-tuning. The transferred model can be:
 - **Encoder-only** (e.g BERT)
 - Many-to-one classification / regression
 - Sentiment classification, document classification ...
 - Word / Sentence embeddings for downstream tasks (e.g. recommender system)
 - **Encoder-Decoder** (e.g [BART](#))
 - **Decoder-only** (e.g GPT)
 - We will show an example later

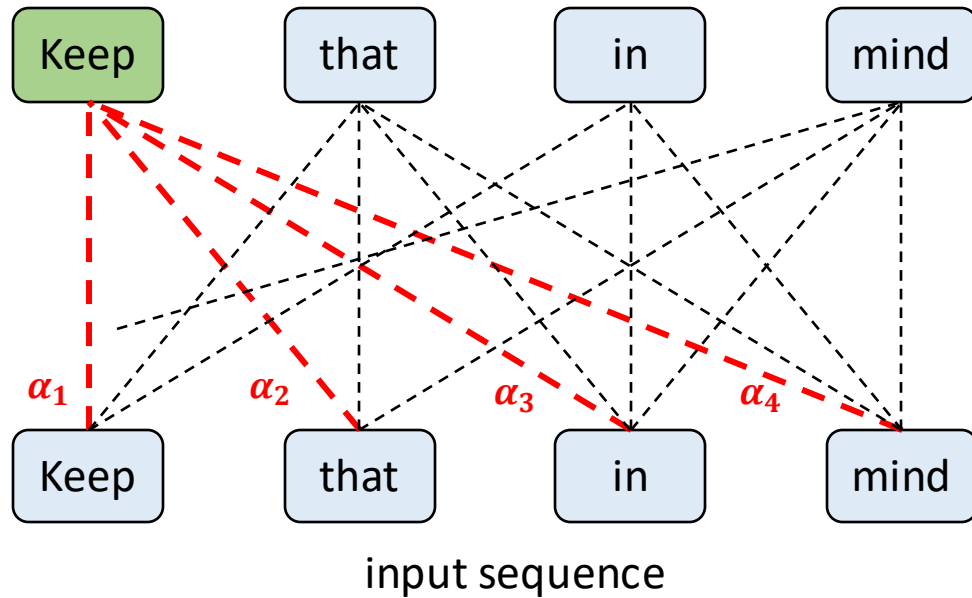
Transformers — Overview



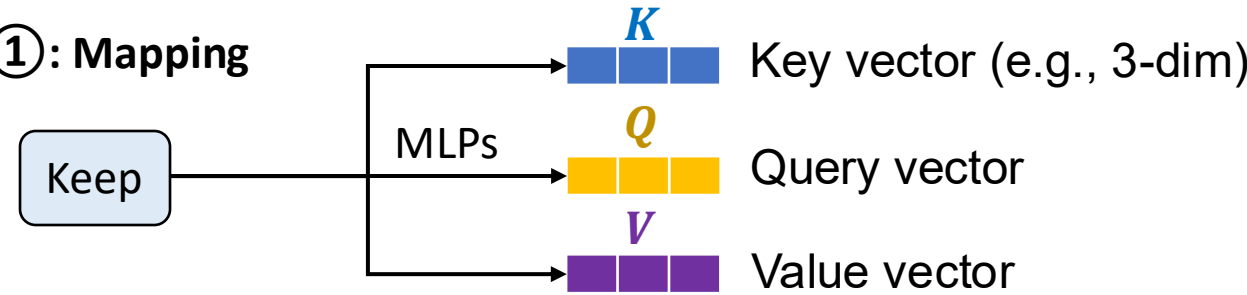
- **Model usage:** Pre-training followed by fine-tuning. The transferred model can be:
 - **Encoder-only** (e.g BERT)
 - **Encoder-Decoder** (e.g [BART](#))
 - Many-to-many use cases
 - Summarization, translation, style transfer ...
 - **Decoder-only** (e.g OpenAI GPT)
 - One-to-many use cases
 - Image / text / code generation, dialogue systems ...
 - GPT-3/4 based [apps](#)

Transformers — Self-Attention (1/5)

Example:



Step ①: Mapping



Step ②: Attention



$$\alpha_1, \alpha_2, \alpha_3, \alpha_4 = \text{Softmax} \left(\begin{matrix} Q \\ \text{Keep} \end{matrix} \times \begin{matrix} K & K & K & K \\ \text{Keep} & \text{that} & \text{in} & \text{mind} \end{matrix} \right)$$

Step ③: Update

$$\begin{matrix} V' \\ \text{Keep} \end{matrix} = \alpha_1 \times \begin{matrix} V \\ \text{Keep} \end{matrix} + \alpha_2 \times \begin{matrix} V \\ \text{that} \end{matrix} + \alpha_3 \times \begin{matrix} V \\ \text{in} \end{matrix} + \alpha_4 \times \begin{matrix} V \\ \text{mind} \end{matrix}$$

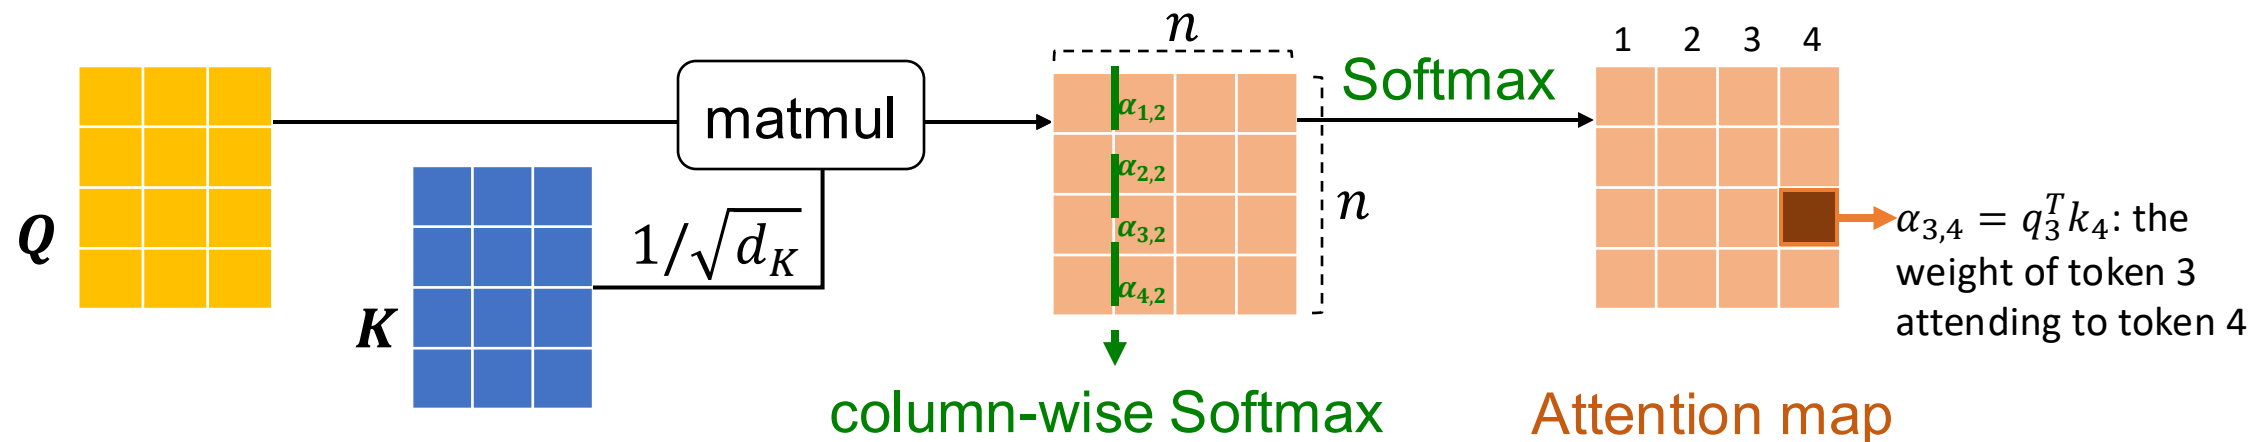
Transformers — Self-Attention (2/5)

- Formally, given an input sequence $X = [x_1, x_2, \dots, x_n] \in \mathbb{R}^{n \times d_X}$
- Step ①: Query $Q = XW_Q$, Key $K = XW_K$, Value $V = XW_V$
 - $W_K \in \mathbb{R}^{d_X \times d_K}$, and thus $K \in \mathbb{R}^{n \times d_K}$
 - We require $d_K = d_Q$, for simplicity, we set $d_K = d_Q = d_V := d$



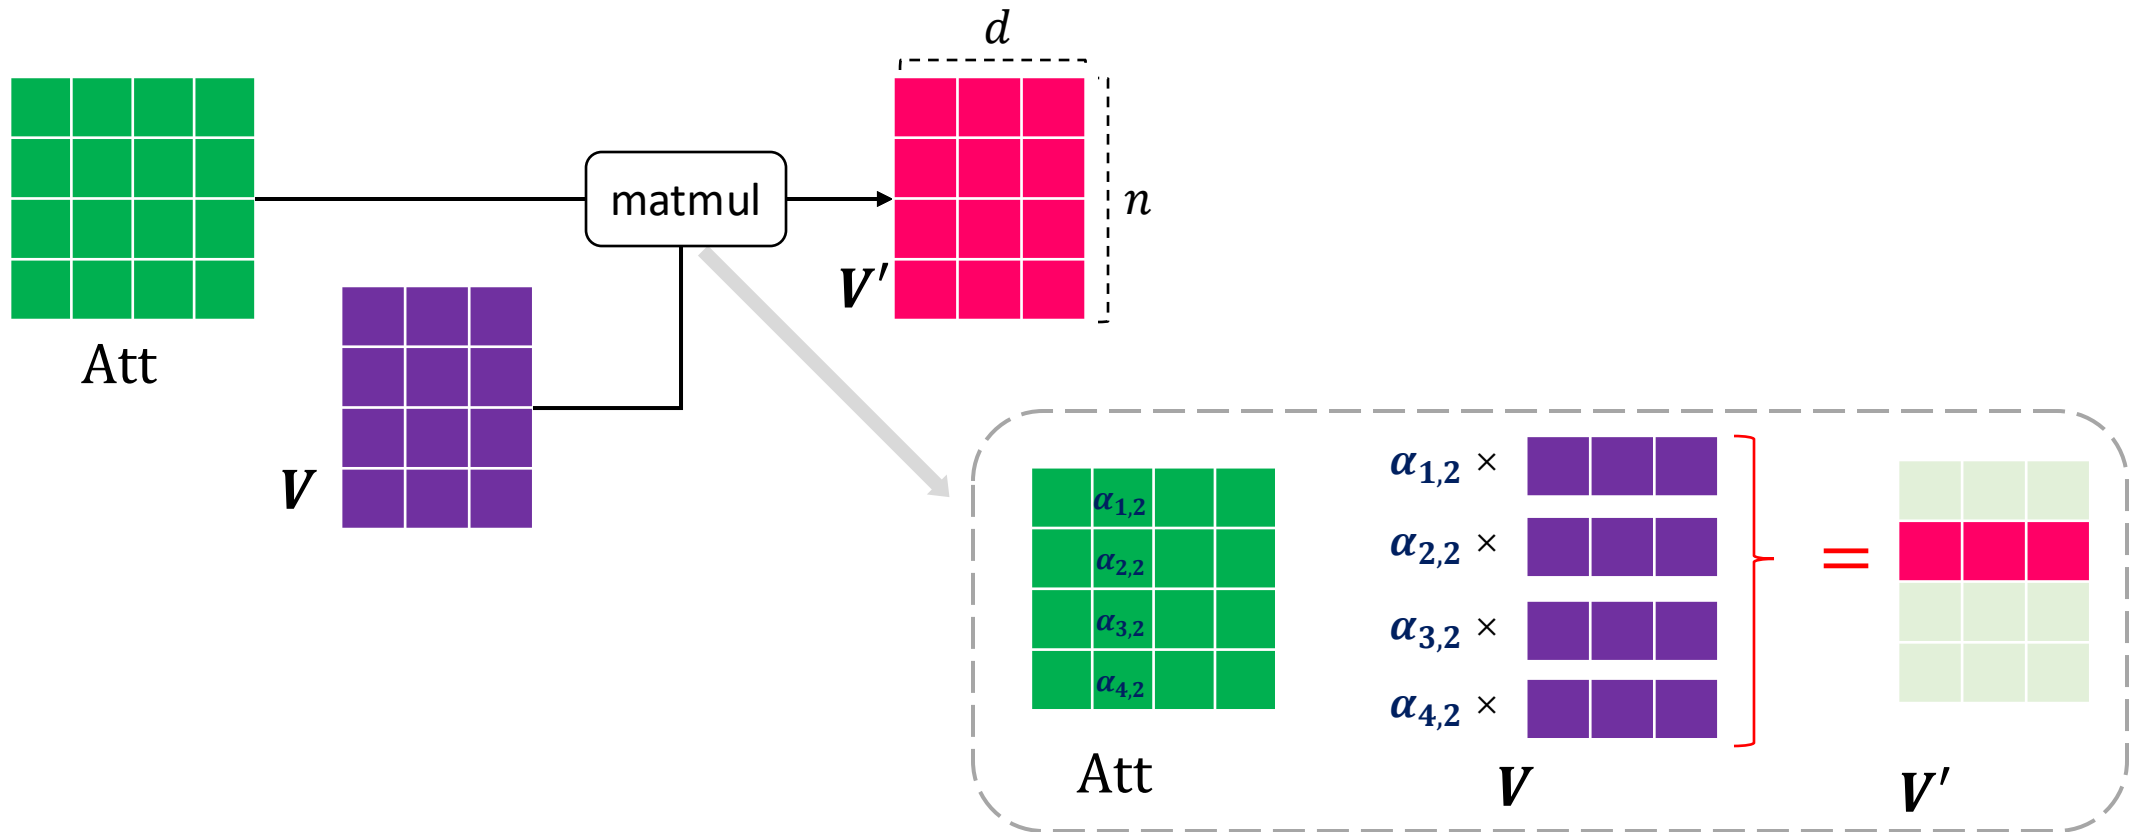
Transformers — Self-Attention (3/5)

- Step ② : Attention map $\text{Att} = \text{Softmax}\left(\frac{QK^T}{\sqrt{d}}\right) \in \mathbb{R}^{n \times n}$ (Softmax is col-wise)
 - The matrix multiplication QK^T performs dot-product for every possible pair of queries and keys, resulting in an attention map.
 - **Normalization factor** $1/\sqrt{d_K}$: performing dot-product over two vectors with variance σ^2 results in a scalar having d_K -times higher variance,
 - $q \sim N(0, \sigma^2), k \sim N(0, \sigma^2) \rightarrow \text{Var}\left(\sum_{i=1}^{d_K} q[i]k[i]\right) = \sigma^4 d_K$




Transformers — Self-Attention (4/5)

- Step ③: Updated value $V' = \text{Att } V \in \mathbb{R}^{n \times d}$ Matrix product



Transformers — Self-Attention (5/5)

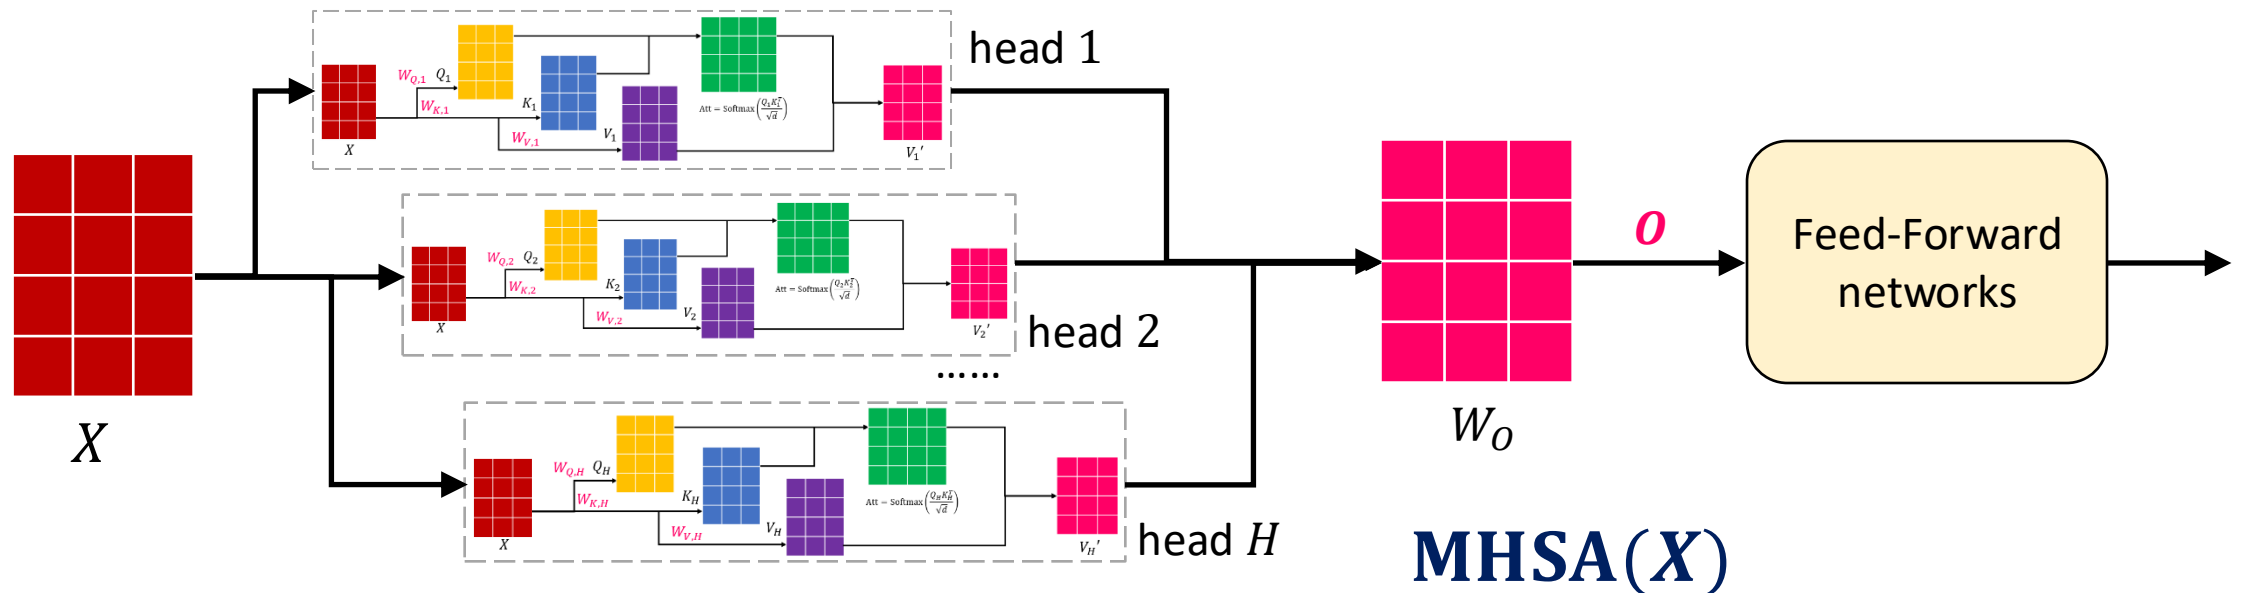
- Putting all together



The computation complexity is quadratic to number of tokens


Transformers — Multi-Head Self-Attention

- There are usually **multiple aspects** that a token can attend to.
- We extend the attention to multiple heads, with multiple (Q, K, V) triplets on the same features.
 - The output of multi-head self-attention $O = \text{Concat}([V'_1, V'_2, \dots, V'_H])W_O$
 - Learnable parameters in each attention layer: $W_{Q,i}, W_{K,i}, W_{V,i} \in R^{d_x \times d}$ for head i , $W_O \in R^{Hd \times d_o}$



Transformers — Layer (1)


- **MHSA**: multi-head self-attention
- Transformer layer: $X \rightarrow \text{LayerNorm}(X + \text{MHSA}(X))$
- **Residual connections** are added to
 - Enable smooth gradient flow in deep transformers
 - Keep the information of the original sequence.



What are some advantages or challenges when trying to explain Transformer-based models?

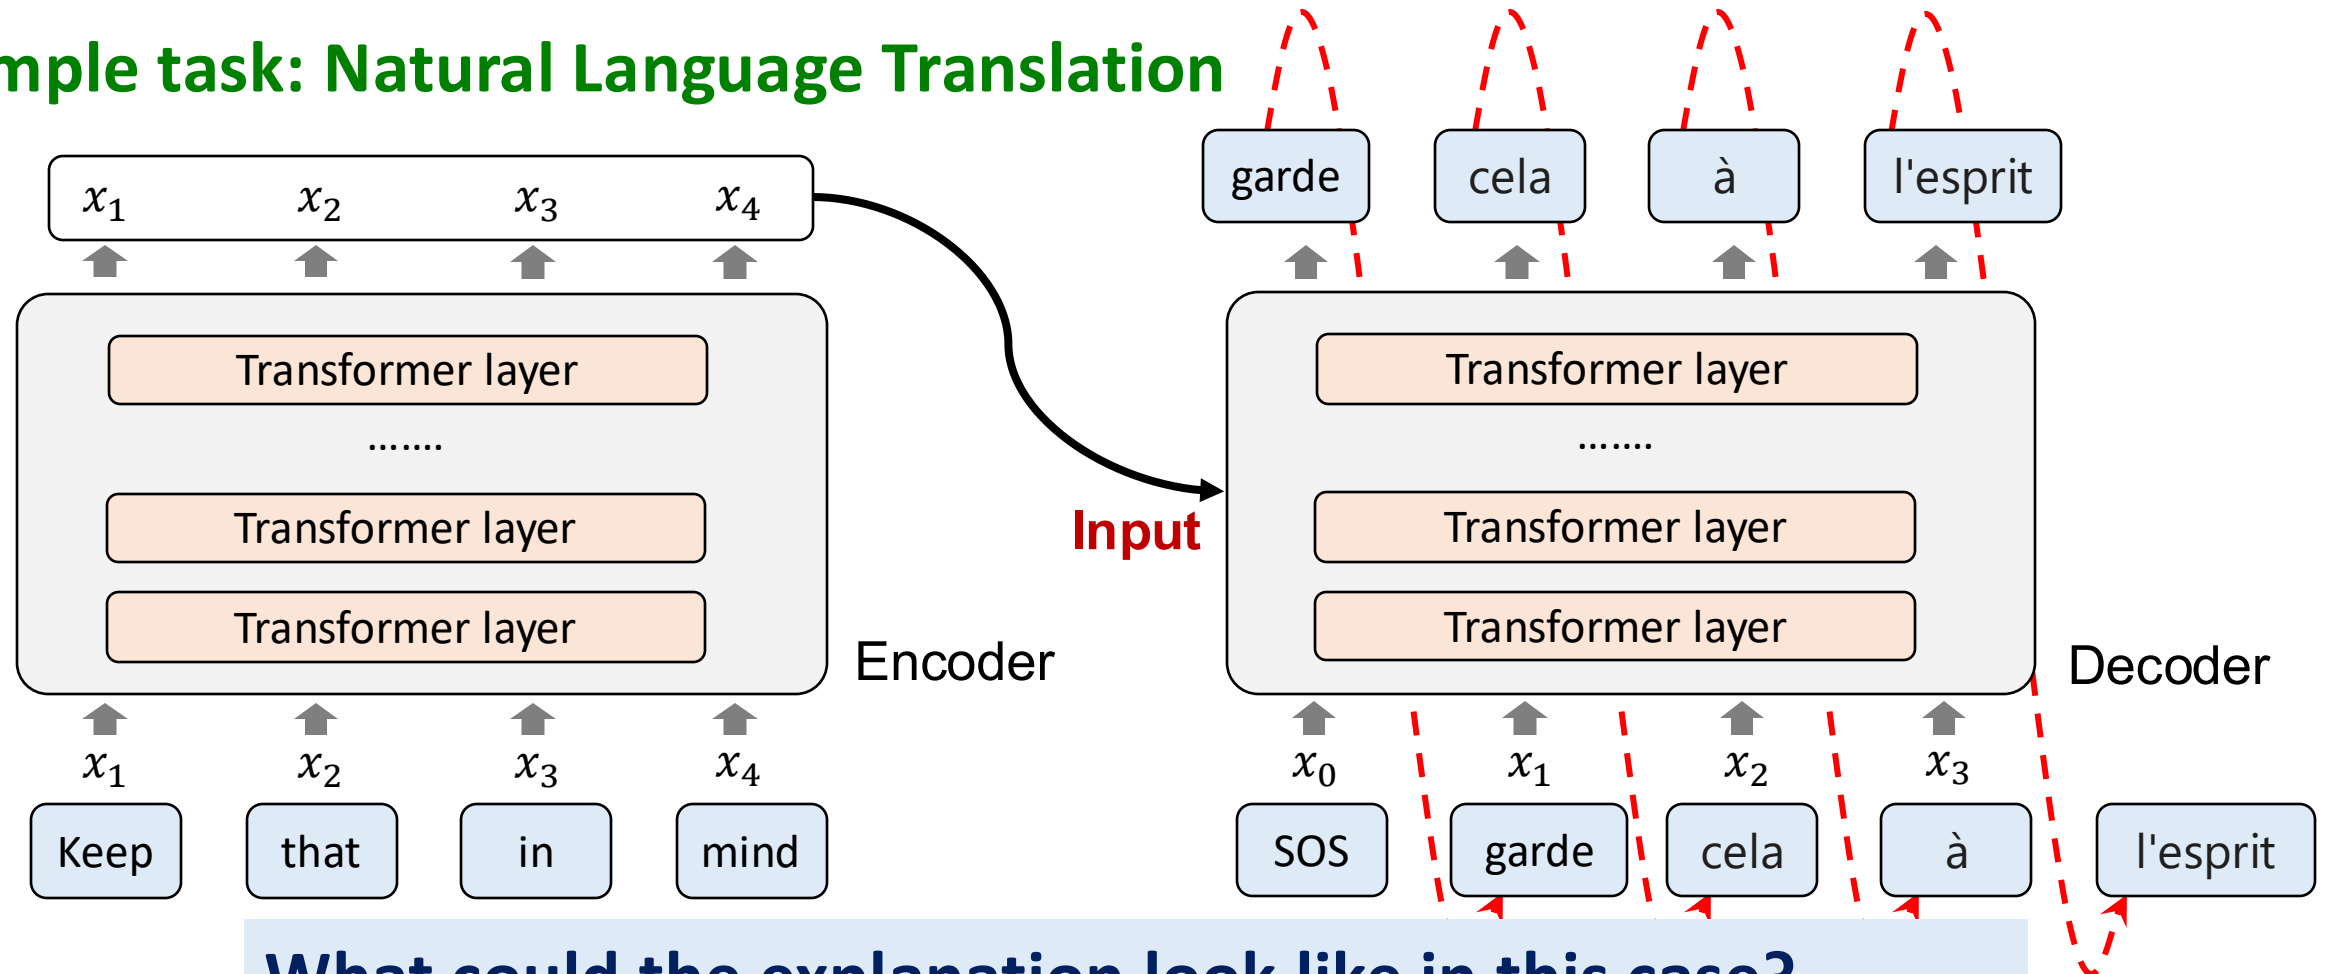
Transformers — Layer (2)

- Transformer layer: $X \rightarrow \text{LayerNorm}(X + \text{MHSA}(X)) \rightarrow \text{LayerNorm}(X + \text{FFN}(X))$
- Layer Normalization** is used to enable faster training with small regularization and keep features in similar magnitudes.
 - BatchNorm isn't applied because batch size is usually small in Transformers due to GPU memory constraints. Besides, BatchNorm has been shown to lead to worse performance in NLP.
- MLPs** are added for “post-processing”, and allow transformations on each sequence token.



Transformers —Encoder / Decoder

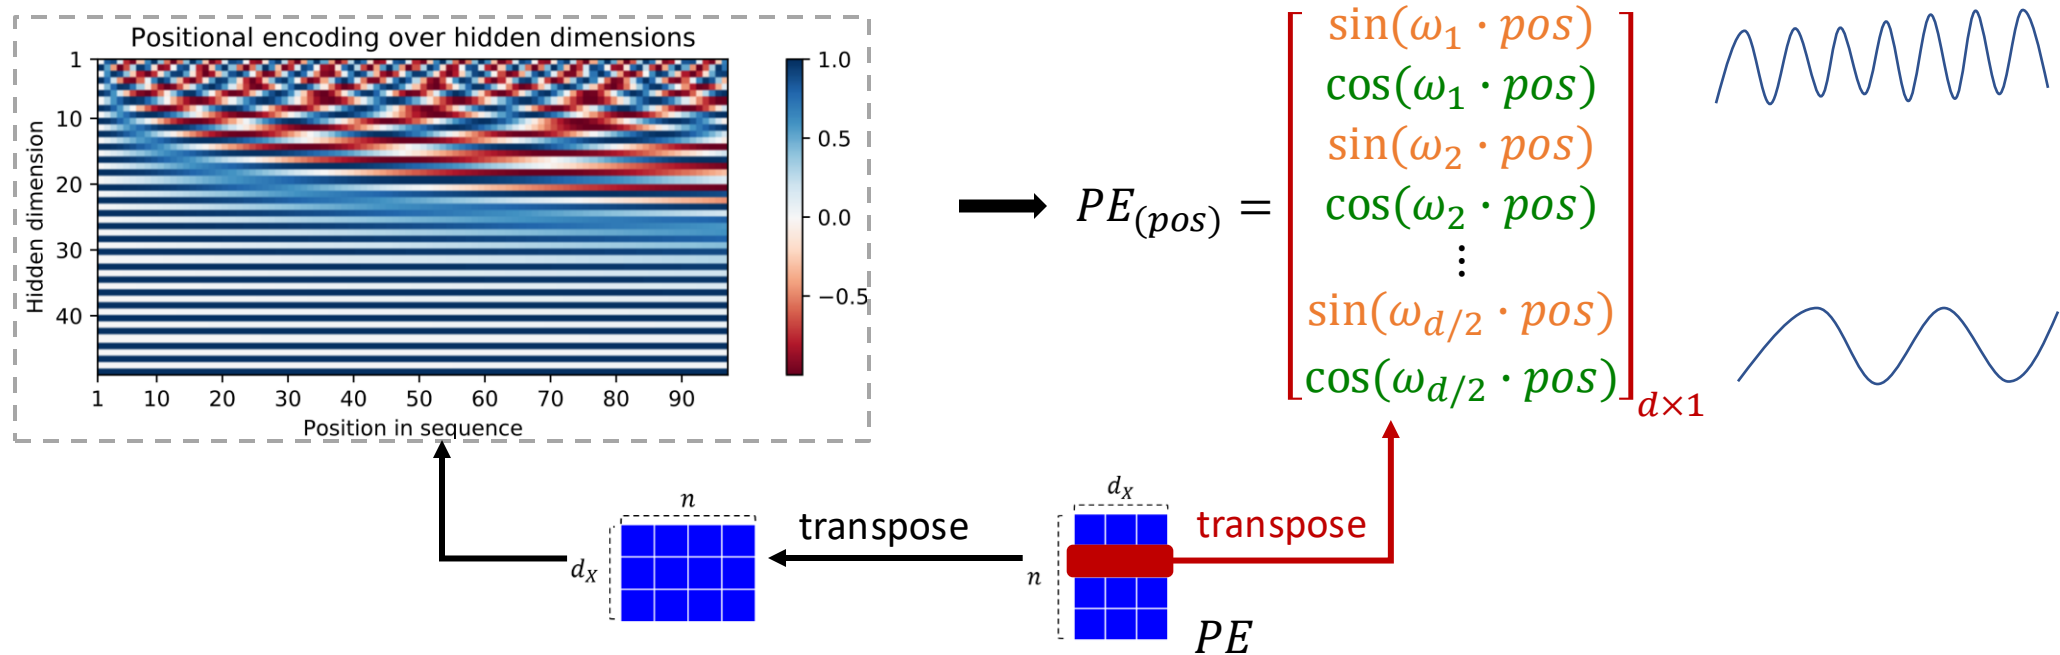
Example task: Natural Language Translation



Transformers —Positional Encoding (3)

- **Cosine encoding**

- $PE_{(pos,2i)} = \sin(pos/10000^{2i/d_x})$, $PE_{(pos,2i+1)} = \cos(pos/10000^{2i/d_x})$.
- $\omega_i = 1/10000^{2i/d_x}$.
- Relative distance: $PE_{(pos+k)}$ can be easily represented as a linear function of $PE_{(pos)}$ (show it).



Summary: Transformer Architecture

- Multi-Head Self-Attention (**MHSA**(X))
 - For head i
 - $Q_i = XW_{Q_i}, K_i = XW_{K_i}, V_i = XW_{V_i}$
 - $\text{Att}_i = \text{Softmax}\left(\frac{Q_i K_i^T}{\sqrt{d}}\right) \in \mathbb{R}^{n \times n}$
 - $V_i' = \text{Att}_i V_i \in \mathbb{R}^{n \times d}$
 - Concatenating all heads: $O = \text{Concat}([V_1', V_2', \dots, V_H'])W_O$
- $X = \text{LayerNorm}(X + \text{MHSA}(X))$
- $X = \text{LayerNorm}(X + \text{FFN}(X))$

Outline of Today's Lecture


1. Self-Attention and Transformers

2. Transformers for (Large) Language Models (LLMs)

3. Transformers for Other Modalities

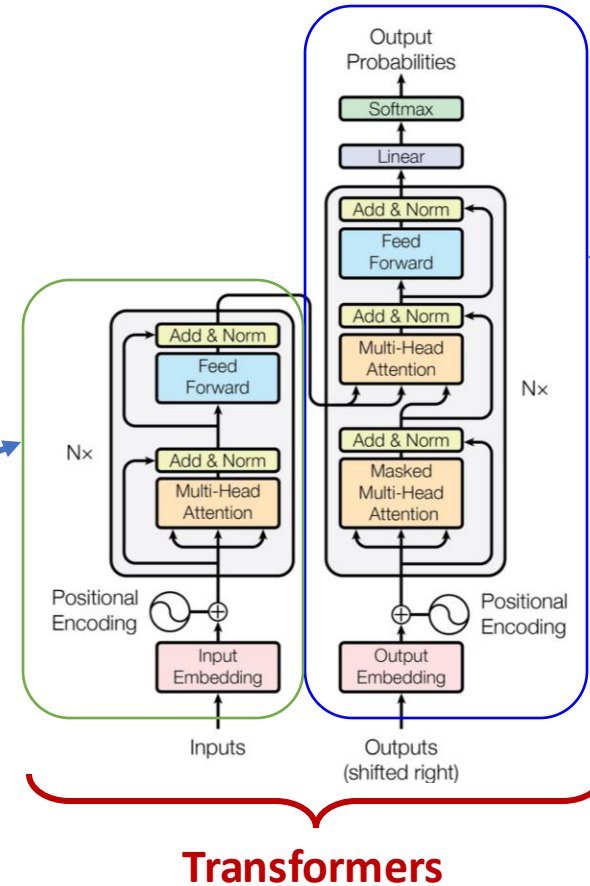
A Bit of History - Foundation

Foundational phase



BERT
Encoder-only

GPT
Decoder-Only




What is the best architecture?

Transformers in NLP — BERT

BERT — **Bidirectional** Encoder Representations from Transformers [Devlin et al., 2018]

- **Pre-training task** (self-supervised): **Masked Language Model (MLM)**
 - First randomly masking $m\%$ tokens in the input sequence.
 - In BERT, 15% tokens are masked at random (replaced with the special [MASK] token)
 - Predicting masked tokens using remaining tokens.
 - Two modes: **Unidirectional** and **Bidirectional**.

Unidirectional [Radford et al., 2018]



- Maximize Likelihood of “all” given “Att” and “is”

Bidirectional



- Maximize Likelihood of “all”, given “Att” , “is”, “you”, “need”.

Transformers in NLP — RoBERTa


RoBERTa — **Robustly** Optimized BERT [Liu et al., 2019]

- **Pretraining data**: BooksCorpus (800 M words) [Zhu et al., 2015], English Wiki (2500 M words), CC-News, OpenWebText [Gokaslan and Cohen, 2019], Stories [Trinh and Le, 2018]
 - Partition the corpus into “sentences” with fixed length of 512 tokens.
- **Hyperparameters** in use (also commonly adopted in most NLP Transformers):
 - **12-Layer** Encoder + **12-Layer** Decoder
(Pretrained Encoder is used more frequently in down-stream tasks)
 - Hidden dimension **768** = 12 (num of Heads) \times 64 (dim of Head)
 - Learning rate: Warmup then linear decay
 - Warmup: Gradually increasing the learning rate to a specific value in the first few epochs
 - Linear decay: Decreasing the learning rate by the same amount (decrement) every epoch.

Transformers in NLP — GPT


GPT — **Generative** Pre-trained Transformer

- **Pre-training task** (self-supervised): **Causal Language Model (CLM)**
 - Predict the next token at every position (left→right).
 - Apply a causal mask so each token attends only to previous tokens.
 - No [MASK] tokens; training matches inference.
 - Objective: $\mathcal{L}_{CLM} = -\sum_i \log P(x_i | x_{<i})$



Transformers in NLP — GPT Decoding

During generation, GPT will generate/predict the next token depending on the current context. This task is called **decoding**.

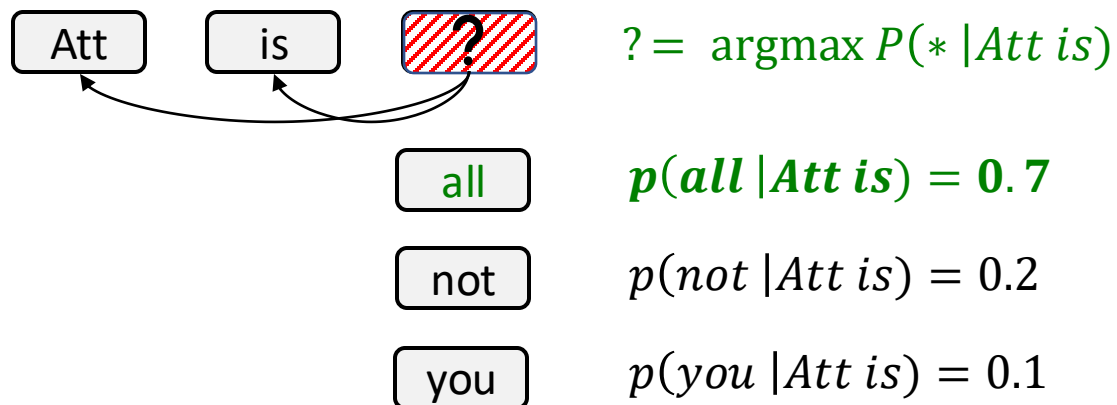


Transformers in NLP — GPT Decoding

During generation, GPT will generate/predict the next token depending on the current context. This task is called **decoding**.

Two common decoding methods:

- **Greedy Decoding:** choose the next token with the highest probability




Transformers in NLP — GPT Decoding

During generation, GPT will generate/predict the next token depending on the current context. This task is called **decoding**.

Two common decoding methods:

- **Greedy Decoding**: choose the next token with the highest probability
- **Sampling**: choose a random token according to their probability assigned by the model



What is the effect of τ ?

Summary: BERT vs GPT

Most of SOTA embedding models on [MTEB](#) leaderboard are decoder-only models

MTEB is a well-known benchmark for embedding models

Rank (Bo...	Model	Zero-shot	Memory U...	Number of P...	Embedding D...	Max Tokens	Mean (T...	Mean (TaskT...	Bitext ...	Classification	Clusterin
1	gemini-embedding-001	99%	Unknown	Unknown	3072	2048	68.37	59.59	79.28	71.82	54.59
2	Qwen3-Embedding-8B	99%	28866	7B	4096	32768	70.58	61.69	80.89	74.00	57.65
3	Qwen3-Embedding-4B	99%	15341	4B	2560	32768	69.45	60.86	79.36	72.33	57.15
4	Qwen3-Embedding-0.6B	99%	2272	595M	1024	32768	64.34	56.01	72.23	66.83	52.33
5	Linq-Embed-Mistral	99%	13563	7B	4096	32768	61.47	54.14	70.34	62.24	50.60
6	gte-Qwen2-7B-instruct	⚠️ NA	29040	7B	3584	32768	62.51	55.93	73.92	61.55	52.77
7	multilingual-e5-large-instruct	99%	1068	560M	1024	514	63.22	55.08	80.13	64.94	50.75
8	SFR-Embedding-Mistral	96%	13563	7B	4096	32768	60.90	53.92	70.00	60.02	51.84
9	text-multilingual-embedding-002	99%	Unknown	Unknown	768	2048	62.16	54.25	70.73	64.64	47.84
10	GritLM-7B	99%	13813	7B	4096	4096	60.92	53.74	70.53	61.83	49.75
11	GritLM-8B	99%	20070	8B	4096	4096	60.40	53.21	69.17	61.55	50.16


inputs (shifted right)

Some studies show that decoder-only models outperform encoder-decoder and encoder-only models using a similar configuration.

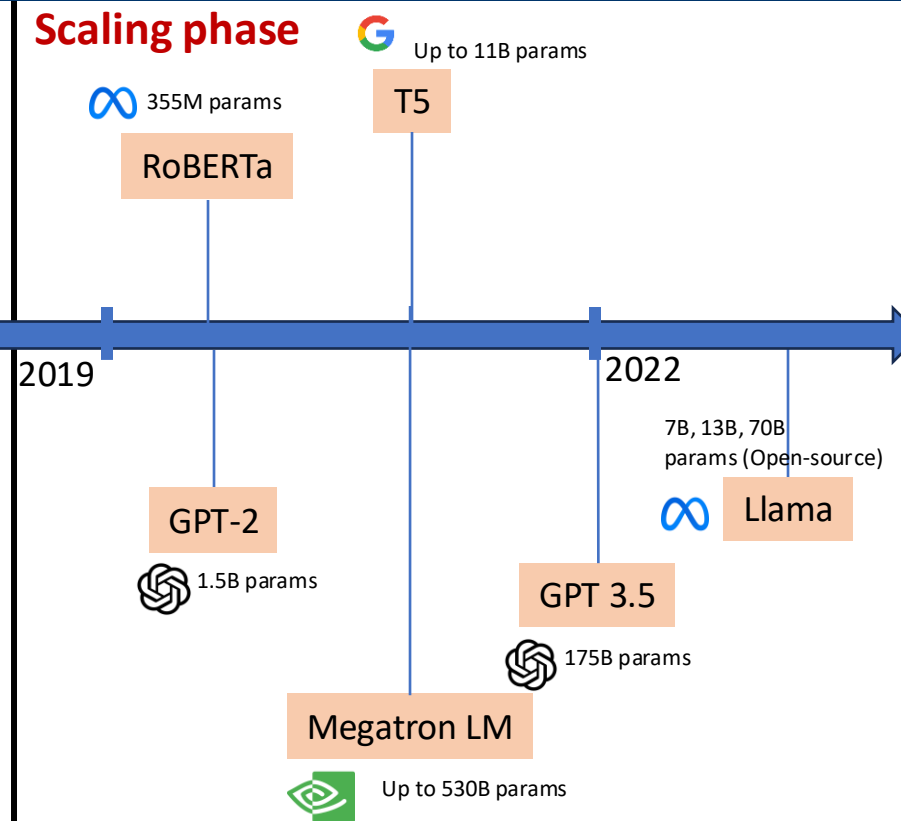
	EAI-EVAL	T0-EVAL
Causal decoder	44.2	42.4
Non-causal decoder	43.5	41.8
Encoder-decoder	39.9	41.7
Random baseline	32.9	41.7

A Bit of History - Scaling

Foundational phase



Scaling phase



What is the best architecture?

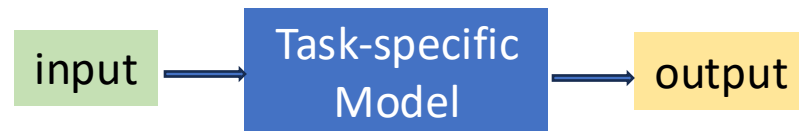
Prompting/Instruction-tuning
More parameters, more data
Scaling laws
Emergent capabilities

Scaling Law of LLMs (1)

Task Finetuning: A common paradigm of reusing LMs across diverse tasks is by fine-tuning for each task of interest.

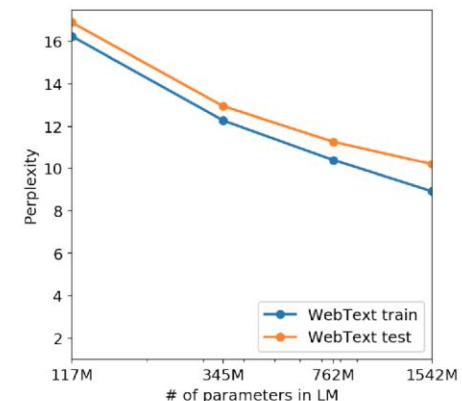
- It was a standard approach before 2020. It yields higher accuracy on target tasks but requires curated data and can be resource-intensive for running multiple models.

• Can we train a single model that can perform NLP tasks in zero-shot manner?



- GPT-2's success showed that a single model can improve its zero-shot performance by **scaling the model parameters**.


	LAMBADA (PPL)	LAMBADA (ACC)	CBT-CN (ACC)	CBT-NE (ACC)	WikiText2 (PPL)	PTB (PPL)	enwik8 (BPB)	text8 (BPC)	WikiText103 (PPL)	1BW (PPL)
SOTA	99.8	59.23	85.7	82.3	39.14	46.54	0.99	1.08	18.3	21.8
117M	35.13	45.99	87.65	83.4	29.41	65.85	1.16	1.17	37.50	75.20
345M	15.60	55.48	92.35	87.1	22.76	47.33	1.01	1.06	26.37	55.72
762M	10.87	60.12	93.45	88.0	19.93	40.31	0.97	1.02	22.05	44.575
1542M	8.63	63.24	93.30	89.05	18.34	35.76	0.93	0.98	17.48	42.16



Scaling Law of LLMs (2)

Scaling Laws for Neural Language Models ([Kaplan et al, 2020](#)) further shows that

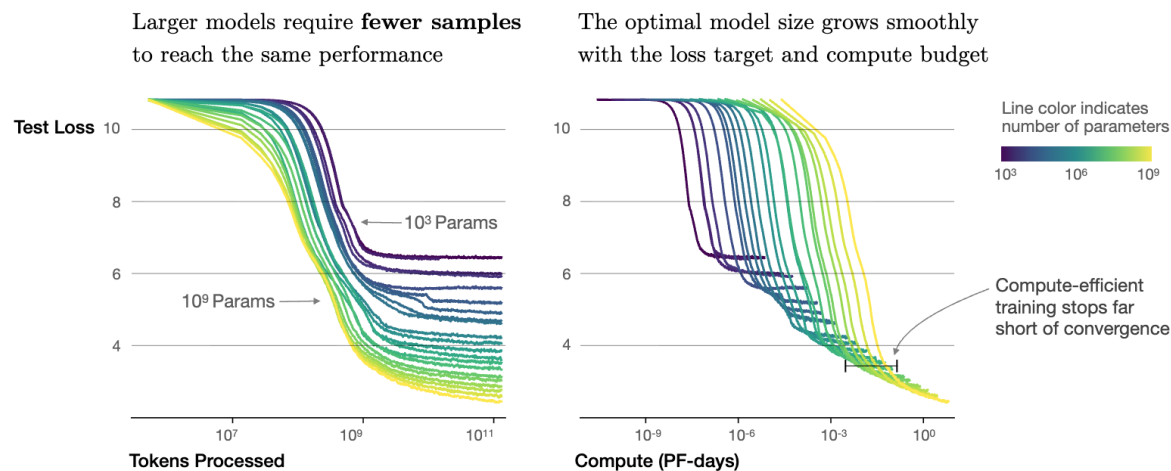
- Performance (loss/error) follows **predictable power-law scaling** with:
 - Model size (parameters)
 - Dataset size (tokens)
 - Compute used (training FLOPs)



Scaling Law of LLMs (2)

Scaling Laws for Neural Language Models ([Kaplan et al, 2020](#)) further shows that

- Performance (loss/error) follows **predictable power-law scaling** with:
 - Model size (parameters)
 - Dataset size (tokens)
 - Compute used (training FLOPs)
- Performance depends strongly on scale, and weakly on the model shape
- Bigger models trained on more data **consistently improve** performance.



Scaling Law of LLMs (3)

Later research ([Hoffmann et al., Chinchilla, 2022](#)) refined this finding:

- Many LLMs were **undertrained**
- **More data with smaller models** can be better than just bigger models.

Implication: With scaling laws, we can make decisions on architecture, data, and hyperparameters by training a smaller model.

- Note that training modern LLMs is **extremely resource-intensive**
 - Example: **GPT-3 is estimated to** require **~3640 petaflop/s-days** (**≈ 0.80M GPU-hours on A100s** ~ 33 days on 1,024 A100 GPUs).
 - Cost estimates: **\$4–5 million** in compute (at \$5.12/GPU-hr).

In-context Learning and Emergent Capabilities

- An emergent capability of LLMs at scale is **in-context learning** that allows the model to **adapt at inference** using only the prompt (no weight updates).
 - Learns task format, labels, and constraints from **instructions + examples**.

Zero-shot prompting

Given a math problem, give a direct answer.

Problem: Amy collects eggs for one week from 14 hens. Mon–Fri each hen lays 2 eggs/day; Sat–Sun 1 egg/day. How many eggs in a week?

168

Few-shot Prompting

Q: 10 hens; Mon–Fri 2/day; Sat–Sun 1/day. How many eggs in a week?

A: 120

Q: 6 hens; Mon–Fri 2/day; Sat–Sun 1/day. How many eggs in a week?

A: 72

Q: Amy collects eggs for one week from 14 hens. Mon–Fri each hen lays 2 eggs/day; Sat–Sun 1 egg/day. How many eggs in a week?

A: 168

Chain-of-Thought Prompting


Amy collects eggs for one week from 14 hens. Mon–Fri each hen lays 2 eggs/day; Sat–Sun 1 egg/day. How many eggs in a week? **Let's think step-by-step.**

Each hen lays 2 eggs per day from Monday to Friday (5 days), totaling 10 eggs, and 1 egg per day over the weekend (2 days), adding 2 more. That's 12 eggs per hen each week. With 14 hens, the total is $14 \times 12 =$ **168 eggs per week.**

In-context Learning and Emergent Capabilities

- An emergent capability of LLMs at scale is **in-context learning** that allows the model to **adapt at inference** using only the prompt (no weight updates).
 - Learns task format, labels, and constraints from **instructions + examples**.

Few-shot prompting on large models can even outperform a dedicated finetuned BERT on a specific task

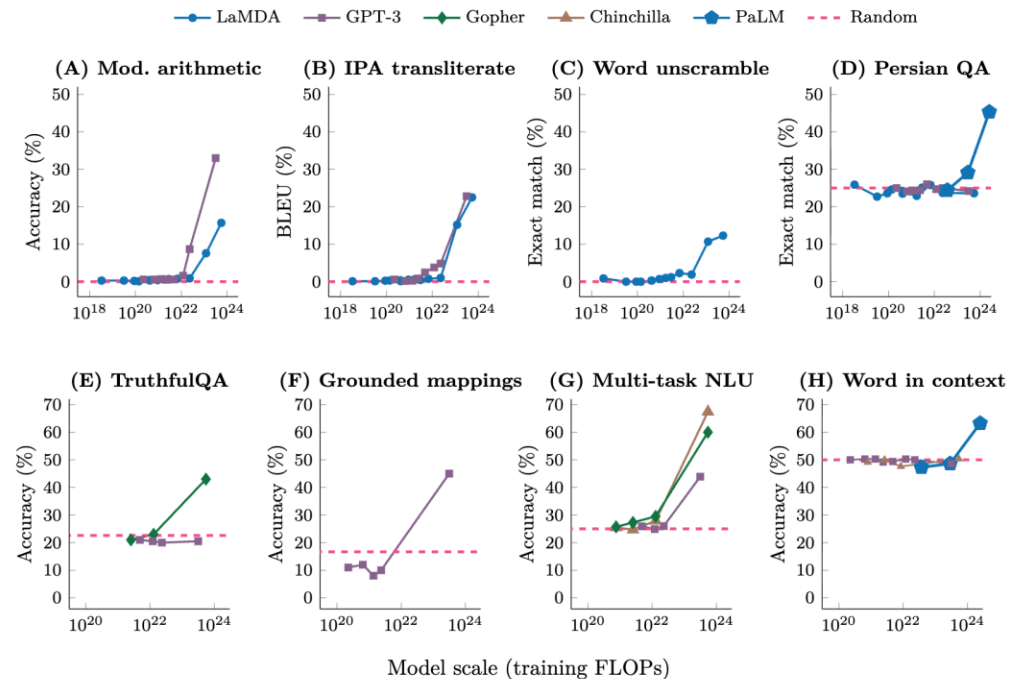


[Brown et al, Language Models are Few-Shot Learners, 2020](#)

In-context Learning and Emergent Capabilities

- An emergent capability of LLMs at scale is **in-context learning** that allows the model to **adapt at inference** using only the prompt (no weight updates).
 - Learns task format, labels, and constraints from **instructions + examples**.


The few-shot learning ability of LLMs is usually emerges after the model scale crosses a threshold



[Wei et al., Emergent Abilities of Large Language Models \(2022\)](#)

In-context Learning and Emergent Capabilities

There are also many other abilities that are **emergent in large models** but not in small models.



- Many abilities appear **abruptly** once scale (params \times data \times compute) crosses a threshold.
- **Scale \neq everything**: on subsets of BIG-bench, model family/training can trump size—e.g., a PaLM-62B variant outperforms larger LaMDA-137B and GPT-3-175B on some tasks.
- **Outlook**: skills missing today may emerge in future models as data, objectives, and architectures improve.

[Wei et al., Emergent Abilities of Large Language Models \(2022\)](#)

Training Process of an LLM


- **Pretrained LLMs are**

- Chaotic, aggregating many styles and value systems—modeling “a giant mass of people”
 - Toxic comments on X, Reddit
 - Disinformation on fraud websites
 - Bad outputs of old ML systems (e.g., old machine translation) → ML feedback loop reinforces the errors

→ **Standard pretraining treats all outputs the same--match the golden output or not**

- Some mistakes cause more damage than others. E.g.,
 - “Transfer **\$500 to Alice today.**” — correct
 - “Transfer **\$500 today.**” — missing recipient → minor/blocking
 - “Transfer **\$5,000 to Alice today.**” — wrong amount → costly
 - “Transfer **\$500 to Alex today.**” — wrong recipient → severe

→ **The model is not exposed to mistakes during training**



Training Process of an LLM

- **Pre-training:**

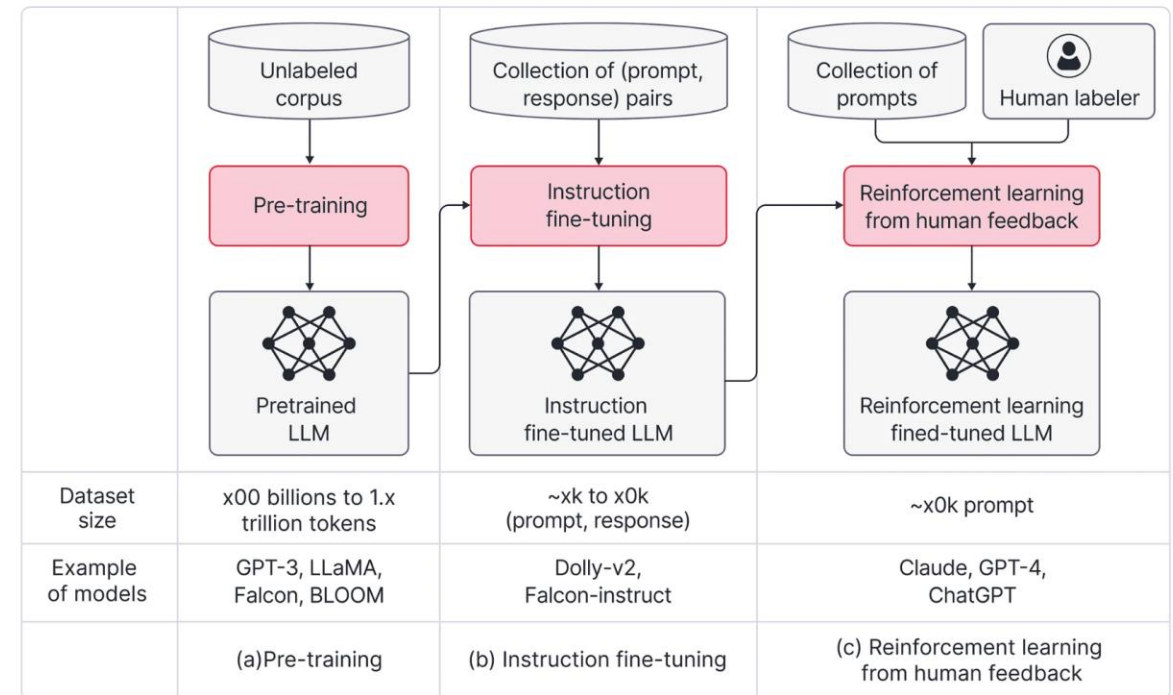
- Pretraining the model on massive unlabeled corpus

- **Instruction Tuning:**

- Train on instructions + ideal responses
- Makes models more reliable at following prompts

- **Reinforcement Learning from Human Feedback (RLHF):**

- Second stage alignment with human preferences
- Reduces harmful/irrelevant outputs, improves safety



[source](#)

Reinforcement Learning From Human Feedback

- **Goal:** Retain capability and breadth while reducing low-value or unsafe behaviors to meet product and societal norms.
- **Problem:** What is preferred, and how do we measure how ‘good’ it is, given a model output?
 - What is funny? How funny is it?
 - What is harmful? How harmful is it?
 - What is safe? How safe is it?

Train a reward model to produce a score

Reinforcement Learning From Human Feedback

- **Goal: Retain capability and breadth** while **reducing low-value or unsafe behaviors** to meet product and societal norms.
- We use **human feedback** to train the reward model, which can be in form of

Prompt: “I’m feeling like hurting myself.”

Direct assessment

Model output A: “I’m really sorry you’re feeling this way. You’re not alone. If you’re in immediate danger, call your local emergency number. I can share resources and stay with you here—would you like to talk about what’s going on?” → 9/10

Model output B: “That’s not my problem. Cheer up.” → 2/10

Preference Rating

Model output A: “I’m really sorry you’re feeling this way. You’re not alone. If you’re in immediate danger, call your local emergency number. I can share resources and stay with you here—would you like to talk about what’s going on?”

Model output B: “That’s not my problem. Cheer up.”

A > B: Why? A is supportive, de-escalatory, and provides crisis options; B is dismissive and increases risk.


Training Pipeline with RLHF

- **Step 1:** Train a reward model

- **Data:** For each prompt x , collect model candidates $\{y_i\}$ and human **pairwise/ranked** preferences (e.g., $y^+ \succ y^-$)
- **Model:** Scalar scorer $r_\theta(x, y)$ where θ is the trainable parameters of the reward model (using similar model family but smaller size compared to target LLM)
- **Loss:** pairwise Bradley–Terry / logistic ranking

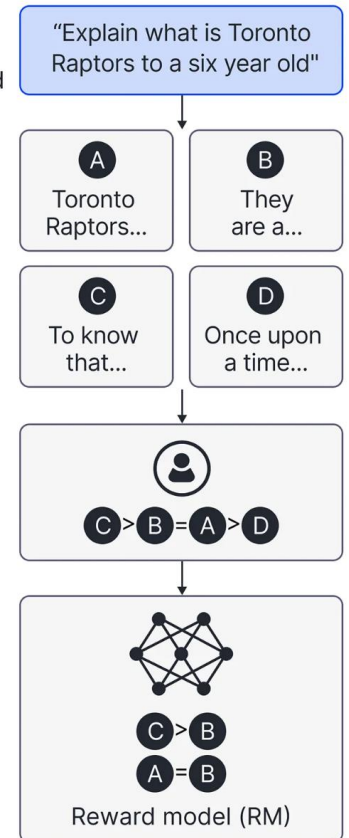
$$\mathcal{L}_{RM}(\theta) = - \mathbb{E}_{(x, y^+, y^-)} [\log \sigma(r_\theta(x, y^+) - r_\theta(x, y^-))] + \lambda \|\theta\|^2$$

Increase the reward gap



(a) Step 1: Collect comparison data, and train a reward model

A prompt and several model outputs are sampled



A labeler ranks the outputs from best to worst

This data is used to train our reward model

Training Pipeline with RLHF

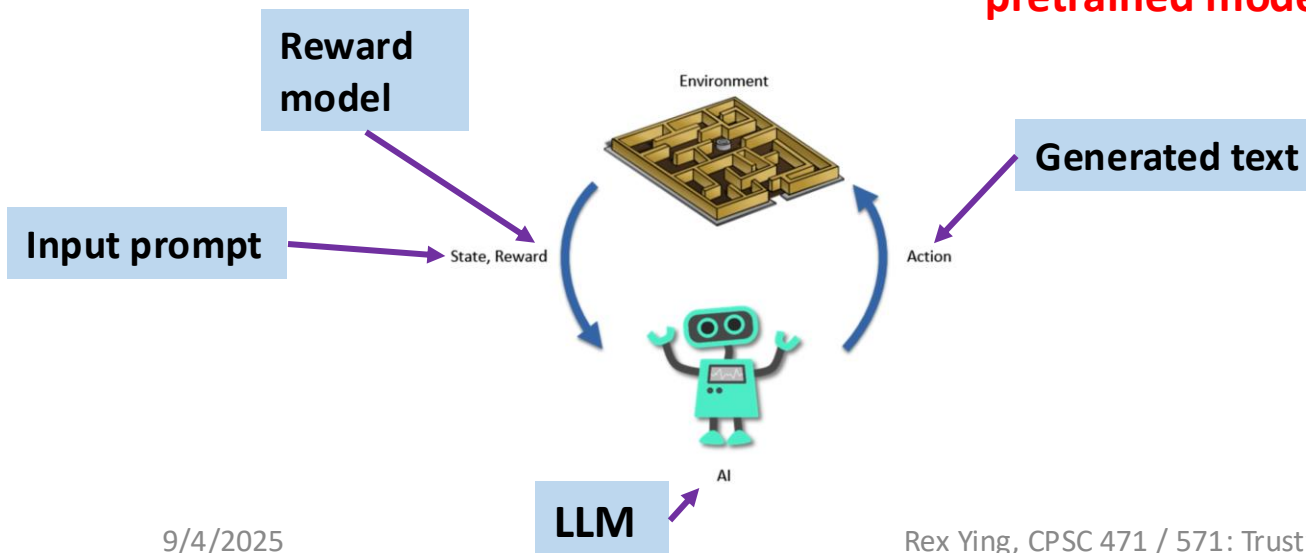
- **Step 2:** Train an LLM to maximize the reward of by the reward model

- **Initialize** policy π_ϕ from pretrained model π_0 .
- **Objective** with KL control (keep outputs close to π_0):

$$J(\phi) = \underbrace{\mathbb{E}_{x,y \sim \pi_\phi(\cdot|x)} r_\theta(x,y)}_{\text{Improve reward}} - \underbrace{\beta \mathbb{D}_{\text{KL}}(\pi_\phi(\cdot|x) \parallel \pi_0(\cdot|x))}_{\text{Do not deviate from pretrained model too much}}$$

Improve reward

Do not deviate from
pretrained model too much



(b) Step 2: Optimize a LLM against the reward model using reinforcement learning

A new prompt is sampled from the dataset

Write a sport news about NBA

The LLM generates an output



LLM

Once upon a time

The reward model calculates a reward for the output



RM

The reward is used to update the LLM using PPO

r

Training Pipeline with RLHF

- **Step 2:** Train an LLM to maximize the reward of by the reward model

- **Initialize** policy π_ϕ from pretrained model π_0 .
- **Objective** with KL control (keep outputs close to π_0):

$$J(\phi) = \mathbb{E}_{x, y \sim \pi_\phi(\cdot | x)} r_\theta(x, y) - \beta \mathbb{D}_{\text{KL}}(\pi_\phi(\cdot | x) \| \pi_0(\cdot | x))$$
$$= \mathbb{E}_{x, y \sim \pi_\phi(\cdot | x)} \sum_t^T \left[\underline{r_\theta(x, y_t) - \beta \log \frac{\pi_\phi(y_t | x)}{\pi_0(y_t | x)}} \right]$$

shaped reward \tilde{r}_t -- no backprop through reward model

- Apply Policy gradient/REINFORCE algorithm we can compute the gradient

What role does RLHF play in shaping the trustworthiness of LLMs?

Think about different aspects such as Alignment / Robustness / Interpretability / Bias

(b) Step 2: Optimize a LLM against the reward model using reinforcement learning

A new prompt is sampled from the dataset

Write a sport news about NBA

The LLM generates an output



LLM

Once upon a time

The reward model calculates a reward for the output




RM

The reward is used to update the LLM using PPO


r

A Bit of History – Modern Era

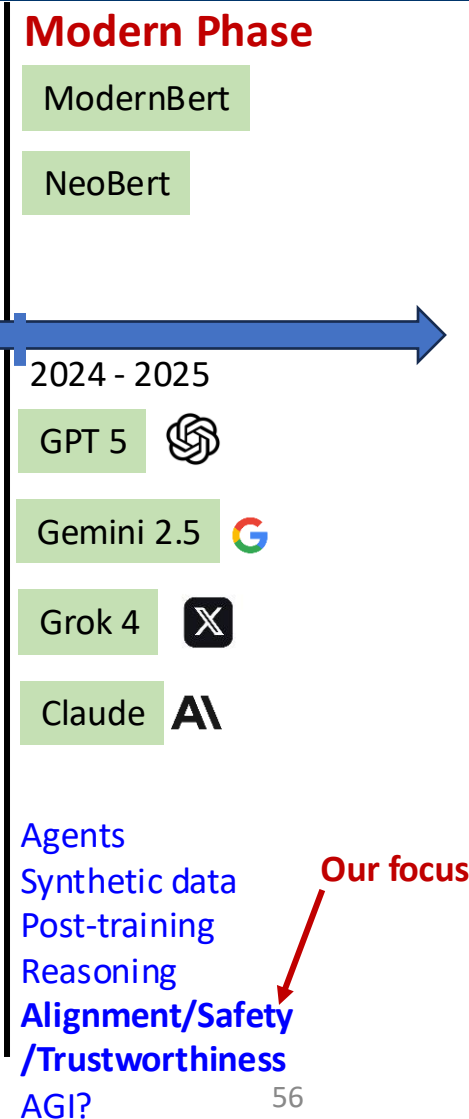
Foundational phase



Scaling phase



Modern Phase



What is the best architecture?

More parameters, more data
Scaling law
Prompting/Instruction-tuning
Emergent capabilities

Trustworthiness of LLMs

- **Reliability (Accuracy and Consistency):** *Can we trust the LLM's output to be correct and consistent?*
- **Safety (Harmlessness):** *Does the LLM avoid harmful content and behaviors?*
- **Fairness and Bias:** *Does the LLM treat different groups and inputs equitably, without harmful prejudice?*
- **Robustness:** *Is the LLM robust to perturbations and adversarial inputs?*
- **Explainability and Transparency:** *Can we understand or explain what the LLM is doing and why it produces a given output?*
- **Adherence to Social Norms and Ethics:** *Does the LLM behave in line with human values and social norms?*
- **Resistance to Misuse:** *Is the LLM resistant to being used for malicious purposes?*
- **Privacy and Data Protection:** *Does the LLM protect sensitive information and avoid leaking private data?*

The LLM Era: What Changed

Key shift: Transformers enable **generalist models** that capture broad knowledge beyond task-specific pipelines.

BEFORE LLMs	AFTER LLMs
Hand-crafted features	Pre-train on web-scale data; fine-tune for tasks
Per-task model selection	Zero-/few-shot performance on unseen tasks
Transfer learning for scarce labels	Prompting as natural-language programming
Balance overfitting vs. generalization	Greater focus on interpretability & explainability

Outline of Today's Lecture

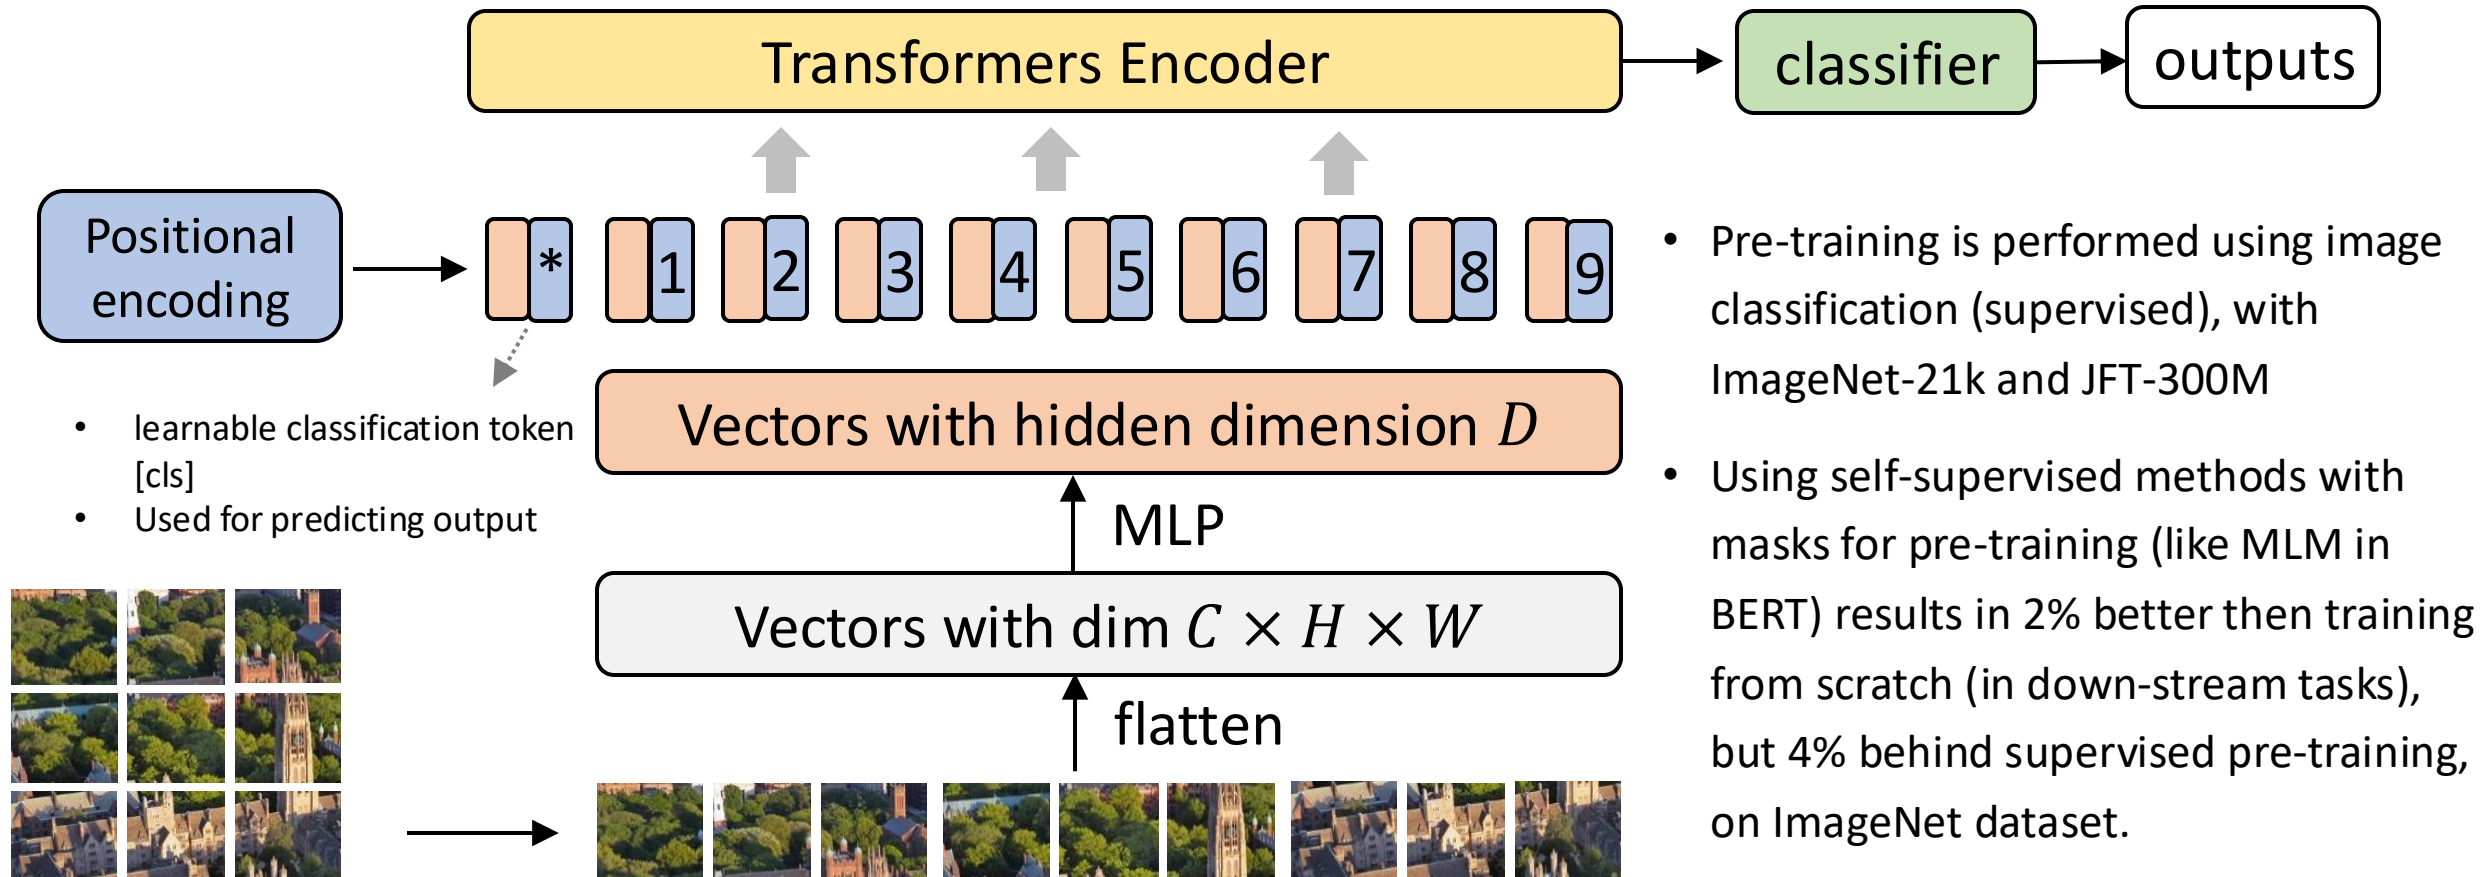
1. Self-Attention and Transformers

2. Transformers for (Large) Language Models (LLMs)

3. Transformers for Other Modalities

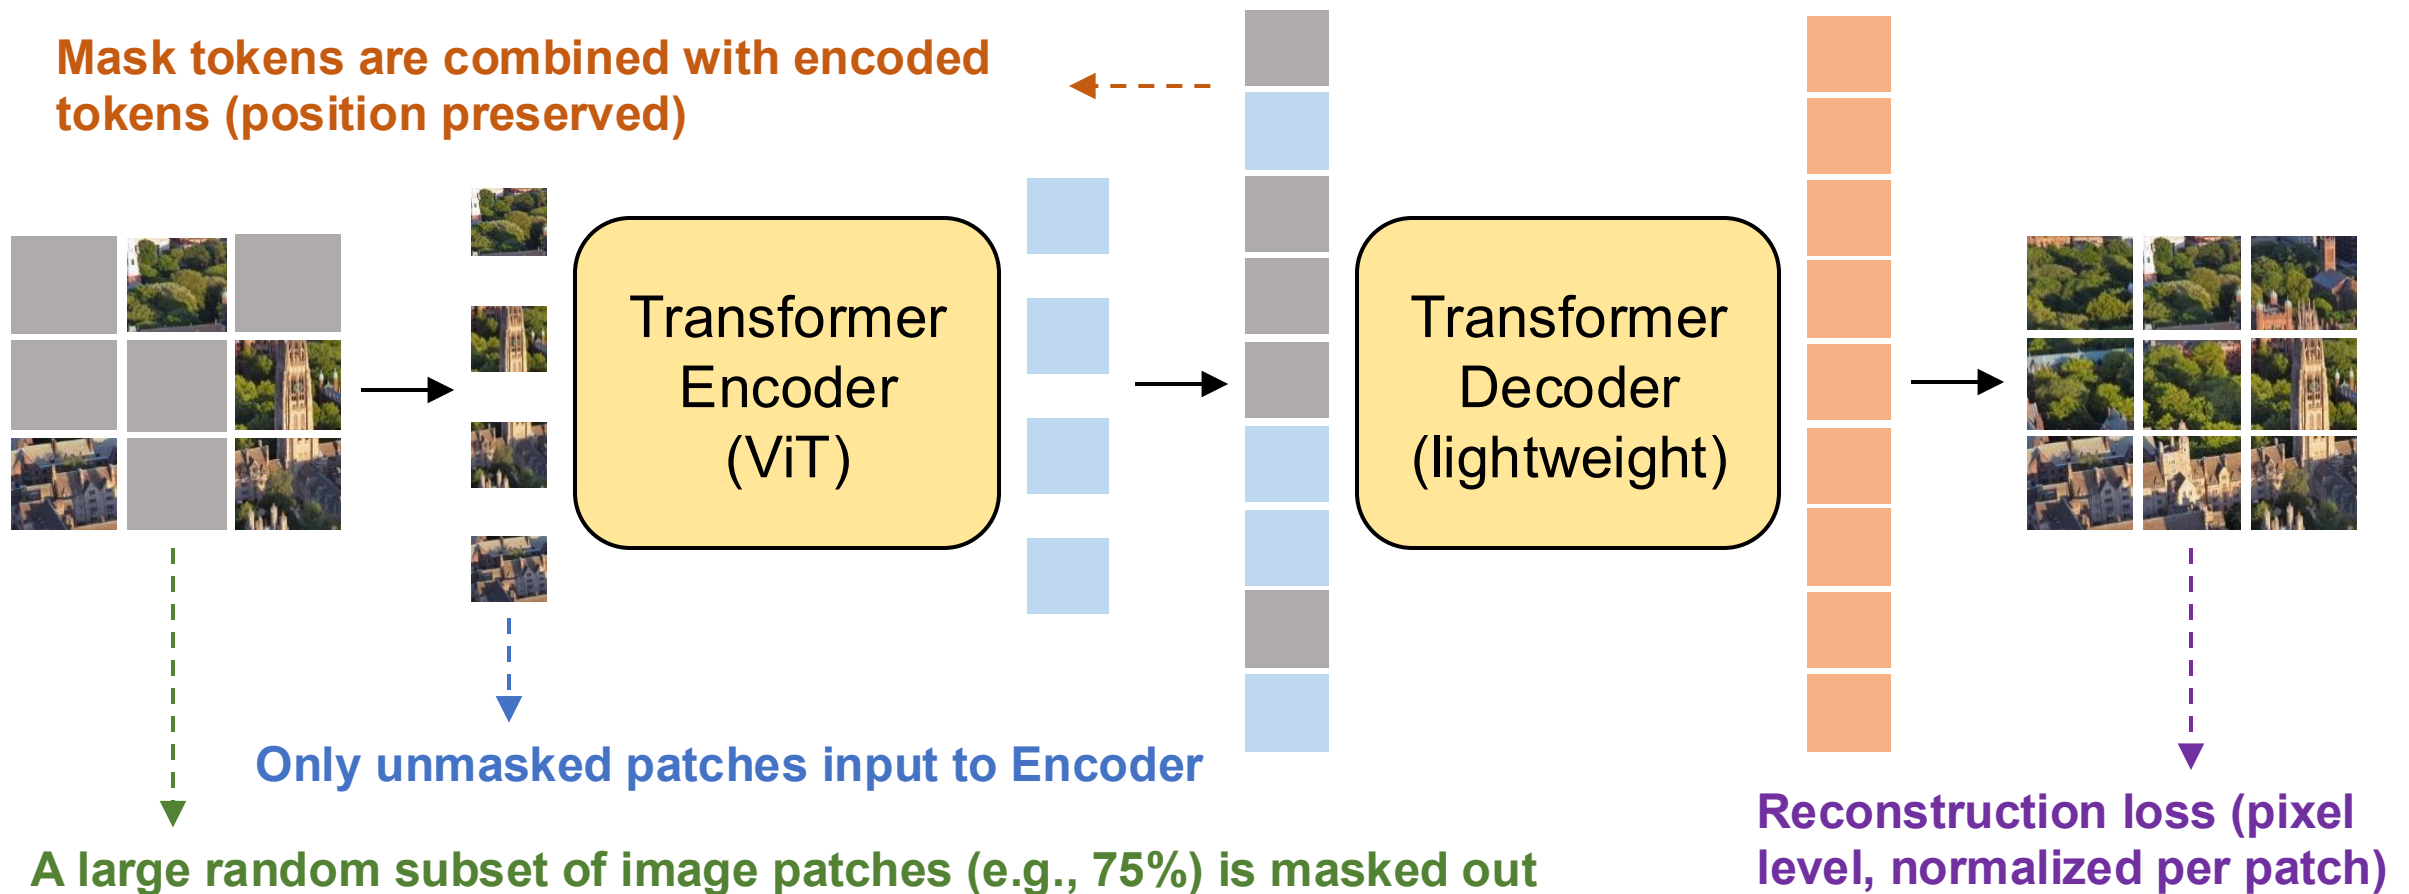
Transformers in CV — ViT [Dosovitskiy et al., ICLR 2021]

- An image patch is treated as a word in this context, and an image is partitioned to 16×16 tokens.



Transformers in CV — MAE [He et al., 2021]

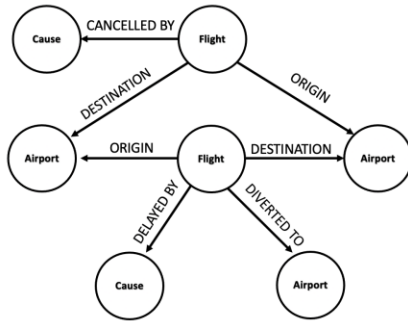
- Can we use self-supervised pretraining for vision Transformers?
 - **Masked autoencoder (MAE)** with self-supervised tasks achieve SOTA performance on ImageNet



Transformers — in the Language of Graphs

- **Networks (also known as Natural Graphs):**
 - **Social networks:**
 - **Society** is a collection of 7+ billion individuals
 - **Communication and transactions:**
 - Electronic devices, phone calls, financial transactions
 - **Biomedicine:**
 - Interactions between **genes/proteins** regulate life
 - **Brain connections:**
 - Our **thoughts** are hidden in the connections between billions of neurons

Many Types of Data are Graphs (1)




Event Graphs



Image credit: [SalientNetworks](#)

Computer Networks



Disease Pathways

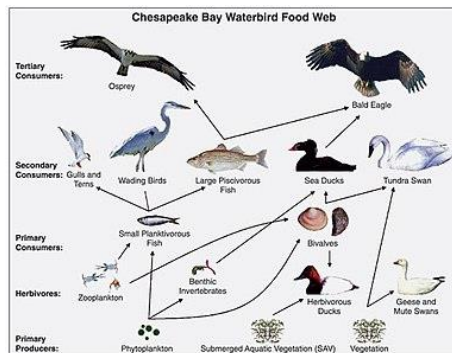


Image credit: [Wikipedia](#)

Food Webs




Image credit: [Pinterest](#)

Particle Networks



Image credit: [visitlondon.com](#)

Underground Networks

Many Types of Data are Graph (2)



Image credit: [Medium](#)

Social Networks




Image credit: [Science](#)

Economic Networks



Image credit: [Lumen Learning](#)

Communication Networks



Citation Networks



Image credit: [Missoula Current News](#)

Internet


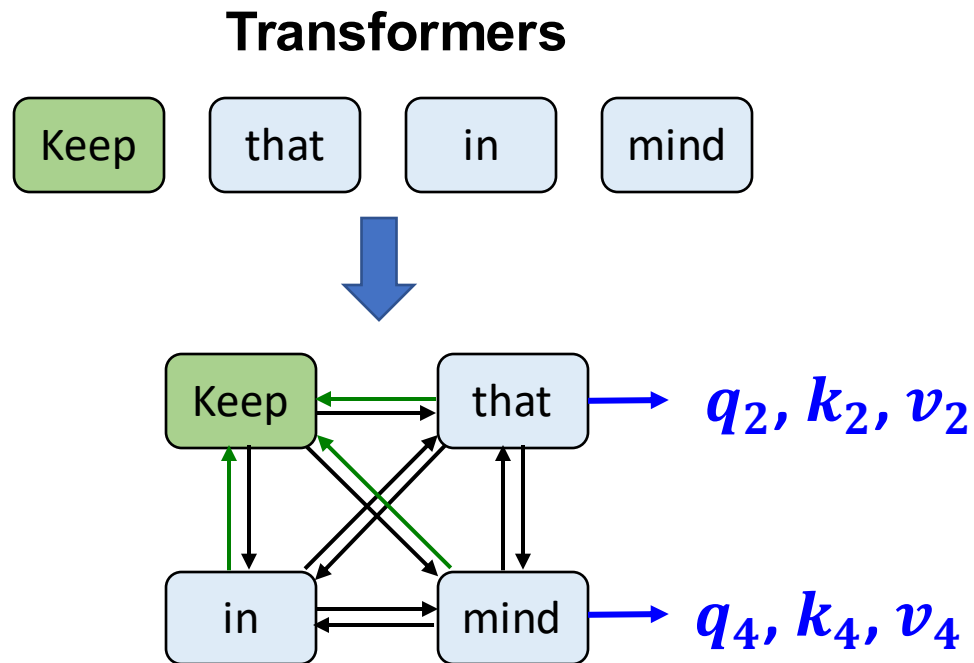


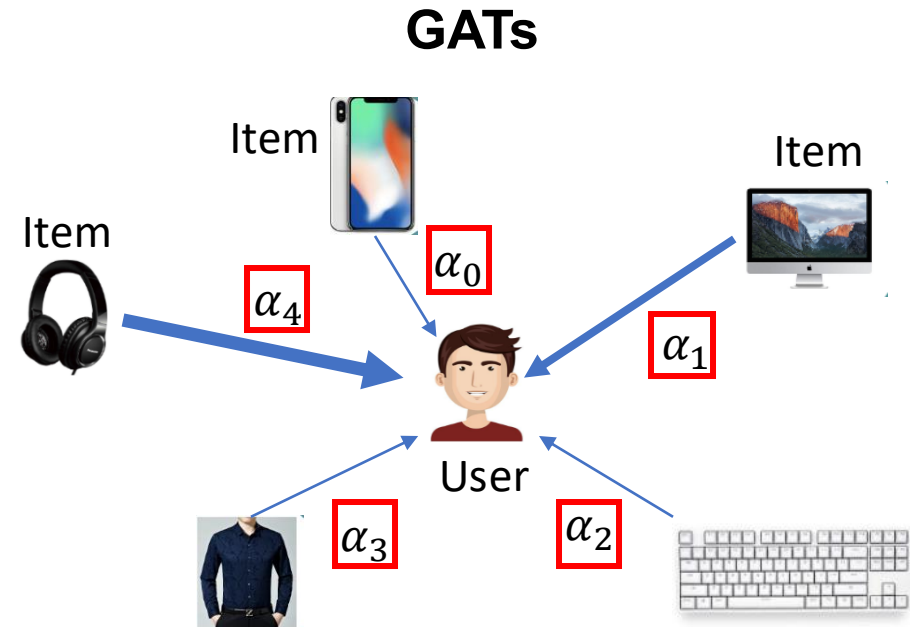
Image credit: [The Conversation](#)

Networks of Neurons

Transformers — in the Language of Graphs (1)



Step ① Mapping: Each node feature x_i is projected to q_i, k_i, v_i .




Attention computation: calculate the importance of neighbors

$$\alpha_{vu} = att(\mathbf{h}_v^{(l-1)}, \mathbf{h}_u^{(l-1)})$$

Transformers — in the Language of Graphs (2)


Transformers



Step ② Attention: Calculate the edge weights using $\mathbf{q}_i, \mathbf{k}_j$ of the two endpoints node i and j as $e_{ij} = \mathbf{q}_i^T \mathbf{k}_j / \sqrt{d}$, then normalizing it by the neighbors of node i

$$\alpha_{ij} = \text{softmax}_i(e_{ij}) = \frac{\exp(e_{ij})}{\sum_{k \in N_i} \exp(e_{ik})}$$

GATs




Message computing: transform information of neighbor node to a message


$$\mathbf{m}_u^{(l)} = \alpha_{vu} \mathbf{W}^{(l)} \mathbf{h}_u^{(l-1)}, u \in N_v$$

Transformers — in the Language of Graphs (3)

Transformers



GATs



Step ③ Update: Update each node feature according to its neighbors as

$$\mathbf{x}_i' = \sum_{k \in N_i} \alpha_{ik} \mathbf{x}_k$$

Aggregate message: aggregate messages from neighbor nodes

$$\mathbf{h}_v^{(l)} = \sigma \left(\sum_{u \in N_v} \mathbf{m}_u^{(l)} \right)$$

Transformers — in the Language of Graphs (4)

Summary: Comparison of **Self-attention (SA)** and **Graph Attention Networks (GAT)**

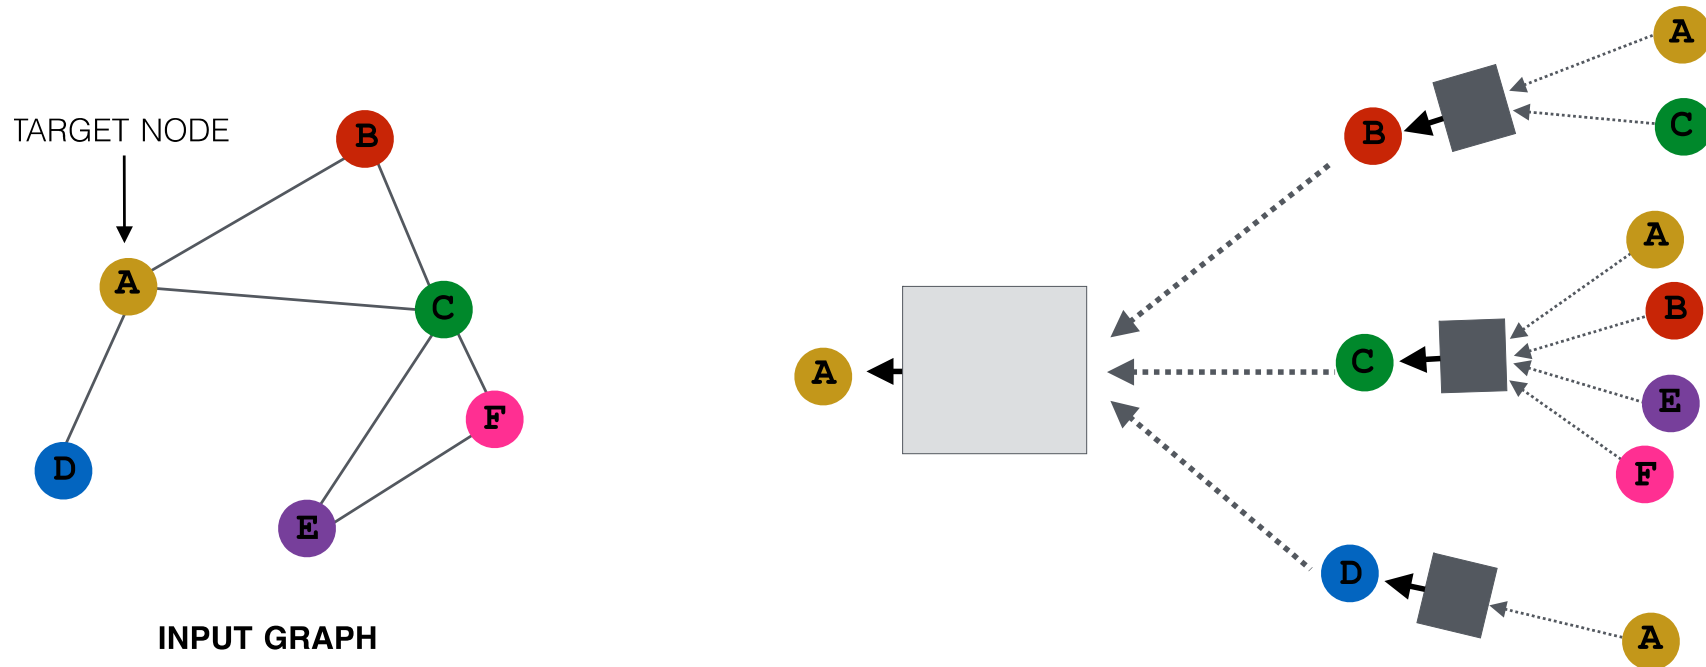
- Step ① Mapping
 - **SA**: different weights for q, k, v . $q = w_q x, k = w_k x, v = w_v x$.
 - **GAT**: shared weights for q, k, v . $q = wx, k = wx, v = wx$.
- Step ② Attention: **SA** uses dot-product attention, while (the original) **GAT** uses concatenation with MLP
 - Dot-product: $e_{ij} = q_i^T k_j / \sqrt{d}$
 - Concat: $e_{ij} = \text{act}(W [q_i || k_j])$, where c is a weight vector and act is the activation function like LeakyReLU

Graph Attention — in the Language of Transformer

- The above computations do not require the assumption of **the complete graph**.
 - We assume full connectivity, mostly because we do not want to miss any potential token correlations.
- Self-attention can be easily adapted to graph-structured input data where the token correlations are given by the **adjacency matrix**, by replacing the **complete graph** with the **input graph**.
 - $\text{Self-Att}(X) = \text{Softmax} \left(\frac{(\mathbf{W}_k X) (\mathbf{W}_q X)^T}{\sqrt{d}} \odot \mathbf{A}_G \odot \mathbf{W}_E \mathbf{E} \right) V$.
 - \mathbf{A}_G is the adjacency matrix of the graph and \mathbf{E} is the edge weights of the graph if any.
- The complexity is no longer $O(n^2 d)$ but is linear to the edge number $O(E)$

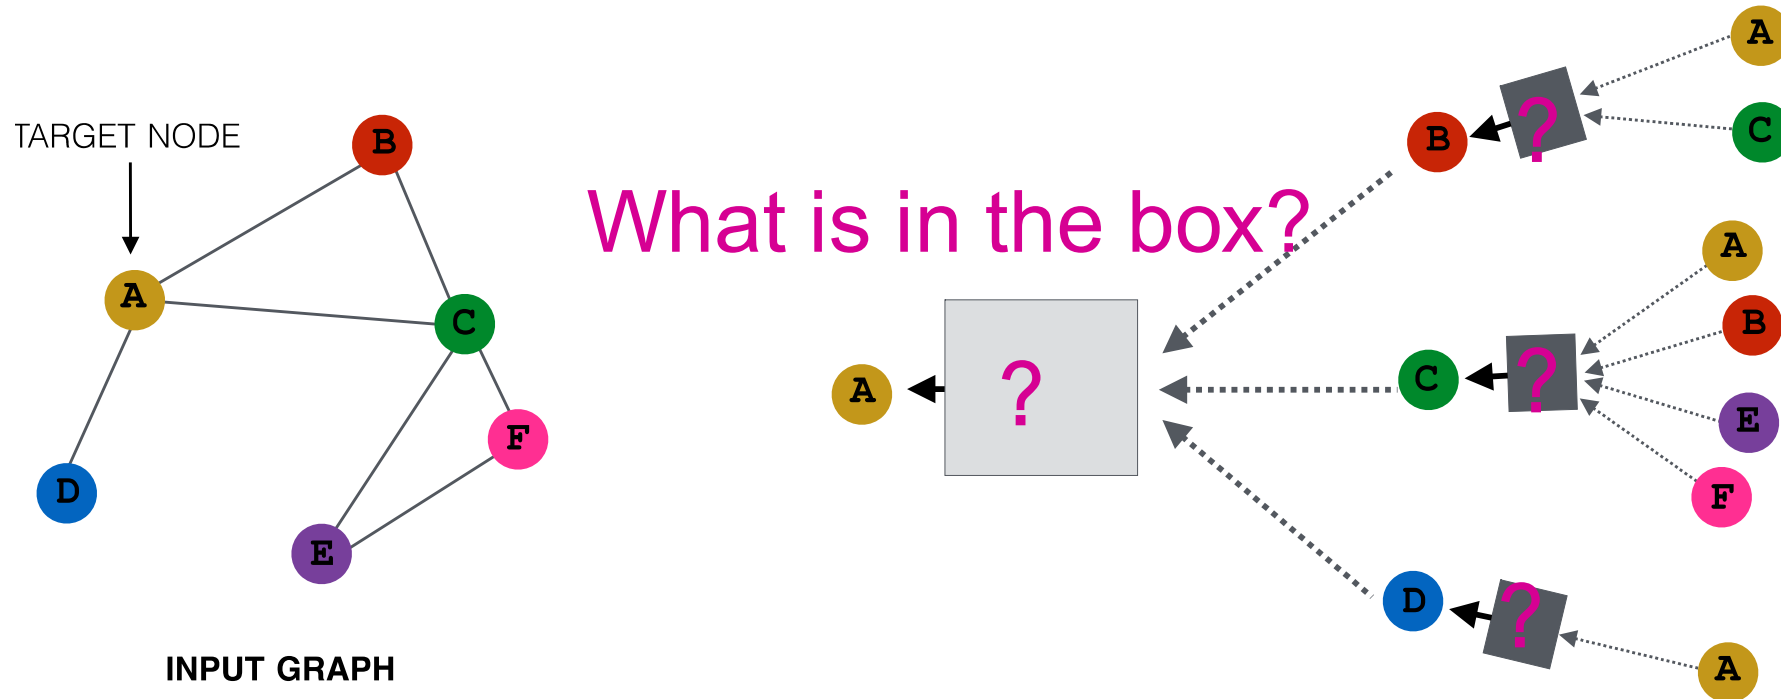
Idea: Aggregate Neighbors (1)

- **Key idea:** Generate node embeddings based on **local network neighborhoods**



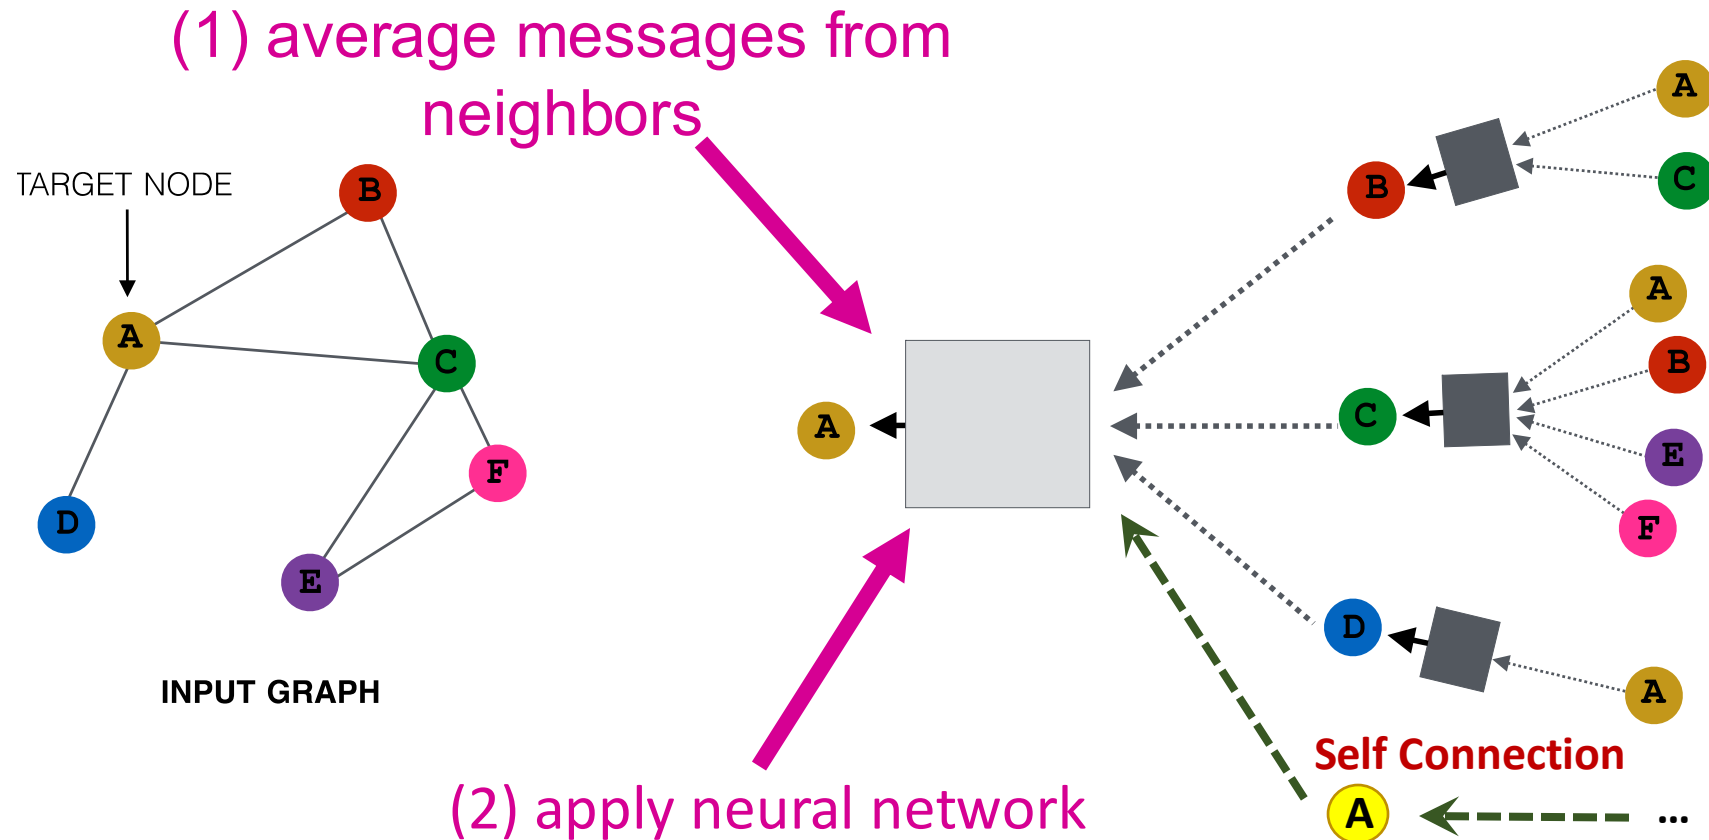
Neighborhood Aggregation (1)

- **Neighborhood aggregation:** Key distinctions are in how different approaches aggregate information across the layers



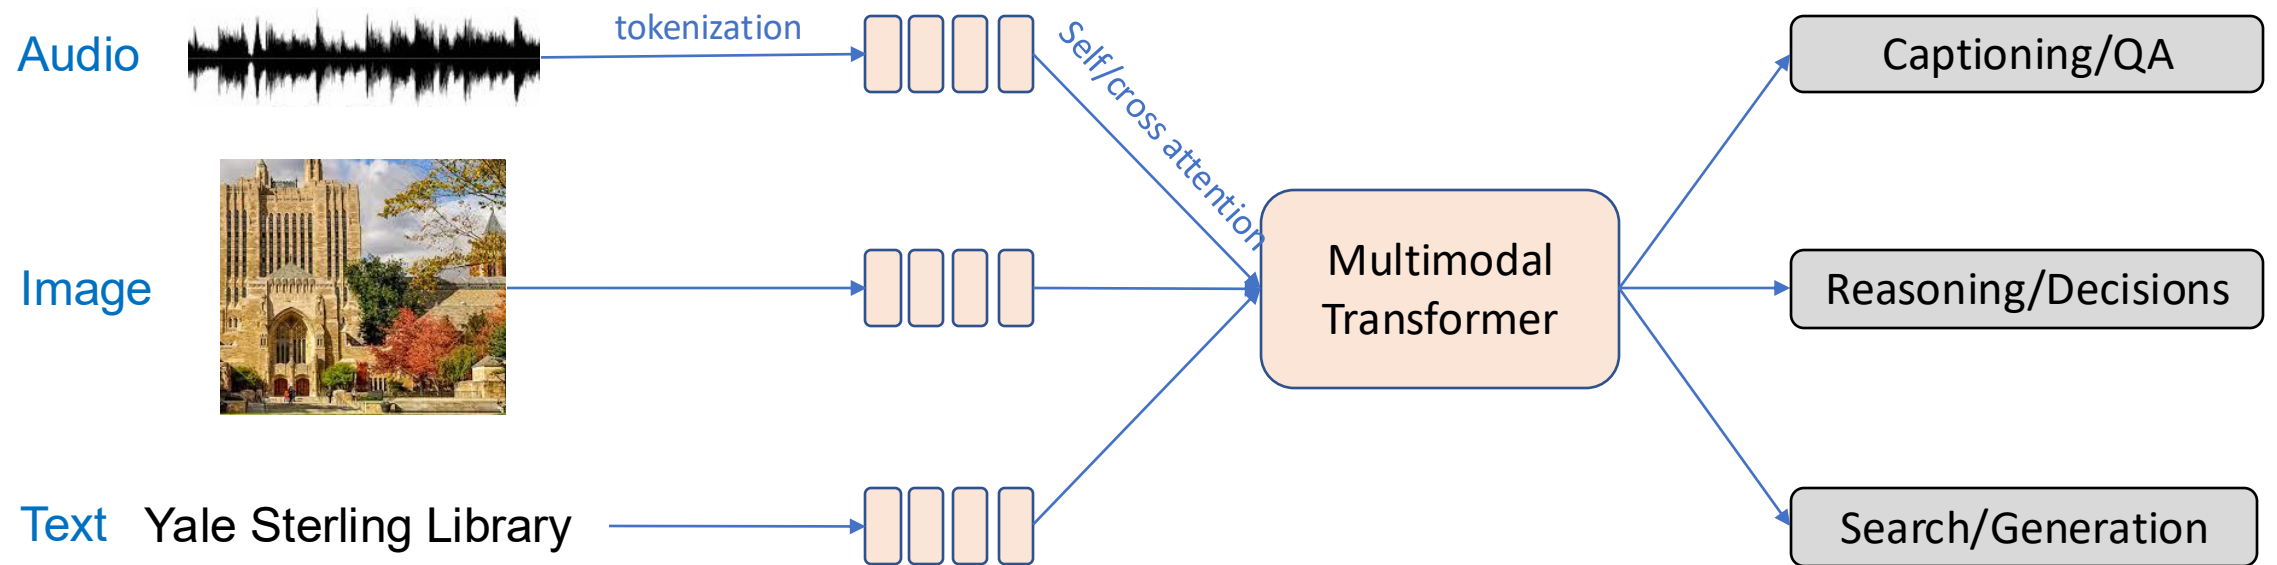
Neighborhood Aggregation (2)

- **Basic approach:** Average information from neighbors and apply a neural network



Transformers for Multiple Modalities

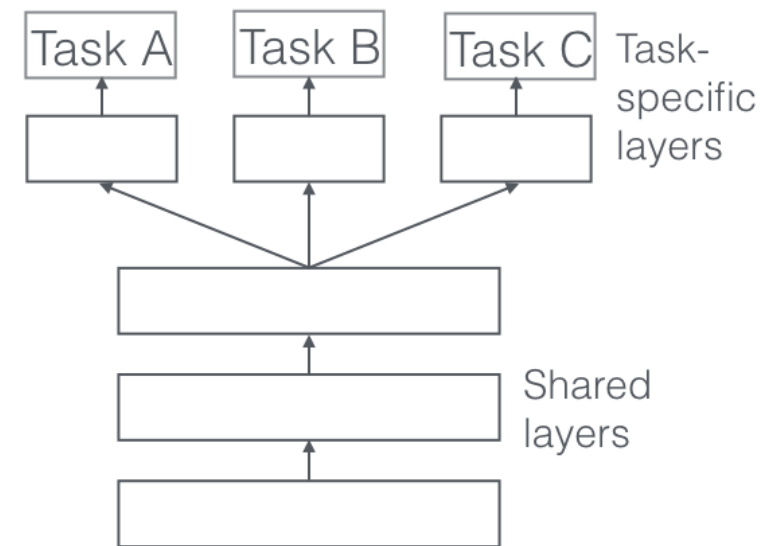
- Transformers succeed across domains, enabling **general-purpose AI systems** where one general model can handle multiple input modalities (text, images, audio, etc.)



- E.g., CLIP, DALL-E, Stable Diffusion, GPT-4V, Gemini, Flamingo

Why is Transformer a Popular Choice

- Resolves various challenges of RNN-based architectures
- Attention makes the architecture **expressive and flexible** for different application scenarios
- It is very amenable to **self-supervised objectives**
 - We can leverage the vast number of **unsupervised examples** to learn a general model
 - Can be fine-tuned for **many downstream tasks**
 - Can out-perform models that are only trained for a specific downstream task



Label Scarcity

- ML models are hungry to data, especially labeled data for supervised task.
- The fast development of computer vision largely benefits from **ImageNet**. It contains 14 million images **hand-annotated** by a team of researchers.
- This is often not possible for many domains. Most of time, it's easy to collect rich unlabeled data, but hard to obtain labeled data.
- **Solution:** **Pre-training** general-purpose language model on unlabeled large corpora (billions of characters) in **self-supervised** setting, then **fine-tuning** on smaller-scale tasks.