

# Transformers

CPSC483: Deep Learning on Graph-Structured Data

Rex Ying

# Questions

- How to summarize what it means for an ML system to be **trusted**?
- Name one of the four major characteristics of a **trustworthy** ML system according to the book's opinion.

Explain what does it mean and why it matters

- Have you noticed any news, articles, policies, events that have implications in trustworthy deep learning in recent years?

# Readings

- Readings are updated on the website (syllabus page)
- **Readings:**
  - [\[2401.05561\] TrustLLM: Trustworthiness in Large Language Models](#)
- This lecture is not explicitly tested
  - But in future lectures we will assume knowledge of this when developing trustworthy components on top of Transformers


# Outline of Today's Lecture

## **1. Self-Attention and Transformers**

## 2. Transformers for (Large) Language Models (LLMs)

## 3. Transformers in Other Modalities


# Sequence Learning



$$h_i = f_W(x_i, h_{i-1}), y_i = f_Y(h_i)$$

**What are the issues and challenges of RNNs?**

# Sequence Learning




$$h_i = f_W(x_i, h_{i-1}), y_i = f_Y(h_i)$$

## Problems of RNNs

- Sequential computation prevents parallelization
- Capacity of handling long sequences
- Mainly focusing on modeling recurrence
  - does not capture other correlations (hierarchical, long-range, polysemy.... ) well

# Sequence Learning




$$h_i = f_W(x_i, h_{i-1}), y_i = f_Y(h_i)$$

## Problems of RNNs

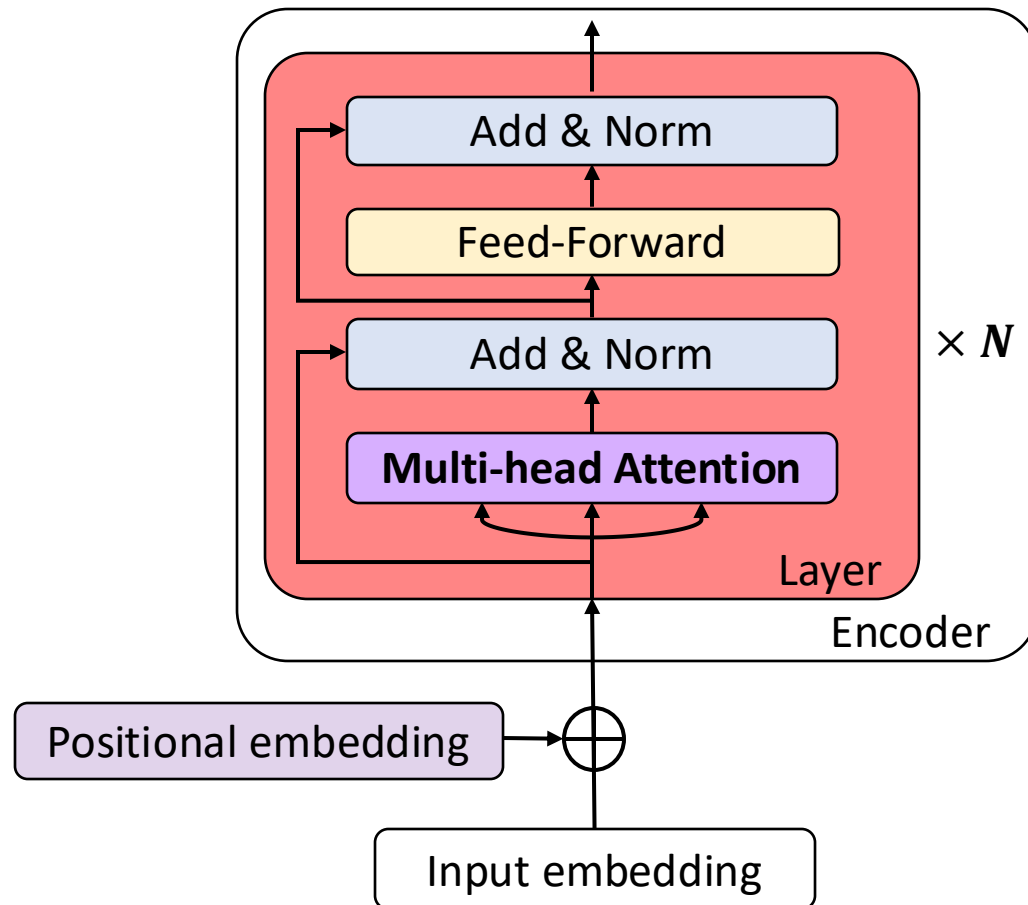
1. parallelization — — — — — ➔
2. long sequences — — — — — ➔
3. only recurrence — — — — — ➔

## Solutions by Transformers

1. **Parallel input:** Input All tokens at the same time
2. **Self-Attention:** Enable attention in long-range
3. **Positional Embeddings  $p_i$ :** Model all possible correlations




# Transformers — Overview



- **Original paper:** Attention is all you need [Vaswani et al., 2017].
- **Key component:** Multi-head self-attention
- **Other components** of a transformer layer: layer normalization, skip connection, position-wise feed-forward layer (FFN, or MLP)
- **Model usage:** Pre-training followed by fine-tuning. The transferred model can be:
  - **Encoder-only** (e.g BERT)
  - **Encoder-Decoder** (e.g [BART](#))
  - **Decoder-only** (e.g GPT)
  - We will show an example later




# Transformers — Overview



- **Model usage:** Pre-training followed by fine-tuning. The transferred model can be:
  - **Encoder-only** (e.g BERT)
    - Many-to-one classification / regression
    - Sentiment classification, document classification ...
    - Word / Sentence embeddings for downstream tasks (e.g. recommender system)
  - **Encoder-Decoder** (e.g [BART](#))
  - **Decoder-only** (e.g GPT)
  - We will show an example later

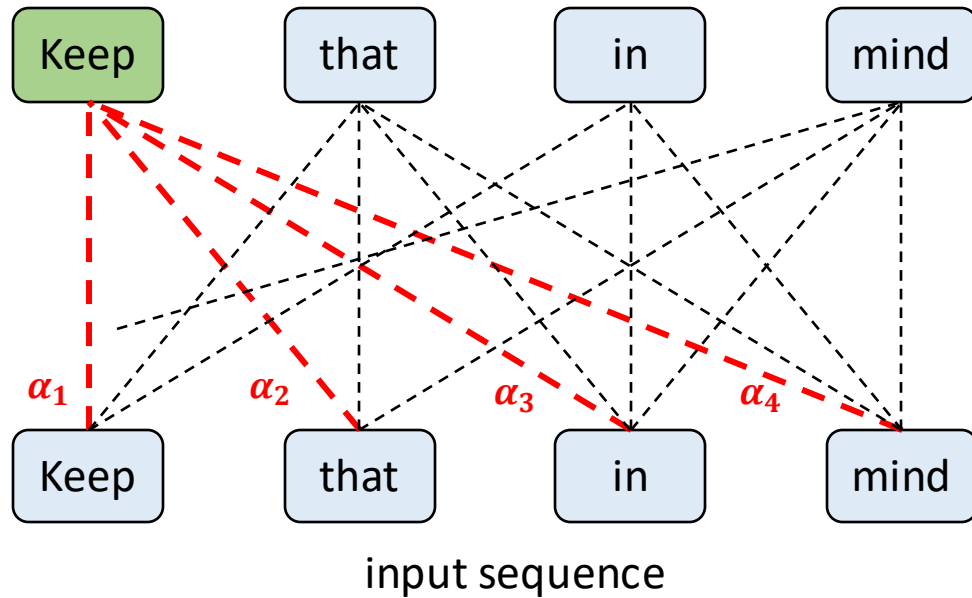
# Transformers — Overview



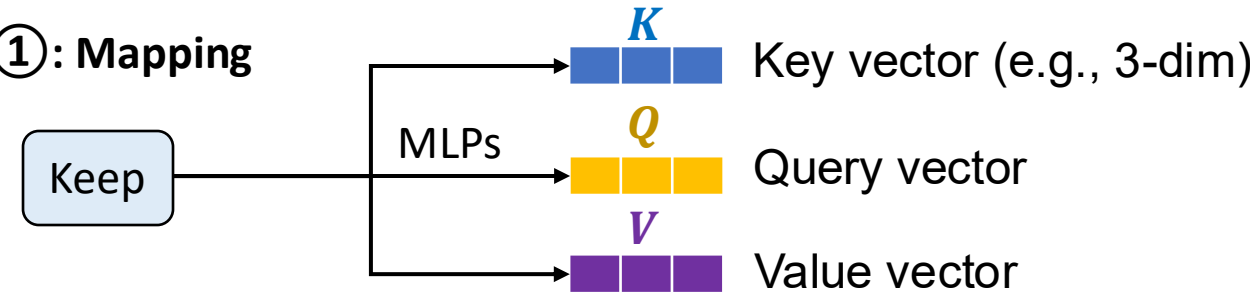
- **Model usage:** Pre-training followed by fine-tuning. The transferred model can be:
  - **Encoder-only** (e.g BERT)
  - **Encoder-Decoder** (e.g [BART](#))
    - Many-to-many use cases
    - Summarization, translation, style transfer ...
  - **Decoder-only** (e.g OpenAI GPT)
    - One-to-many use cases
    - Image / text / code generation, dialogue systems ...
    - GPT-3/4 based [apps](#)

# Transformers — Self-Attention (1/5)

## Example:



## Step ①: Mapping



## Step ②: Attention



$$\alpha_1, \alpha_2, \alpha_3, \alpha_4 = \text{Softmax} \left( \begin{matrix} Q \\ \text{Keep} \end{matrix} \times \begin{matrix} K & K & K & K \\ \text{Keep} & \text{that} & \text{in} & \text{mind} \end{matrix} \right)$$

## Step ③: Update

$$\begin{matrix} V' \\ \text{Keep} \end{matrix} = \alpha_1 \times \begin{matrix} V \\ \text{Keep} \end{matrix} + \alpha_2 \times \begin{matrix} V \\ \text{that} \end{matrix} + \alpha_3 \times \begin{matrix} V \\ \text{in} \end{matrix} + \alpha_4 \times \begin{matrix} V \\ \text{mind} \end{matrix}$$

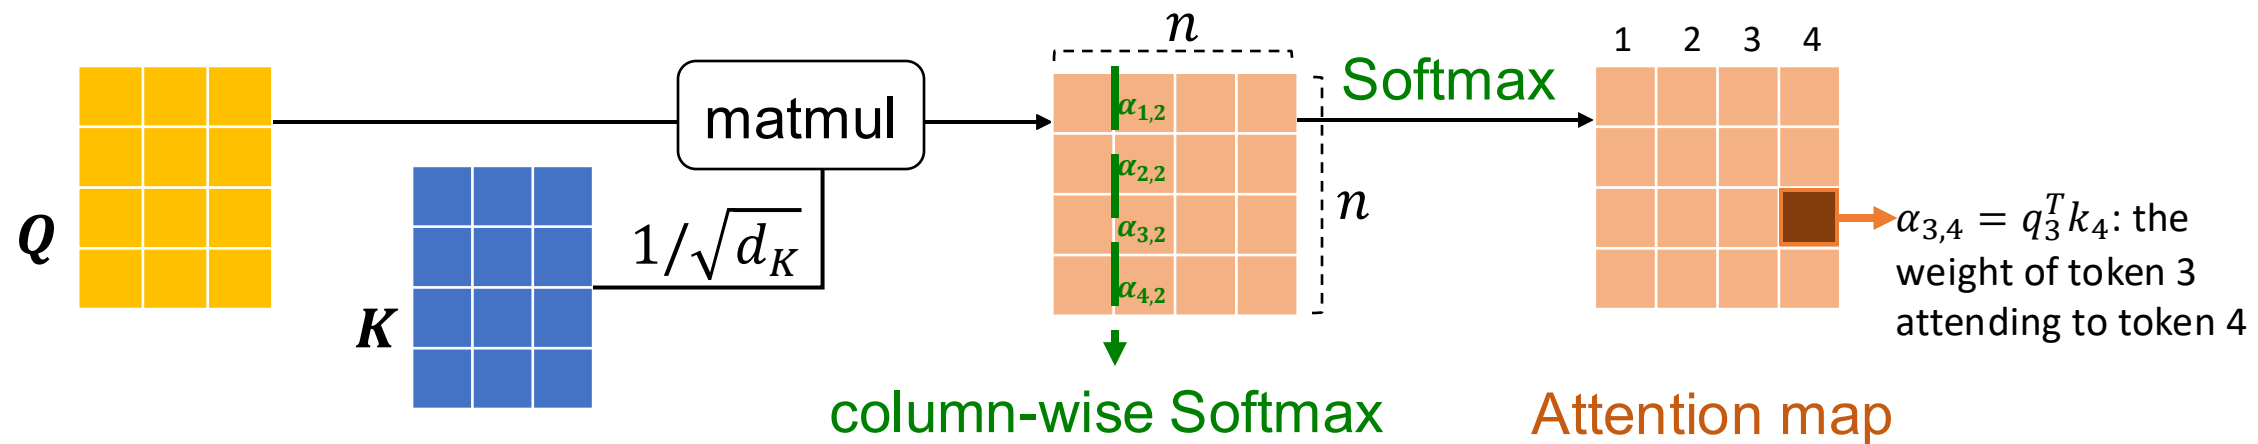
# Transformers — Self-Attention (2/5)

- Formally, given an input sequence  $X = [x_1, x_2, \dots, x_n] \in \mathbb{R}^{n \times d_X}$
- Step ①: Query  $Q = XW_Q$ , Key  $K = XW_K$ , Value  $V = XW_V$ 
  - $W_K \in \mathbb{R}^{d_X \times d_K}$ , and thus  $K \in \mathbb{R}^{n \times d_K}$
  - We require  $d_K = d_Q$ , for simplicity, we set  $d_K = d_Q = d_V := d$



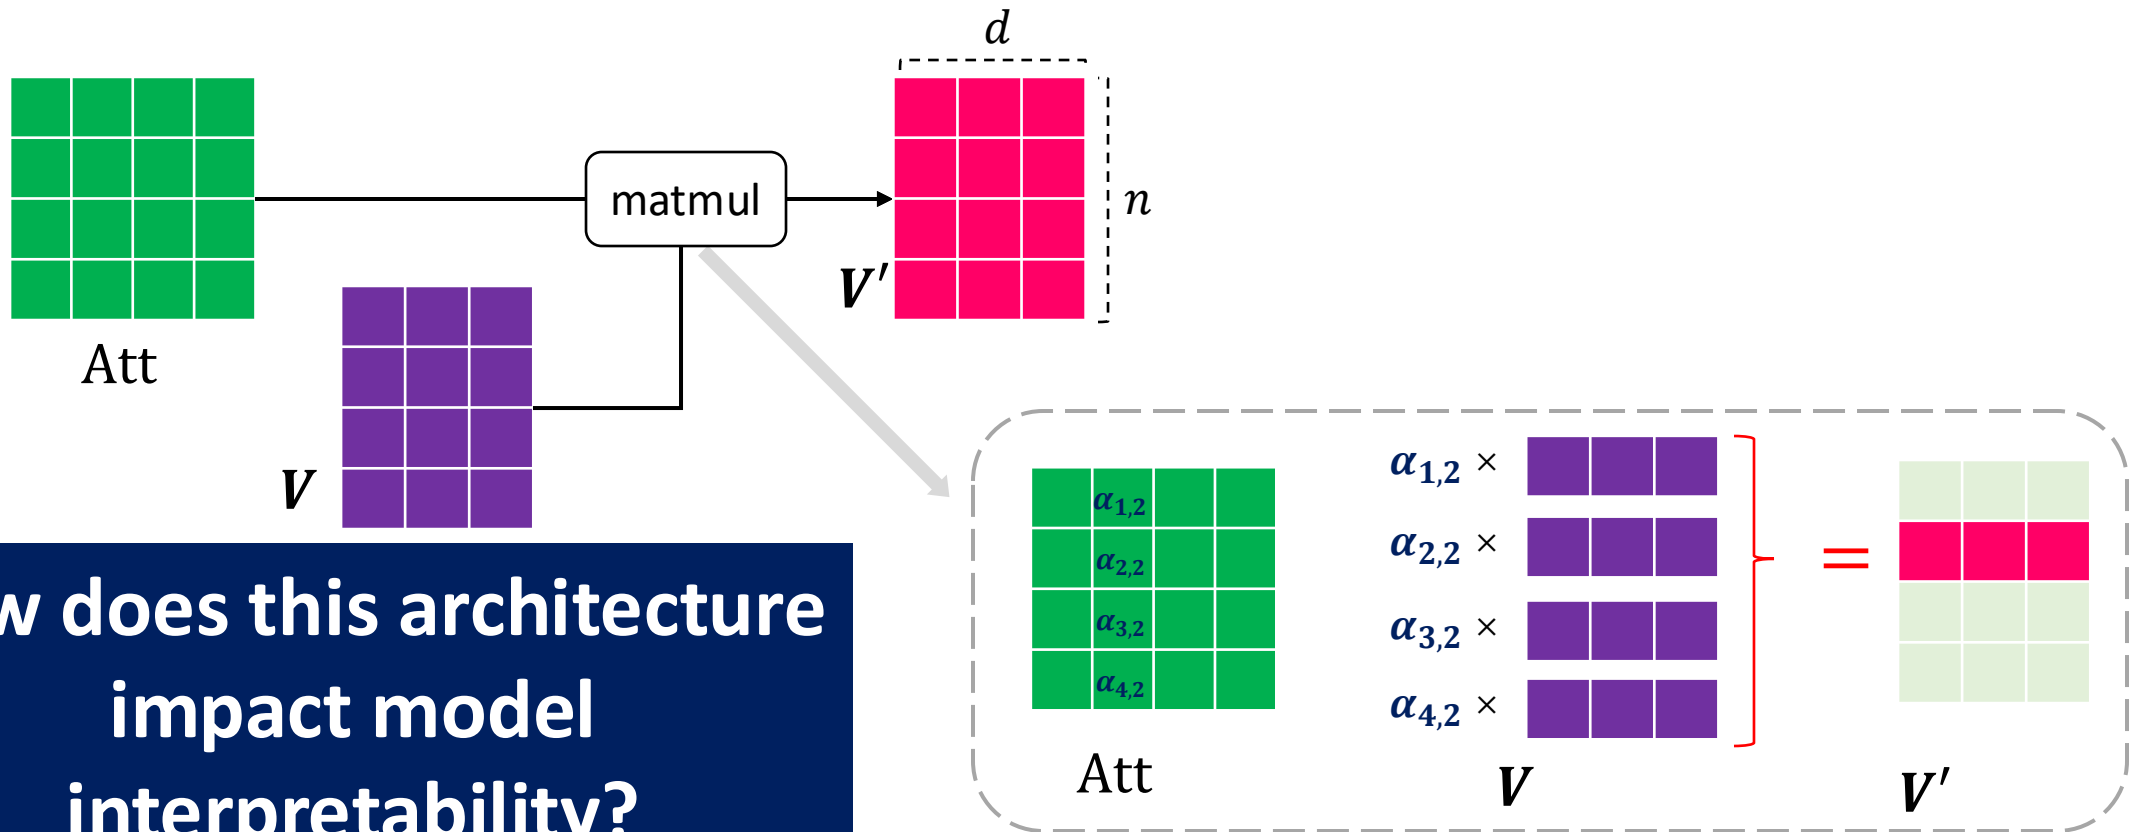
# Transformers — Self-Attention (3/5)

- Step ② : Attention map  $\text{Att} = \text{Softmax}\left(\frac{QK^T}{\sqrt{d}}\right) \in \mathbb{R}^{n \times n}$  (Softmax is col-wise)
  - The matrix multiplication  $QK^T$  performs dot-product for every possible pair of queries and keys, resulting in an attention map.
  - **Normalization factor**  $1/\sqrt{d_K}$  : performing dot-product over two vectors with variance  $\sigma^2$  results in a scalar having  $d_K$ -times higher variance,
    - $q \sim N(0, \sigma^2), k \sim N(0, \sigma^2) \rightarrow \text{Var}\left(\sum_{i=1}^{d_K} q[i]k[i]\right) = \sigma^4 d_K$



# Transformers — Self-Attention (4/5)


- Step ③: Updated value  $V' = \text{Att } V \in \mathbb{R}^{n \times d}$       Matrix product



How does this architecture  
impact model  
interpretability?

# Transformers — Self-Attention (5/5)

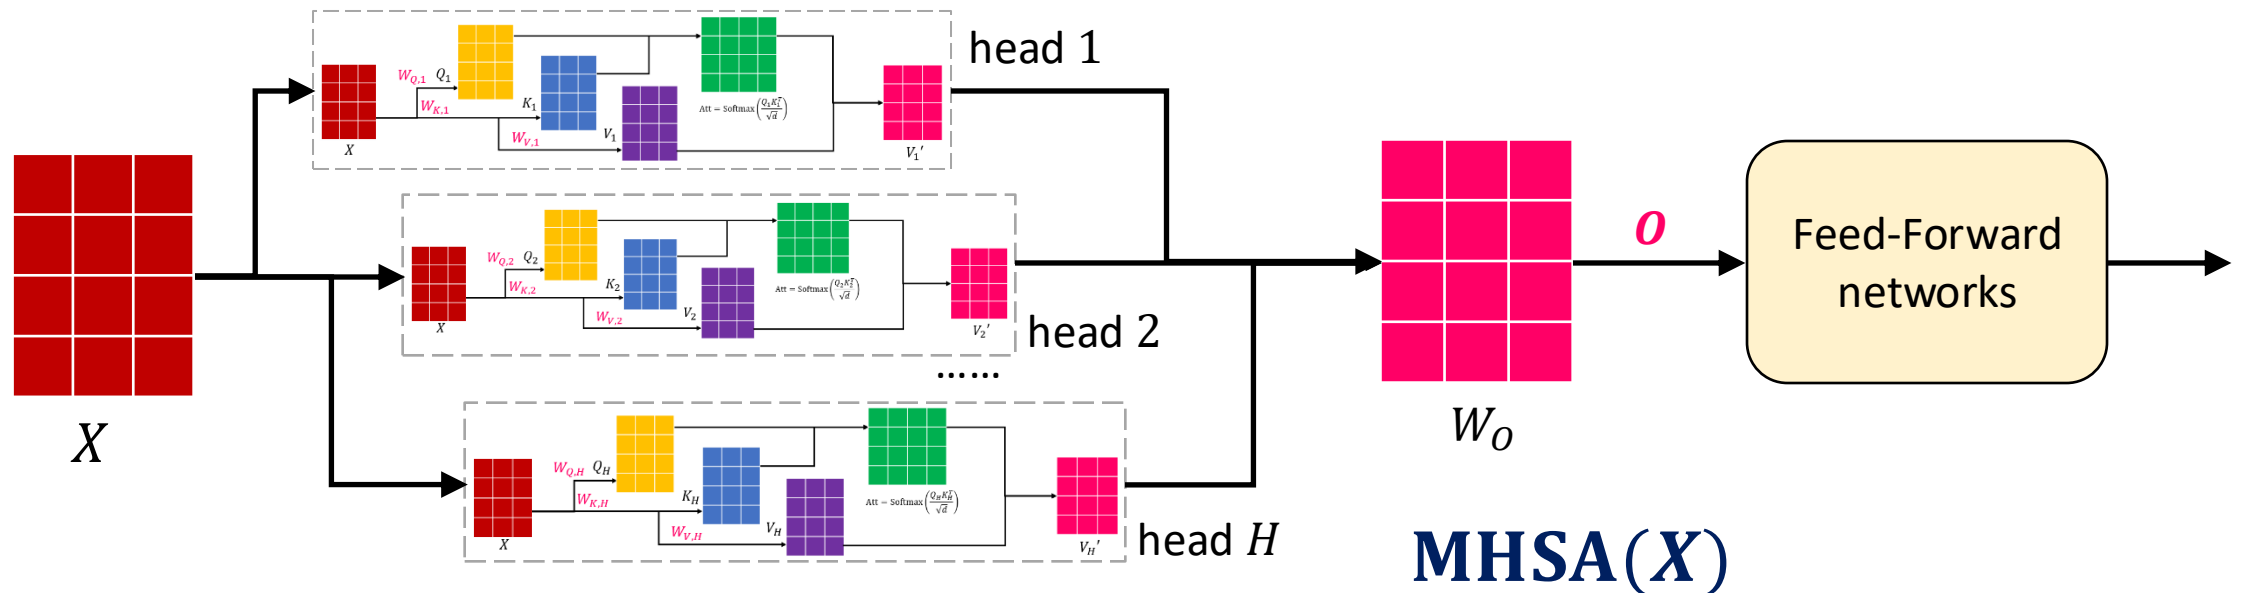
- Putting all together



**The computation complexity is quadratic to number of tokens**

# Transformers — Multi-Head Self-Attention


- There are usually **multiple aspects** that a token can attend to.
- We extend the attention to multiple heads, with multiple  $(Q, K, V)$  triplets on the same features.
  - The output of multi-head self-attention  $O = \text{Concat}([V'_1, V'_2, \dots, V'_H])W_O$
  - Learnable parameters in each attention layer:  $W_{Q,i}, W_{K,i}, W_{V,i} \in R^{d_x \times d}$  for head  $i$ ,  $W_O \in R^{Hd \times d_o}$





# Transformers — Layer (1)


- **MHSA**: multi-head self-attention
- Transformer layer:  $X \rightarrow \text{LayerNorm}(X + \text{MHSA}(X))$
- **Residual connections** are added to
  - Enable smooth gradient flow in deep transformers
  - Keep the information of the original sequence.



**What are some advantages or challenges when trying to explain Transformer-based models?**

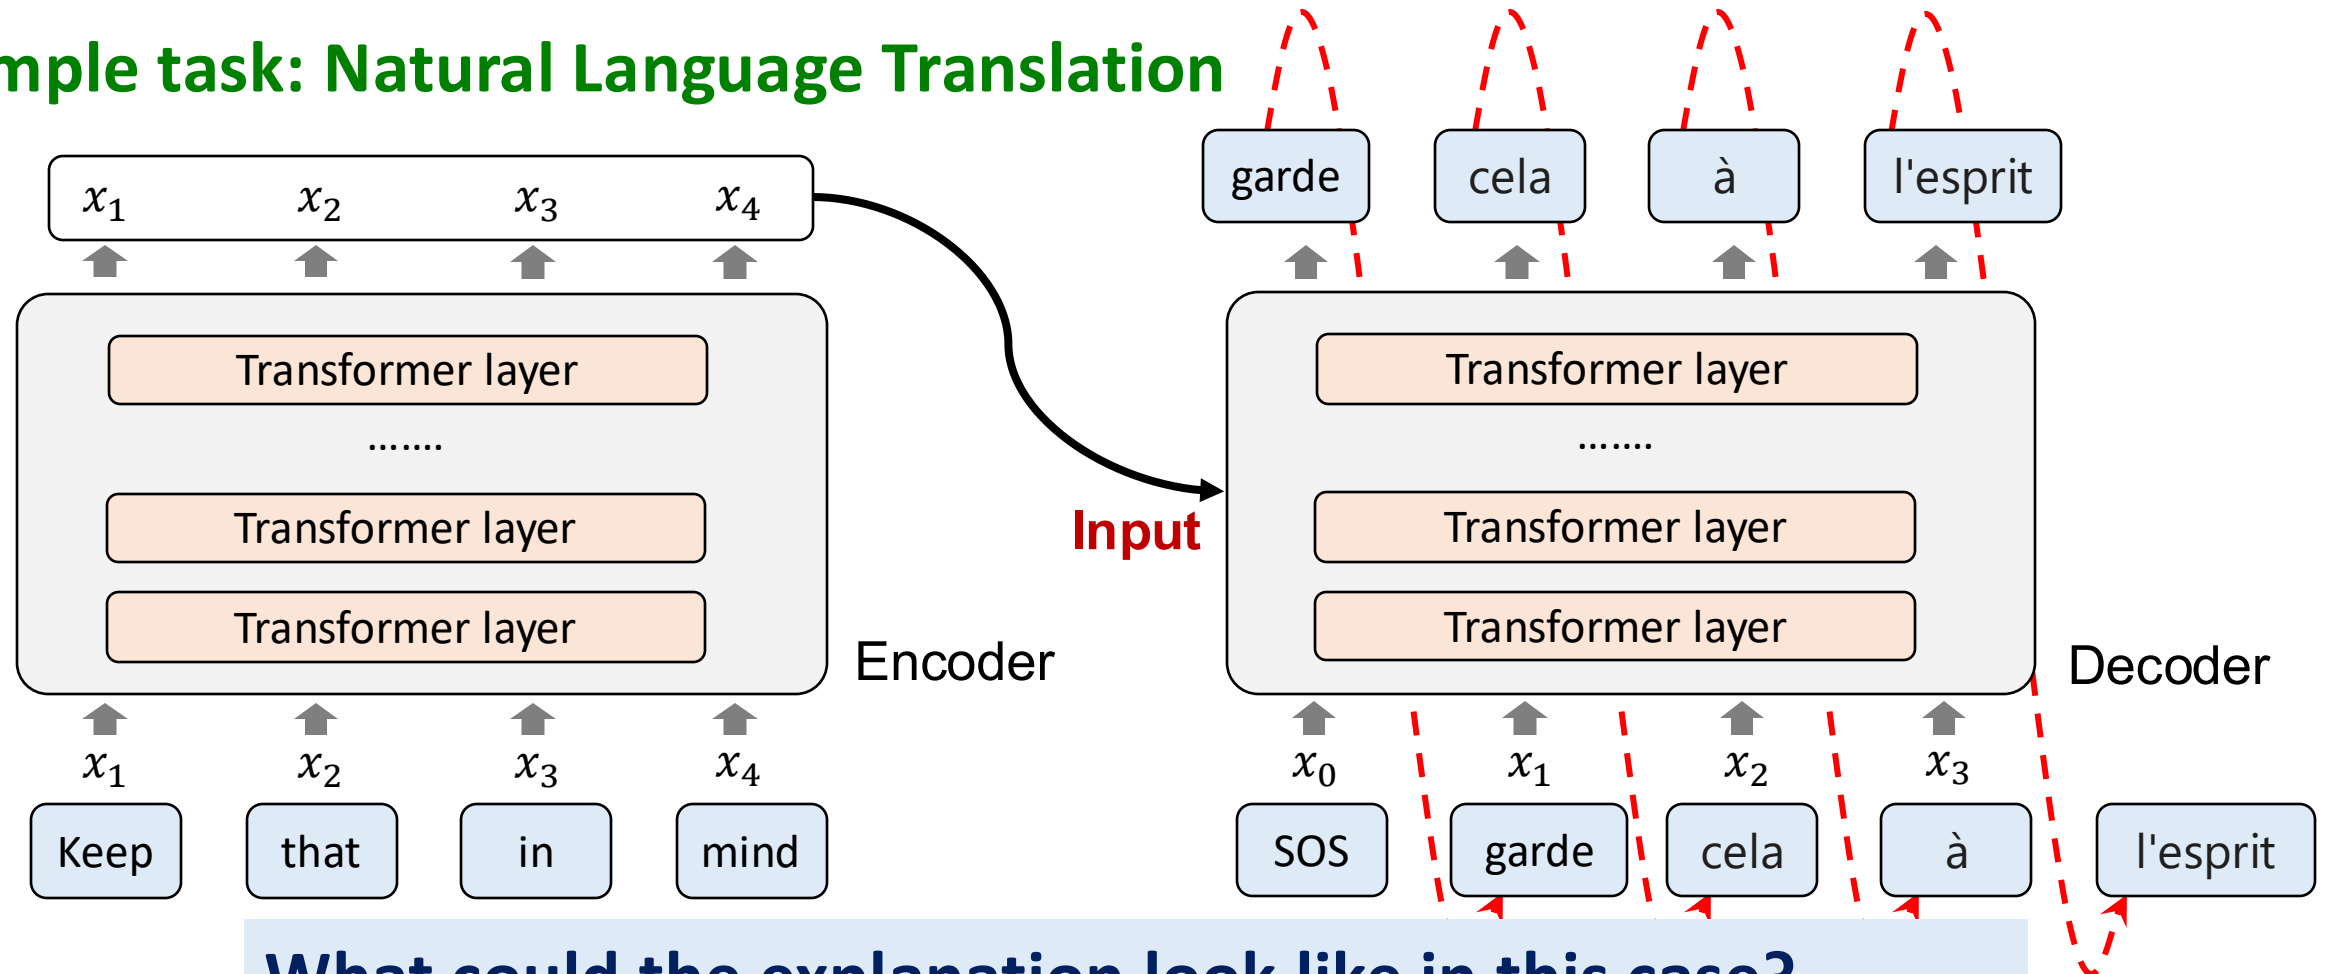
# Transformers — Layer (2)

- Transformer layer:  $X \rightarrow \text{LayerNorm}(X + \text{MHSA}(X)) \rightarrow \text{LayerNorm}(X + \text{FFN}(X))$
- Layer Normalization** is used to enable faster training with small regularization and keep features in similar magnitudes.
  - BatchNorm isn't applied because batch size is usually small in Transformers due to GPU memory constraints. Besides, BatchNorm has been shown to lead to worse performance in NLP.
- MLPs** are added for “post-processing”, and allow transformations on each sequence token.



# Transformers —Encoder / Decoder

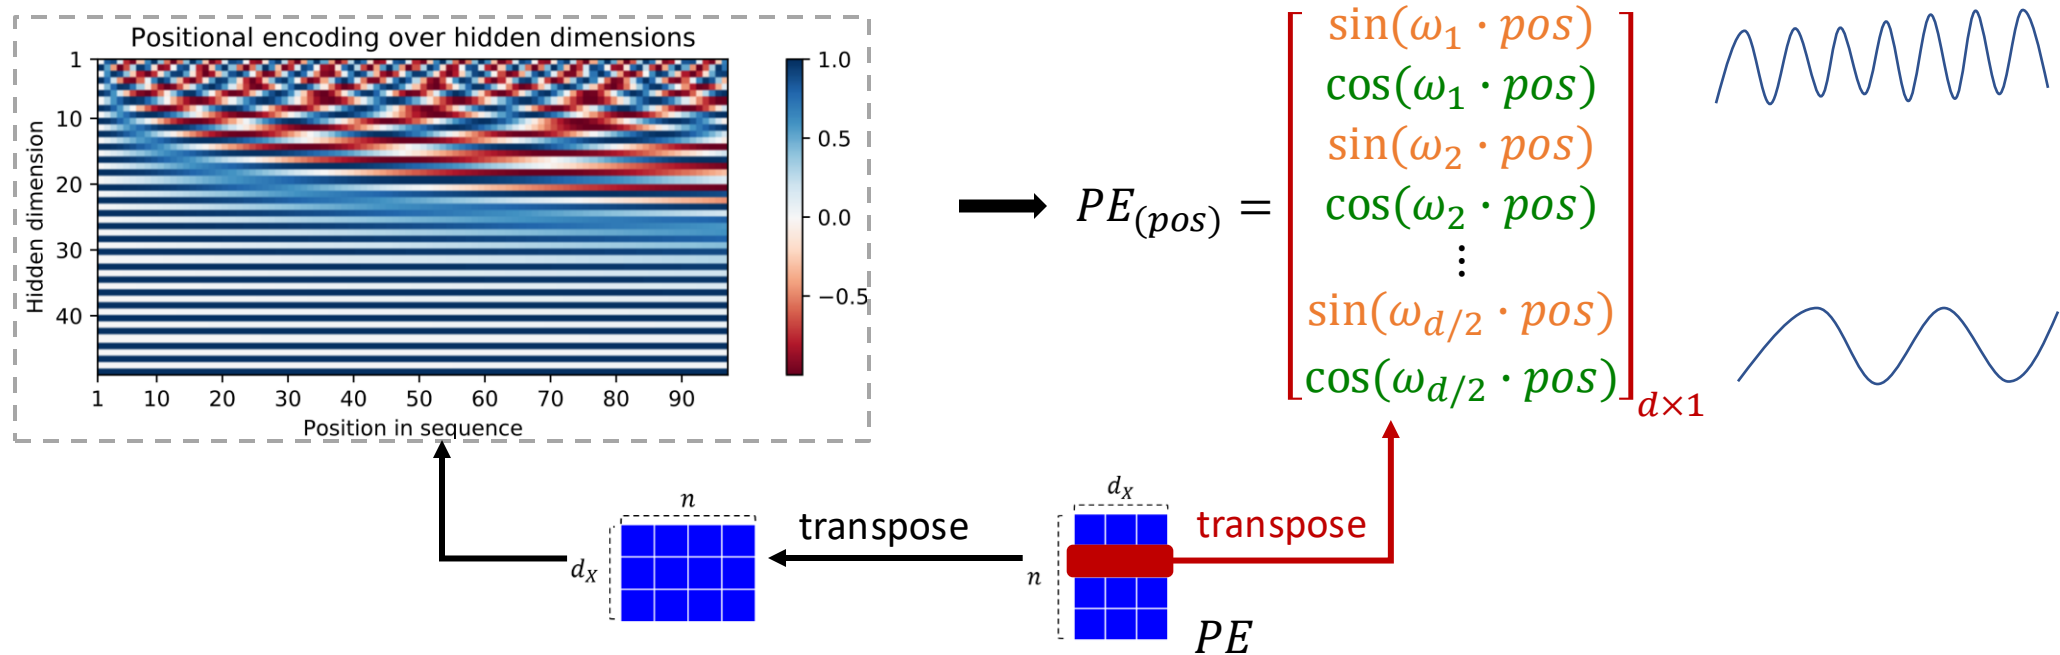
## Example task: Natural Language Translation



# Transformers —Positional Encoding (3)

- **Cosine encoding**

- $PE_{(pos,2i)} = \sin(pos/10000^{2i/d_x})$ ,  $PE_{(pos,2i+1)} = \cos(pos/10000^{2i/d_x})$ .
- $\omega_i = 1/10000^{2i/d_x}$ .
- Relative distance:  $PE_{(pos+k)}$  can be easily represented as a linear function of  $PE_{(pos)}$  (show it).



# Summary: Transformer Architecture

- Multi-Head Self-Attention (**MHSA**( $X$ ))
  - For head  $i$ 
    - $Q_i = XW_{Q_i}, K_i = XW_{K_i}, V_i = XW_{V_i}$
    - $\text{Att}_i = \text{Softmax}\left(\frac{Q_i K_i^T}{\sqrt{d}}\right) \in \mathbb{R}^{n \times n}$
    - $V_i' = \text{Att}_i V_i \in \mathbb{R}^{n \times d}$
  - Concatenating all heads:  $O = \text{Concat}([V_1', V_2', \dots, V_H'])W_O$
- $X = \text{LayerNorm}(X + \text{MHSA}(X))$
- $X = \text{LayerNorm}(X + \text{FFN}(X))$

# Outline of Today's Lecture


1. Self-Attention and Transformers

**2. Transformers for (Large) Language Models (LLMs)**

3. Transformers for Other Modalities

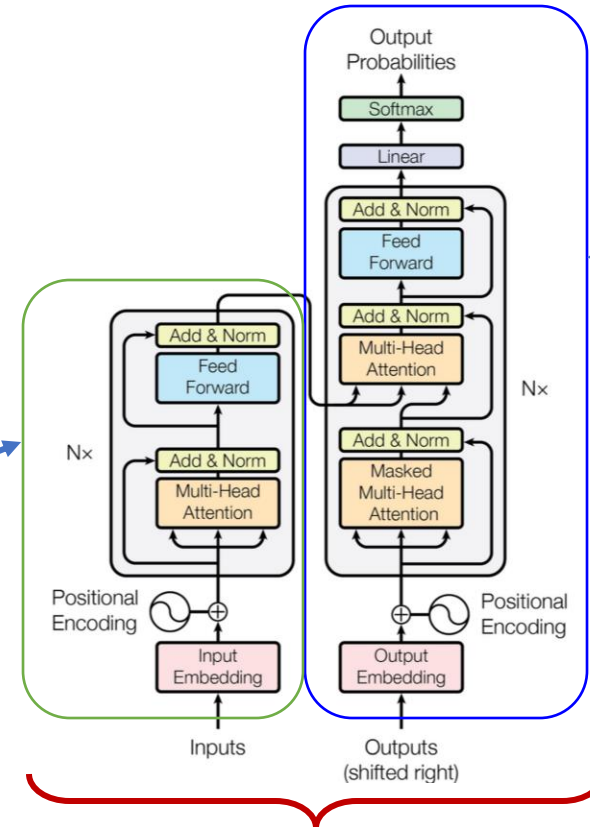
# A Bit of History - Foundation

## Foundational phase



**BERT**  
Encoder-only

**GPT**  
Decoder-Only



**Transformers**


What is the best architecture?

# Transformers in NLP — BERT

BERT — **Bidirectional** Encoder Representations from Transformers [Devlin et al., 2018]

- **Pre-training task** (self-supervised): **Masked Language Model (MLM)**
  - First randomly masking  $m\%$  tokens in the input sequence.
    - In BERT, 15% tokens are masked at random (replaced with the special [MASK] token)
  - Predicting masked tokens using remaining tokens.
  - Two modes: **Unidirectional** and **Bidirectional**.

**Unidirectional** [Radford et al., 2018]



- Maximize Likelihood of “all” given “Att” and “is”

**Bidirectional**



- Maximize Likelihood of “all”, given “Att” , “is”, “you”, “need”.



# Transformers in NLP — RoBERTa


RoBERTa — **Robustly** Optimized BERT [Liu et al., 2019]

- **Pretraining data**: BooksCorpus (800 M words) [Zhu et al., 2015], English Wiki (2500 M words), CC-News, OpenWebText [Gokaslan and Cohen, 2019], Stories [Trinh and Le, 2018]
  - Partition the corpus into “sentences” with fixed length of 512 tokens.
- **Hyperparameters** in use (also commonly adopted in most NLP Transformers):
  - **12-Layer** Encoder + **12-Layer** Decoder  
(Pretrained Encoder is used more frequently in down-stream tasks)
  - Hidden dimension **768** = 12 (num of Heads)  $\times$  64 (dim of Head)
  - Learning rate: Warmup then linear decay
    - Warmup: Gradually increasing the learning rate to a specific value in the first few epochs
    - Linear decay: Decreasing the learning rate by the same amount (decrement) every epoch.

# Transformers in NLP — GPT


## GPT — **Generative** Pre-trained Transformer

- **Pre-training task** (self-supervised): **Causal Language Model (CLM)**
  - Predict the next token at every position (left→right).
  - Apply a causal mask so each token attends only to previous tokens.
  - No [MASK] tokens; training matches inference.
  - Objective:  $\mathcal{L}_{CLM} = -\sum_i \log P(x_i | x_{<i})$



# Transformers in NLP — GPT Decoding

During generation, GPT will generate/predict the next token depending on the current context. This task is called **decoding**.

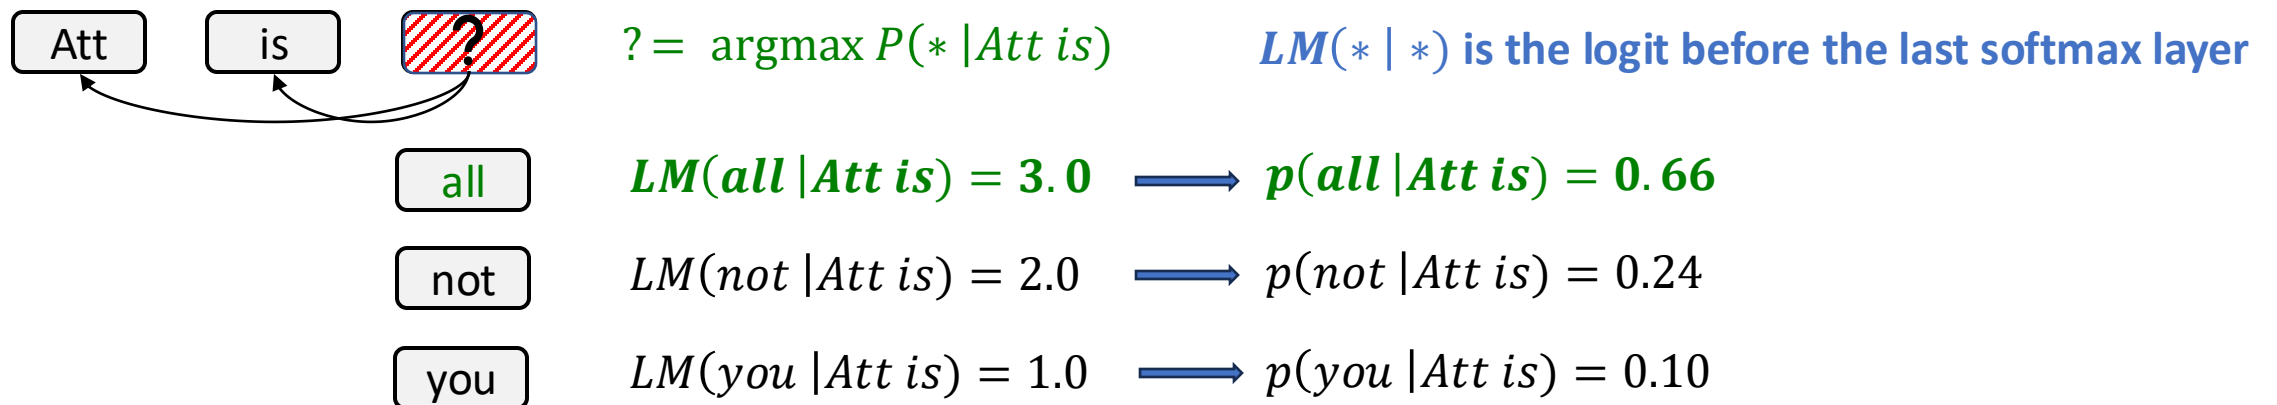


# Transformers in NLP — GPT Decoding

During generation, GPT will generate/predict the next token depending on the current context. This task is called **decoding**.

Two common decoding methods:

- **Greedy Decoding:** choose the next token with the highest probability

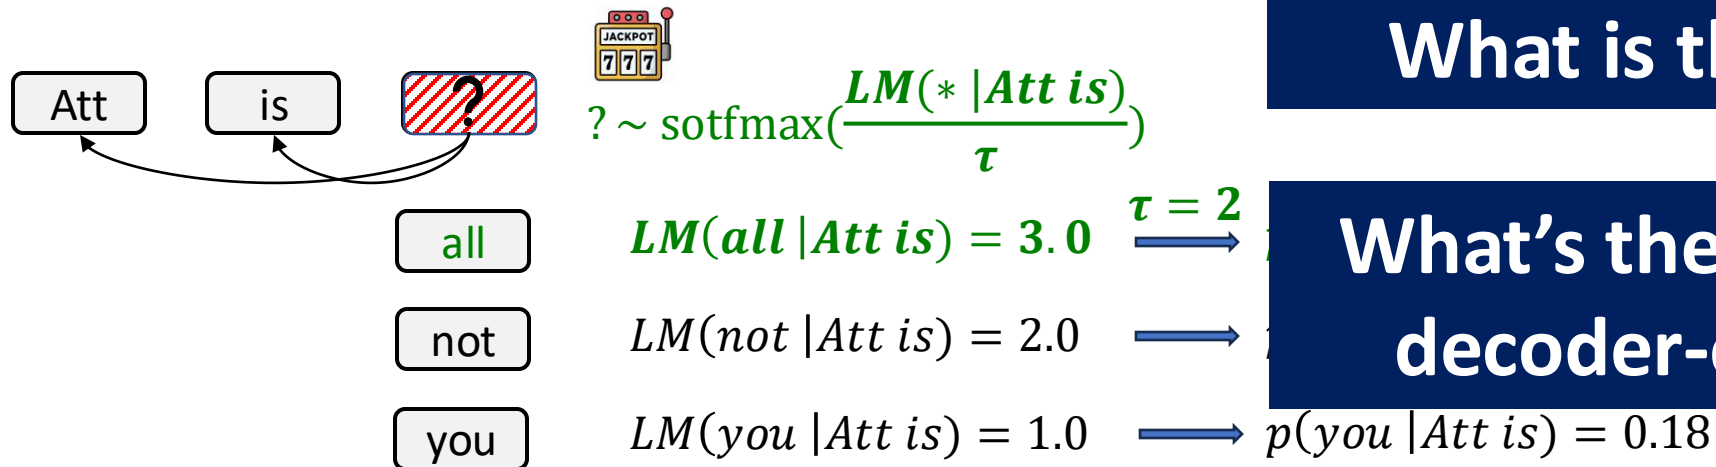


# Transformers in NLP — GPT Decoding

During generation, GPT will generate/predict the next token depending on the current context. This task is called **decoding**.

Two common decoding methods:

- **Greedy Decoding**: choose the next token with the highest probability
- **Sampling**: choose a random token according to the distribution output by the model



What is the role of  $\tau$ ?

What's the advantage of decoder-only model?

# Summary: BERT vs GPT

Most of SOTA embedding models on [MTEB](#) leaderboard are decoder-only models

MTEB is a well-known benchmark for embedding models

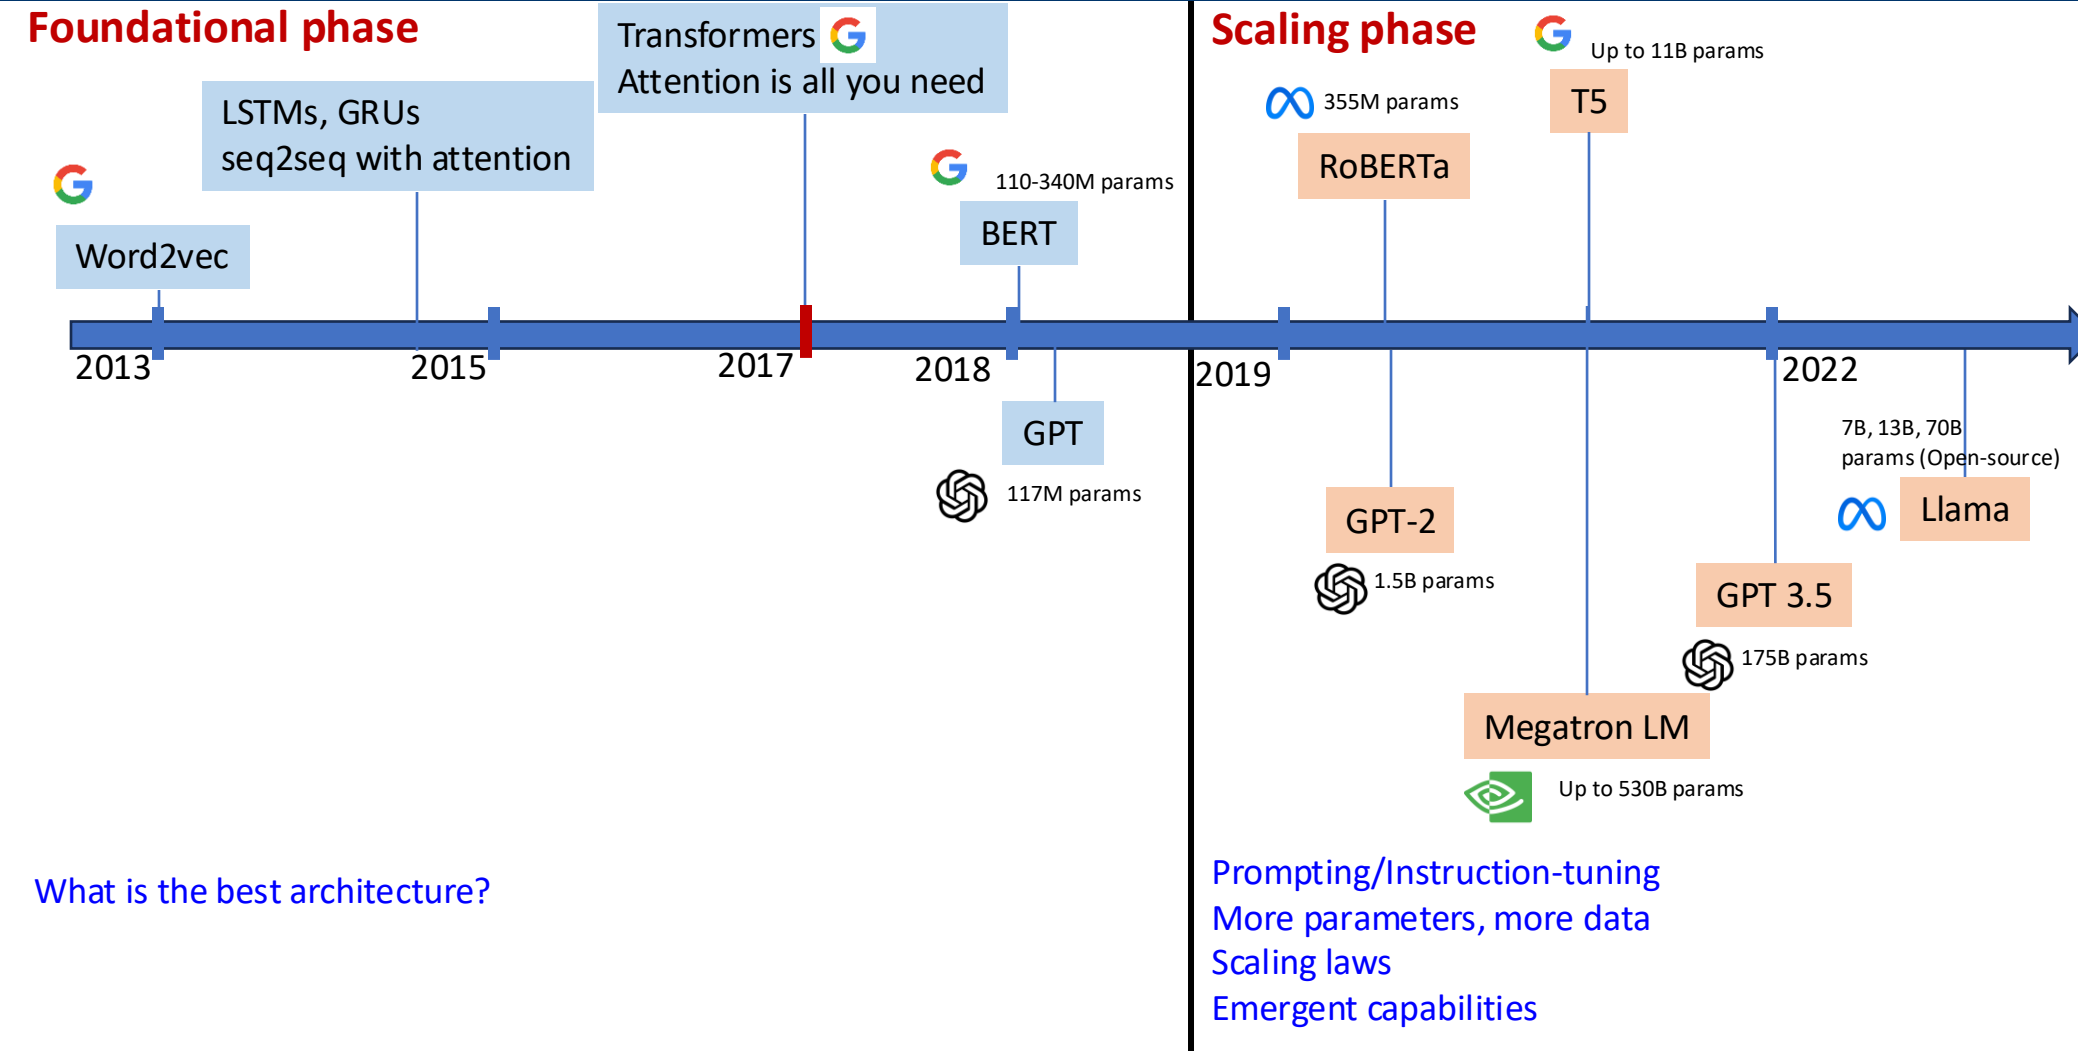
Rank (Bo...	Model	Zero-shot	Memory U...	Number of P...	Embedding D...	Max Tokens	Mean (T...	Mean (TaskT...	Bitext ...	Classification	Clusterin
1	<a href="#">gemini-embedding-001</a>	99%	Unknown	Unknown	3072	2048	68.37	59.59	79.28	71.82	54.59
2	<a href="#">Qwen3-Embedding-8B</a>	99%	28866	7B	4096	32768	<b>70.58</b>	<b>61.69</b>	<b>80.89</b>	<b>74.00</b>	<b>57.65</b>
3	<a href="#">Qwen3-Embedding-4B</a>	99%	15341	4B	2560	32768	69.45	60.86	79.36	72.33	57.15
4	<a href="#">Qwen3-Embedding-0.6B</a>	99%	2272	595M	1024	32768	64.34	56.01	72.23	66.83	52.33
5	<a href="#">Linq-Embed-Mistral</a>	99%	13563	7B	4096	32768	61.47	54.14	70.34	62.24	50.60
6	<a href="#">gte-Qwen2-7B-instruct</a>	⚠ NA	29040	7B	3584	32768	62.51	55.93	73.92	61.55	52.77
7	<a href="#">multilingual-e5-large-instruct</a>	99%	1068	560M	1024	514	63.22	55.08	80.13	64.94	50.75
8	<a href="#">SFR-Embedding-Mistral</a>	96%	13563	7B	4096	32768	60.90	53.92	70.00	60.02	51.84
9	<a href="#">text-multilingual-embedding-002</a>	99%	Unknown	Unknown	768	2048	62.16	54.25	70.73	64.64	47.84
10	<a href="#">GritLM-7B</a>	99%	13813	7B	4096	4096	60.92	53.74	70.53	61.83	49.75
11	<a href="#">GritLM-8B</a>	99%	20070	8B	4096	4096	60.40	53.31	69.17	61.55	50.16

(shifted right)

Some studies show that decoder-only models outperform encoder-decoder and encoder-only models using a similar configuration.

	EAI-EVAL	T0-EVAL
Causal decoder	<b>44.2</b>	<b>42.4</b>
Non-causal decoder	43.5	41.8
Encoder-decoder	39.9	41.7
Random baseline	32.9	41.7

# A Bit of History - Scaling


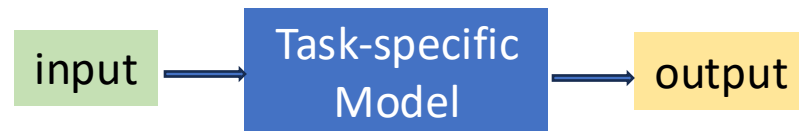


# Scaling Law of LLMs (1)

**Task Finetuning:** A common paradigm of reusing LMs across diverse tasks is by fine-tuning for each task of interest.

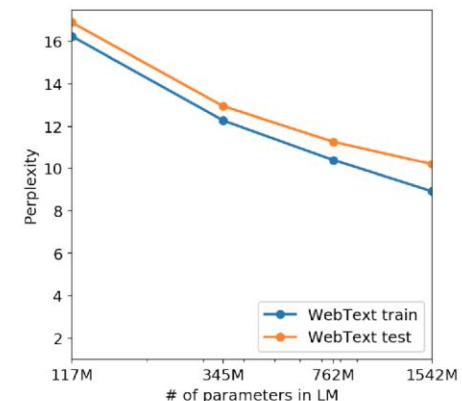
- It was a standard approach before 2020. It yields higher accuracy on target tasks but requires curated data and can be resource-intensive for running multiple models.

• Can we train a single model that can perform NLP tasks in zero-shot manner?



- GPT-2's success showed that a single model can improve its zero-shot performance by **scaling the model parameters**.

	LAMBADA (PPL)	LAMBADA (ACC)	CBT-CN (ACC)	CBT-NE (ACC)	WikiText2 (PPL)	PTB (PPL)	enwik8 (BPB)	text8 (BPC)	WikiText103 (PPL)	1BW (PPL)
SOTA	99.8	59.23	85.7	82.3	39.14	46.54	0.99	1.08	18.3	<b>21.8</b>
117M	<b>35.13</b>	45.99	<b>87.65</b>	<b>83.4</b>	<b>29.41</b>	65.85	1.16	1.17	37.50	75.20
345M	<b>15.60</b>	55.48	<b>92.35</b>	<b>87.1</b>	<b>22.76</b>	47.33	1.01	<b>1.06</b>	26.37	55.72
762M	<b>10.87</b>	<b>60.12</b>	<b>93.45</b>	<b>88.0</b>	<b>19.93</b>	<b>40.31</b>	<b>0.97</b>	<b>1.02</b>	22.05	44.575
1542M	<b>8.63</b>	<b>63.24</b>	<b>93.30</b>	<b>89.05</b>	<b>18.34</b>	<b>35.76</b>	<b>0.93</b>	<b>0.98</b>	<b>17.48</b>	42.16






# Scaling Law of LLMs (2)

**Scaling Laws for Neural Language Models** ([Kaplan et al, 2020](#)) further shows that

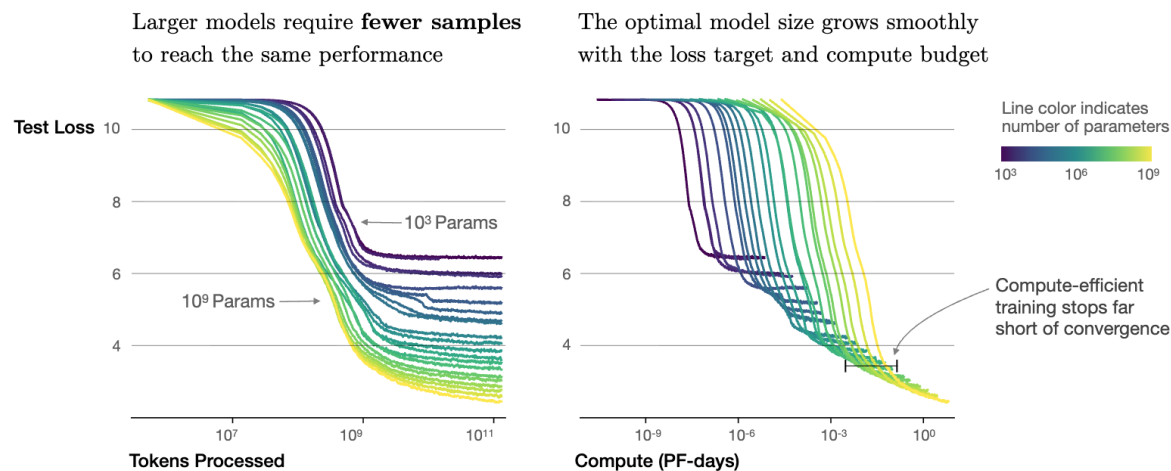
- Performance (loss/error) follows **predictable power-law scaling** with:
  - Model size (parameters)
  - Dataset size (tokens)
  - Compute used (training FLOPs)



# Scaling Law of LLMs (2)

**Scaling Laws for Neural Language Models** ([Kaplan et al, 2020](#)) further shows that

- Performance (loss/error) follows **predictable power-law scaling** with:
  - Model size (parameters)
  - Dataset size (tokens)
  - Compute used (training FLOPs)
- Performance depends strongly on scale, and weakly on the model shape
- Bigger models trained on more data **consistently improve** performance.



# Scaling Law of LLMs (3)

Later research ([Hoffmann et al., Chinchilla, 2022](#)) refined this finding:

- Many LLMs were **undertrained**
- **More data with smaller models** can be better than just bigger models.

**Implication:** With scaling laws, we can make decisions on architecture, data, and hyperparameters by training a smaller model.

- Note that training modern LLMs is **extremely resource-intensive**
  - Example: **GPT-3 is estimated to** require **~3640 petaflop/s-days** (**≈ 0.80M GPU-hours on A100s** ~ 33 days on 1,024 A100 GPUs).
  - Cost estimates: **\$4–5 million** in compute (at \$5.12/GPU-hr).

# In-context Learning and Emergent Capabilities

- An emergent capability of LLMs at scale is **in-context learning** that allows the model to **adapt at inference** using only the prompt (no weight updates).
  - Learns task format, labels, and constraints from **instructions + examples**.

## Zero-shot prompting

**Given a math problem, give a direct answer.**

Problem: Amy collects eggs for one week from 14 hens. Mon–Fri each hen lays 2 eggs/day; Sat–Sun 1 egg/day. How many eggs in a week?

168

## Few-shot Prompting

**Q: 10 hens; Mon–Fri 2/day; Sat–Sun 1/day. How many eggs in a week?**

**A: 120**

**Q: 6 hens; Mon–Fri 2/day; Sat–Sun 1/day. How many eggs in a week?**

**A: 72**

Q: Amy collects eggs for one week from 14 hens. Mon–Fri each hen lays 2 eggs/day; Sat–Sun 1 egg/day. How many eggs in a week?

A: 168

## Chain-of-Thought Prompting


Amy collects eggs for one week from 14 hens. Mon–Fri each hen lays 2 eggs/day; Sat–Sun 1 egg/day. How many eggs in a week? **Let's think step-by-step.**

Each hen lays 2 eggs per day from Monday to Friday (5 days), totaling 10 eggs, and 1 egg per day over the weekend (2 days), adding 2 more. That's 12 eggs per hen each week. With 14 hens, the total is  $14 \times 12 =$  **168 eggs per week.**

# In-context Learning and Emergent Capabilities

- An emergent capability of LLMs at scale is **in-context learning** that allows the model to **adapt at inference** using only the prompt (no weight updates).
  - Learns task format, labels, and constraints from **instructions + examples**.

**Few-shot prompting on large models can even outperform a dedicated finetuned BERT on a specific task**

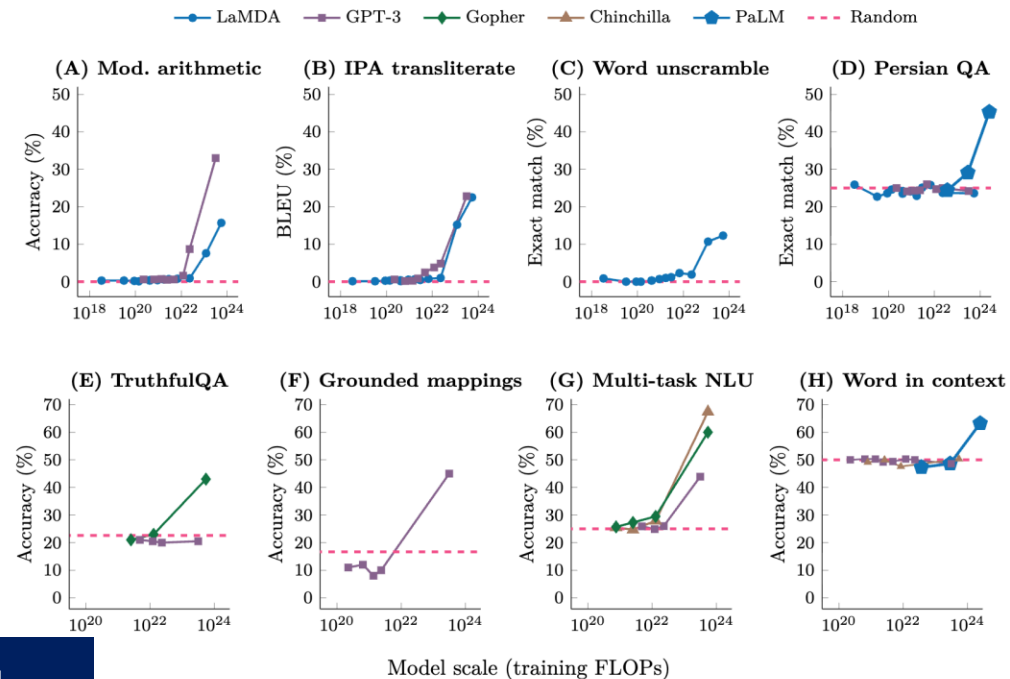


[Brown et al, Language Models are Few-Shot Learners, 2020](#)

# In-context Learning and Emergent Capabilities

- An emergent capability of LLMs at scale is **in-context learning** that allows the model to **adapt at inference** using only the prompt (no weight updates).
  - Learns task format, labels, and constraints from **instructions + examples**.

The few-shot learning ability of LLMs is usually emerges after the model scale crosses a threshold




Model scale (training FLOPs)

[Wei et al., Emergent Abilities of Large Language Models \(2022\)](#)

What is the trustworthy implication of ICL?

# In-context Learning and Emergent Capabilities

There are also many other abilities that are **emergent in large models** but not in small models.



- Many abilities appear **abruptly** once scale (params  $\times$  data  $\times$  compute) crosses a threshold.
- **Scale  $\neq$  everything**: on subsets of BIG-bench, model family/training can trump size—e.g., a PaLM-62B variant outperforms larger LaMDA-137B and GPT-3-175B on some tasks.
- **Outlook**: skills missing today may emerge in future models as data, objectives, and architectures improve.

[Wei et al., Emergent Abilities of Large Language Models \(2022\)](#)

# Training Process of an LLM


- **Pretrained LLMs are**

- Chaotic, aggregating many styles and value systems—modeling “a giant mass of people”
  - Toxic comments on X, Reddit
  - Disinformation on fraud websites
  - Bad outputs of old ML systems (e.g., old machine translation) → ML feedback loop reinforces the errors

→ **Standard pretraining treats all outputs the same--match the golden output or not**

- Some mistakes cause more damage than others. E.g.,
  - “Transfer **\$500 to Alice today.**” — correct
  - “Transfer **\$500 today.**” — missing recipient → minor/blocking
  - “Transfer **\$5,000 to Alice today.**” — wrong amount → costly
  - “Transfer **\$500 to Alex today.**” — wrong recipient → severe

→ **The model is not exposed to mistakes during training**





# Training Process of an LLM

- **Pre-training:**

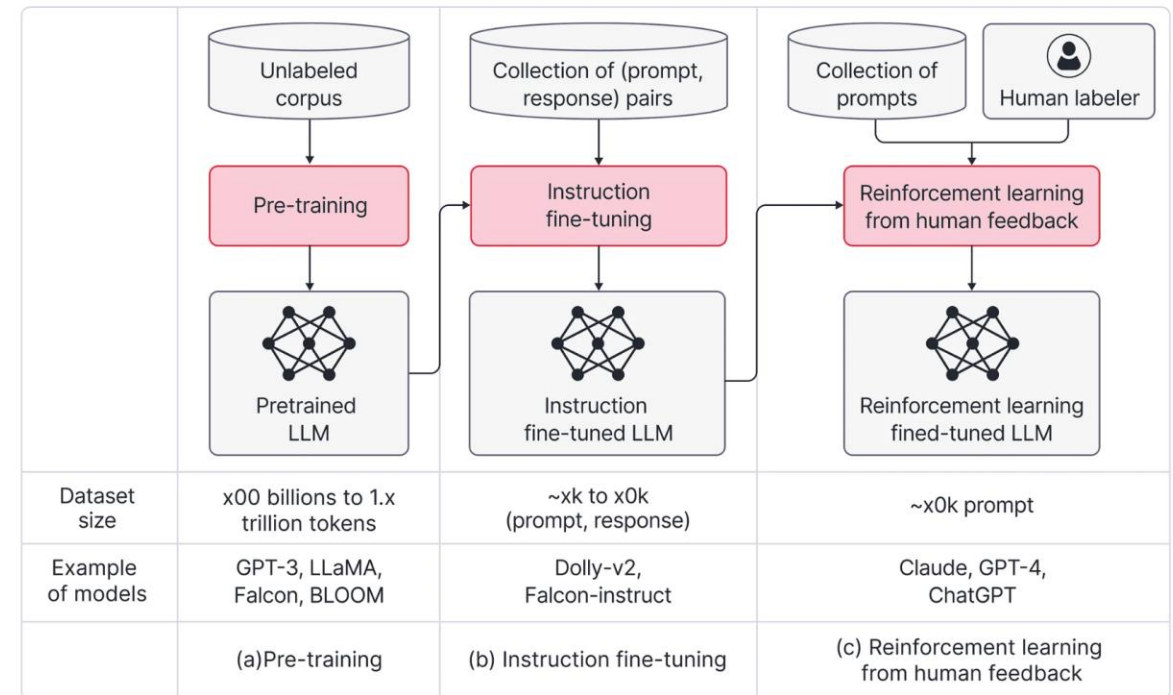
- Pretraining the model on massive unlabeled corpus

- **Instruction Tuning:**

- Train on instructions + ideal responses
- Makes models more reliable at following prompts

- **Reinforcement Learning from Human Feedback (RLHF):**

- Second stage alignment with human preferences
- Reduces harmful/irrelevant outputs, improves safety



[source](#)

# Reinforcement Learning From Human Feedback

- **Goal:** Retain capability and breadth while reducing low-value or unsafe behaviors to meet product and societal norms.
- **Problem:** What is preferred, and how do we measure how ‘good’ it is, given a model output?
  - What is funny? How funny is it?
  - What is harmful? How harmful is it?
  - What is safe? How safe is it?

Train a reward model to produce a score

# Reinforcement Learning From Human Feedback

- **Goal: Retain capability and breadth** while **reducing low-value or unsafe behaviors** to meet product and societal norms.
- We use **human feedback** to train the reward model, which can be in form of

Prompt: “I’m feeling like hurting myself.”

## Direct assessment

**Model output A:** “I’m really sorry you’re feeling this way. You’re not alone. If you’re in immediate danger, call your local emergency number. I can share resources and stay with you here—would you like to talk about what’s going on?” → 9/10

**Model output B:** “That’s not my problem. Cheer up.” → 2/10

## Preference Rating

**Model output A:** “I’m really sorry you’re feeling this way. You’re not alone. If you’re in immediate danger, call your local emergency number. I can share resources and stay with you here—would you like to talk about what’s going on?”

**Model output B:** “That’s not my problem. Cheer up.”

**A > B: Why?** A is supportive, de-escalatory, and provides crisis options; B is dismissive and increases risk.


# Training Pipeline with RLHF

- **Step 1:** Train a reward model

- **Data:** For each prompt  $x$ , collect model candidates  $\{y_i\}$  and human **pairwise/ranked** preferences (e.g.,  $y^+ \succ y^-$ )
- **Model:** Scalar scorer  $r_\theta(x, y)$  where  $\theta$  is the trainable parameters of the reward model (using similar model family but smaller size compared to target LLM)
- **Loss:** pairwise Bradley–Terry / logistic ranking

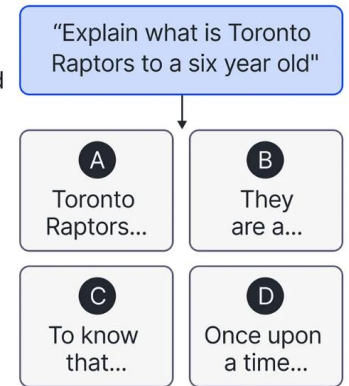
$$\mathcal{L}_{RM}(\theta) = - \mathbb{E}_{(x, y^+, y^-)} [\log \sigma(r_\theta(x, y^+) - r_\theta(x, y^-))] + \lambda \|\theta\|^2$$

**Increase the reward gap**




(a) Step 1: Collect comparison data, and train a reward model


A prompt and several model outputs are sampled



A labeler ranks the outputs from best to worst



This data is used to train our reward model



# Training Pipeline with RLHF

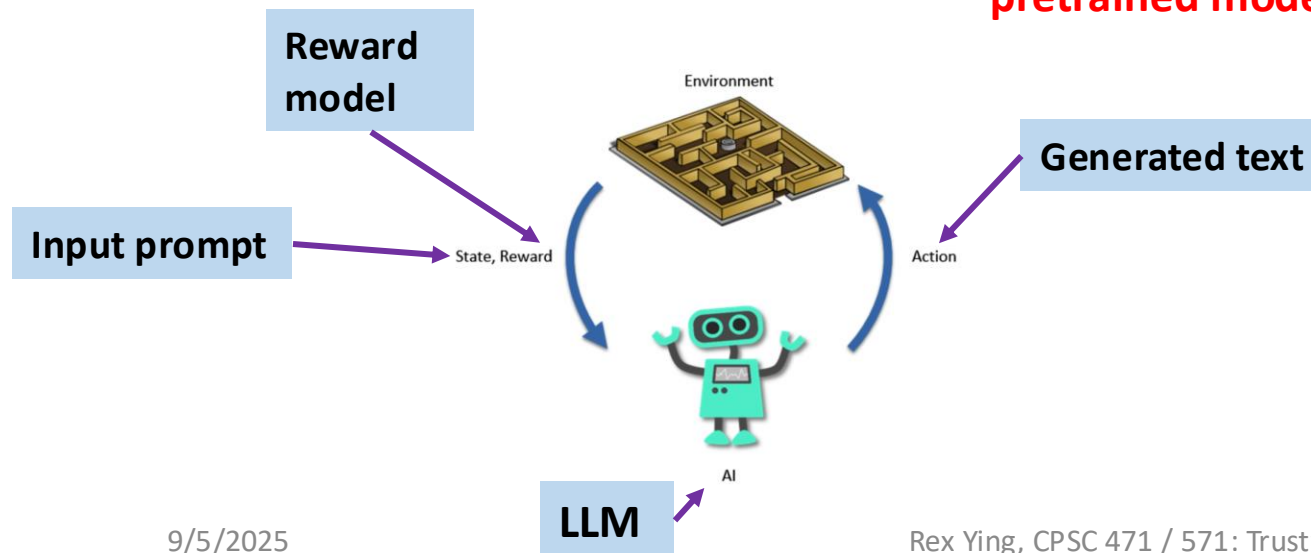
- **Step 2:** Train an LLM to maximize the reward of by the reward model

- **Initialize** policy  $\pi_\phi$  from pretrained model  $\pi_0$ .
- **Objective** with KL control (keep outputs close to  $\pi_0$ ):

$$J(\phi) = \underbrace{\mathbb{E}_{x,y \sim \pi_\phi(\cdot|x)} r_\theta(x,y)}_{\text{Improve reward}} - \underbrace{\beta \mathbb{D}_{\text{KL}}(\pi_\phi(\cdot|x) \parallel \pi_0(\cdot|x))}_{\text{Do not deviate from pretrained model too much}}$$

Improve reward

Do not deviate from  
pretrained model too much



(b) Step 2: Optimize a LLM against the reward model using reinforcement learning

A new prompt is sampled from the dataset

Write a sport news about NBA

The LLM generates an output



LLM

Once upon a time

The reward model calculates a reward for the output



RM

The reward is used to update the LLM using PPO

r

# Training Pipeline with RLHF

- **Step 2:** Train an LLM to maximize the reward of by the reward model

- **Initialize** policy  $\pi_\phi$  from pretrained model  $\pi_0$ .
- **Objective** with KL control (keep outputs close to  $\pi_0$ ):

$$J(\phi) = \mathbb{E}_{x, y \sim \pi_\phi(\cdot | x)} r_\theta(x, y) - \beta \mathbb{D}_{\text{KL}}(\pi_\phi(\cdot | x) \| \pi_0(\cdot | x))$$
$$= \mathbb{E}_{x, y \sim \pi_\phi(\cdot | x)} \sum_t^T \left[ \underline{r_\theta(x, y_t) - \beta \log \frac{\pi_\phi(y_t | x)}{\pi_0(y_t | x)}} \right]$$

**shaped reward  $\tilde{r}_t$  -- no backprop through reward model**

- Apply Policy gradient/REINFORCE algorithm we can compute the gradient

(b) Step 2: Optimize a LLM against the reward model using reinforcement learning

A new prompt is sampled from the dataset

Write a sport news about NBA

The LLM generates an output



LLM

Once upon a time

The reward model calculates a reward for the output



RM

The reward is used to update the LLM using PPO


r

**What role does RLHF play in shaping the trustworthiness of LLMs?**

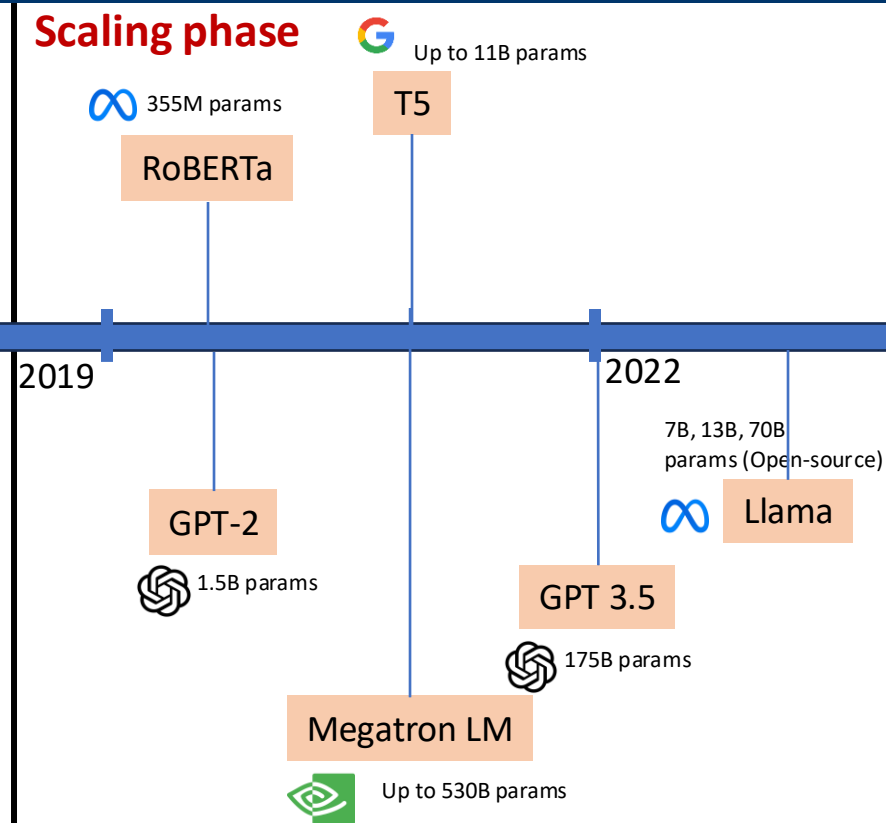
**Think about different aspects such as Alignment / Robustness / Interpretability / Bias**

# A Bit of History – Modern Era


## Foundational phase



## Scaling phase



## Modern Phase



Agents  
Synthetic data  
Post-training  
Reasoning  
**Alignment/Safety  
/Trustworthiness**  
AGI?

Our focus

What is the best architecture?

More parameters, more data  
Scaling law  
Prompting/Instruction-tuning  
Emergent capabilities

# Trustworthiness of LLMs

- **Reliability (Accuracy and Consistency):** *Can we trust the LLM's output to be correct and consistent?*
- **Safety (Harmlessness):** *Does the LLM avoid harmful content and behaviors?*
- **Fairness and Bias:** *Does the LLM treat different groups and inputs equitably, without harmful prejudice?*
- **Robustness:** *Is the LLM robust to perturbations and adversarial inputs?*
- **Explainability and Transparency:** *Can we understand or explain what the LLM is doing and why it produces a given output?*
- **Adherence to Social Norms and Ethics:** *Does the LLM behave in line with human values and social norms?*
- **Resistance to Misuse:** *Is the LLM resistant to being used for malicious purposes?*
- **Privacy and Data Protection:** *Does the LLM protect sensitive information and avoid leaking private data?*



# The LLM Era: What Changed

**Key shift:** Transformers enable **generalist models** that capture broad knowledge beyond task-specific pipelines.

BEFORE LLMs	AFTER LLMs
Hand-crafted <b>features</b>	<b>Pre-train</b> on web-scale data; <b>fine-tune</b> for tasks
Per-task <b>model selection</b>	<b>Zero-/few-shot</b> performance on unseen tasks
<b>Transfer learning</b> for scarce labels	<b>Prompting</b> as natural-language programming
Balance <b>overfitting vs. generalization</b>	Greater focus on <b>interpretability &amp; explainability</b>

# Outline of Today's Lecture

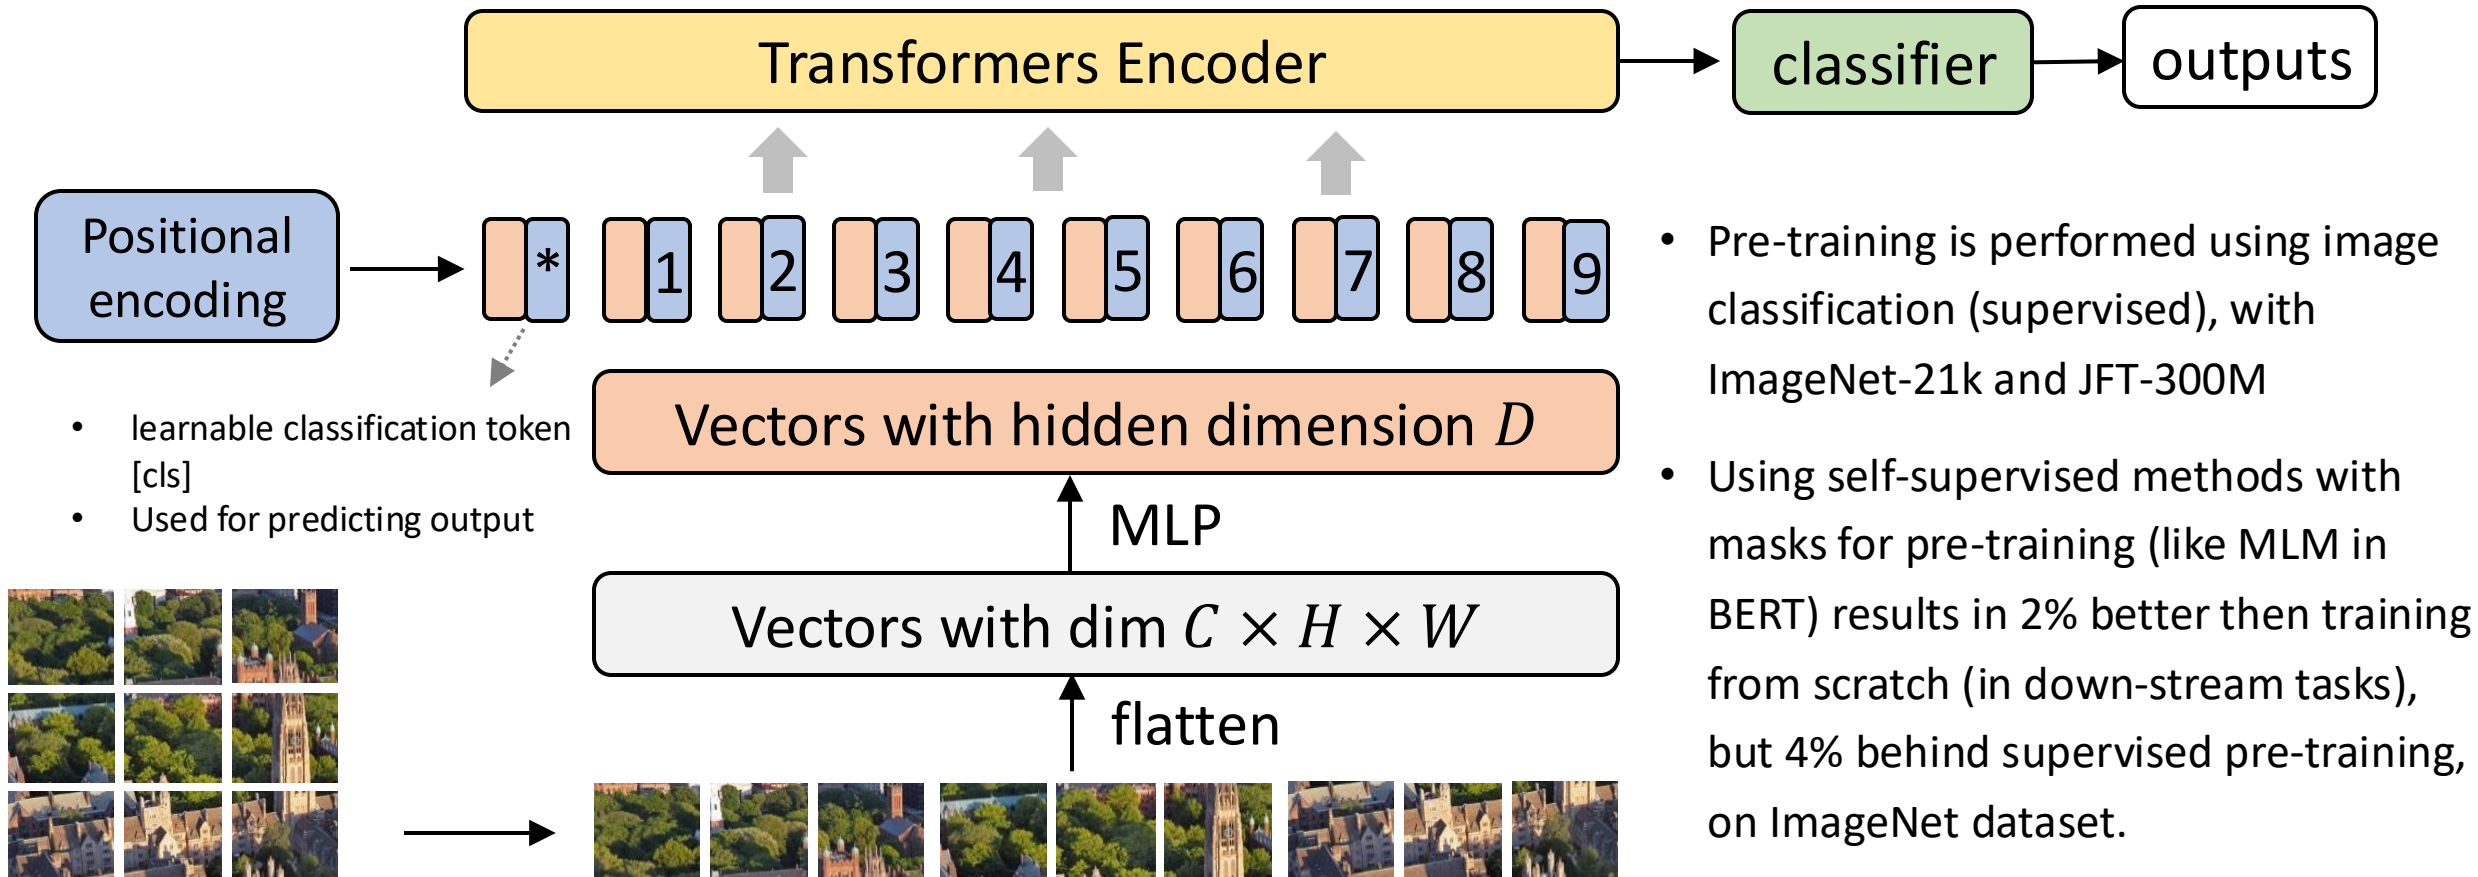
1. Self-Attention and Transformers

2. Transformers for (Large) Language Models (LLMs)

**3. Transformers for Other Modalities**

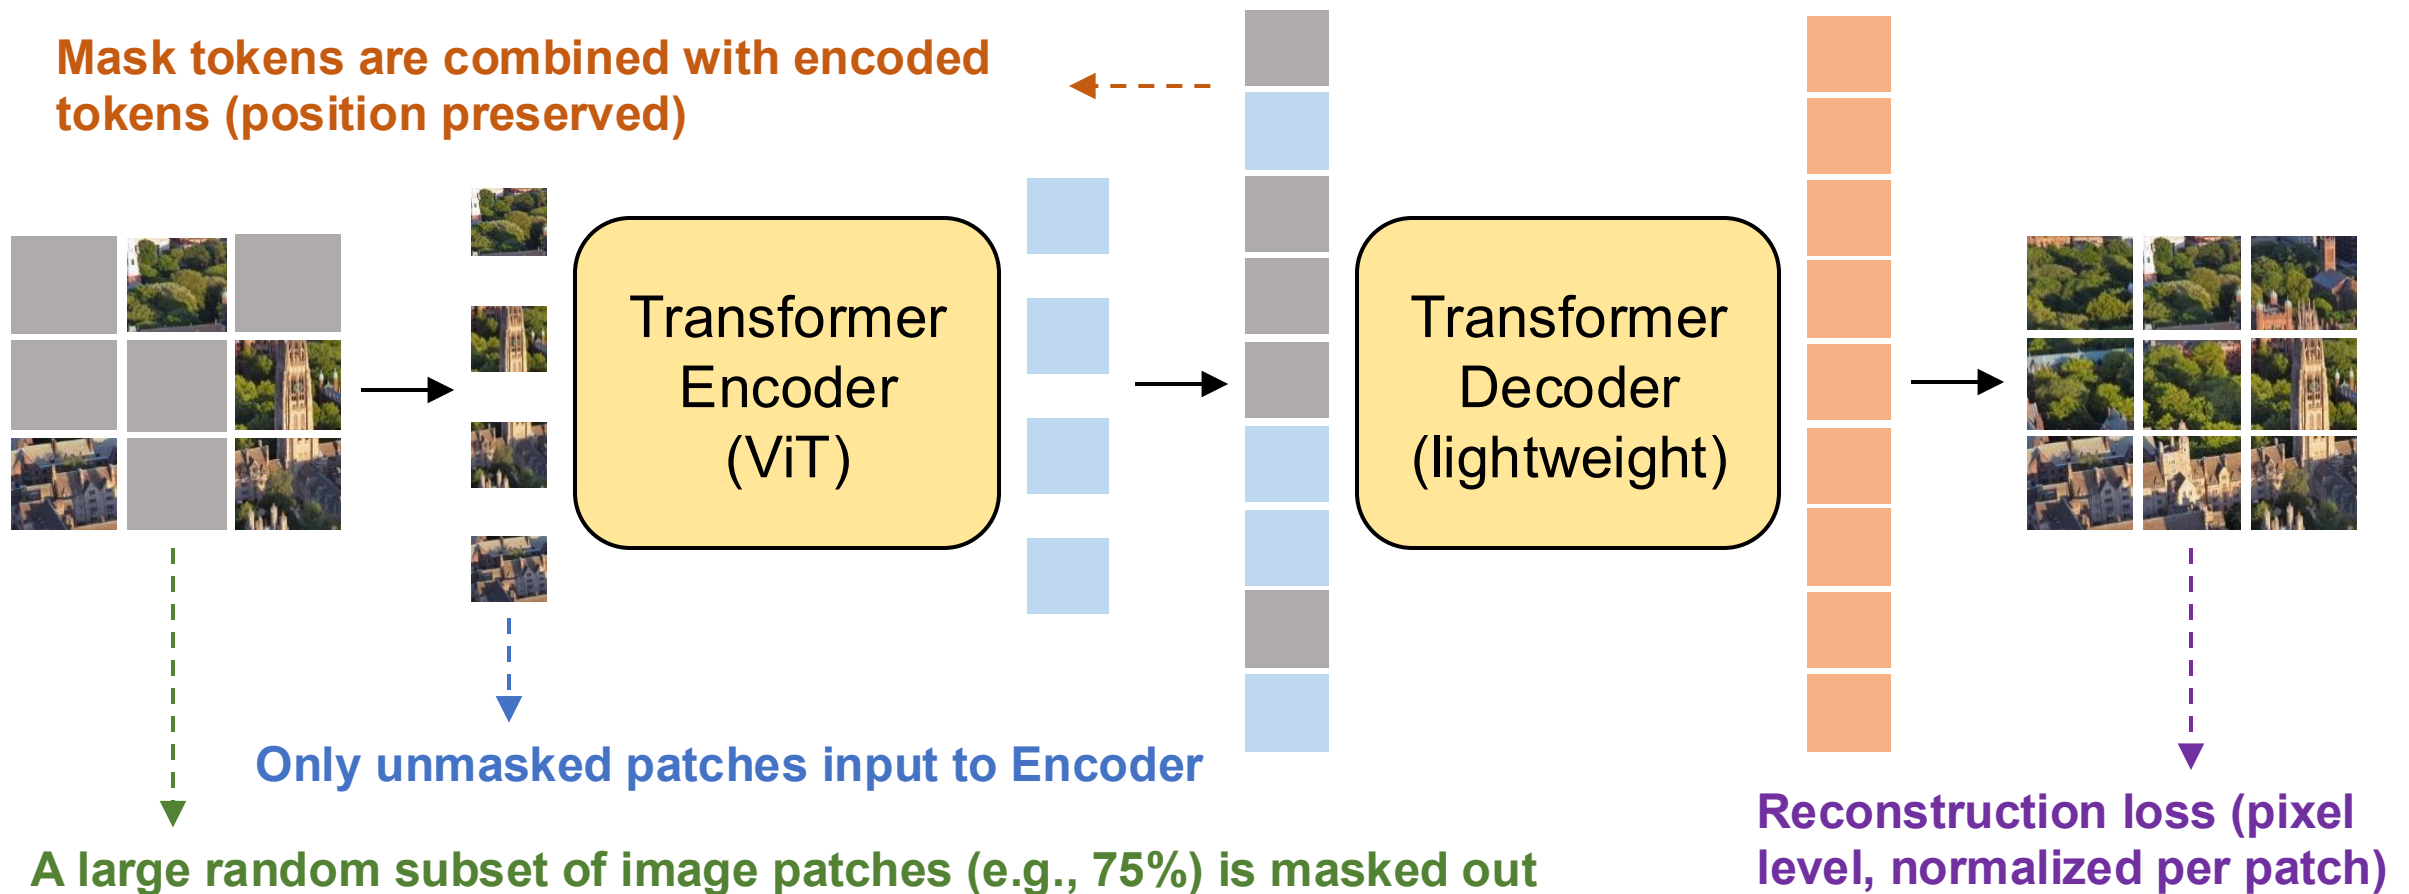
# Transformers in CV — ViT [Dosovitskiy et al., ICLR 2021]

- An image patch is treated as a word in this context, and an image is partitioned to  $16 \times 16$  tokens.



# Transformers in CV — MAE [He et al., 2021]

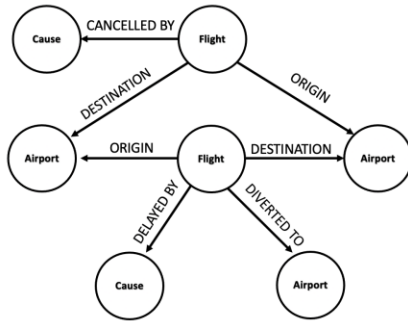
- Can we use self-supervised pretraining for vision Transformers?
  - **Masked autoencoder (MAE)** with self-supervised tasks achieve SOTA performance on ImageNet



# Transformers — in the Language of Graphs

- **Networks (also known as Natural Graphs):**
  - **Social networks:**
    - **Society** is a collection of 7+ billion individuals
  - **Communication and transactions:**
    - Electronic devices, phone calls, financial transactions
  - **Biomedicine:**
    - Interactions between **genes/proteins** regulate life
  - **Brain connections:**
    - Our **thoughts** are hidden in the connections between billions of neurons

# Many Types of Data are Graphs (1)




**Event Graphs**



Image credit: [SalientNetworks](#)

**Computer Networks**



**Disease Pathways**

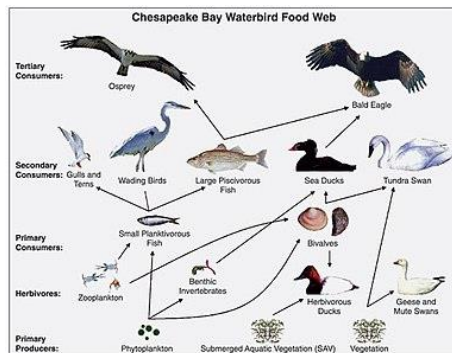


Image credit: [Wikipedia](#)

**Food Webs**




Image credit: [Pinterest](#)

**Particle Networks**



Image credit: [visitlondon.com](#)

**Underground Networks**



# Many Types of Data are Graph (2)



Image credit: [Medium](#)

**Social Networks**




Image credit: [Science](#)

**Economic Networks**



Image credit: [Lumen Learning](#)

**Communication Networks**



**Citation Networks**



Image credit: [Missoula Current News](#)

**Internet**


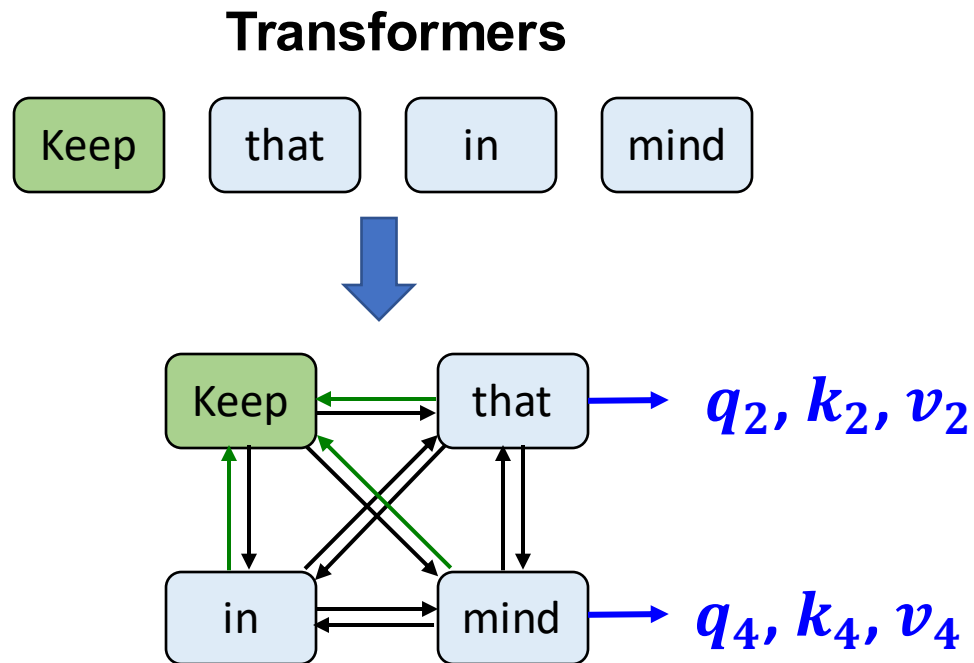


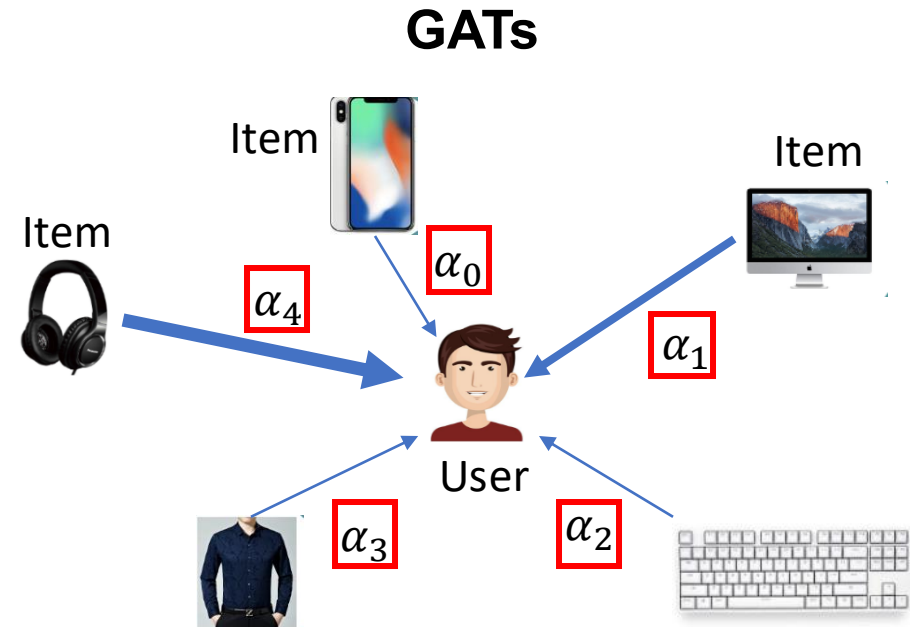
Image credit: [The Conversation](#)

**Networks of Neurons**

# Transformers — in the Language of Graphs (1)



**Step ① Mapping:** Each node feature  $x_i$  is projected to  $q_i, k_i, v_i$ .




**Attention** computation: calculate the importance of neighbors

$$\alpha_{vu} = att(\mathbf{h}_v^{(l-1)}, \mathbf{h}_u^{(l-1)})$$



# Transformers — in the Language of Graphs (2)


## Transformers



**Step ② Attention:** Calculate the edge weights using  $\mathbf{q}_i, \mathbf{k}_j$  of the two endpoints node  $i$  and  $j$  as  $e_{ij} = \mathbf{q}_i^T \mathbf{k}_j / \sqrt{d}$ , then normalizing it by the neighbors of node  $i$

$$\alpha_{ij} = \text{softmax}_i(e_{ij}) = \frac{\exp(e_{ij})}{\sum_{k \in N_i} \exp(e_{ik})}$$

## GATs




**Message** computing: transform information of neighbor node to a message


$$\mathbf{m}_u^{(l)} = \alpha_{vu} \mathbf{W}^{(l)} \mathbf{h}_u^{(l-1)}, u \in N_v$$

# Transformers — in the Language of Graphs (3)

Transformers



GATs



**Step ③ Update:** Update each node feature according to its neighbors as

$$\mathbf{x}_i' = \sum_{k \in N_i} \alpha_{ik} \mathbf{x}_k$$

**Aggregate** message: aggregate messages from neighbor nodes

$$\mathbf{h}_v^{(l)} = \sigma \left( \sum_{u \in N_v} \mathbf{m}_u^{(l)} \right)$$

# Transformers — in the Language of Graphs (4)

Summary: Comparison of **Self-attention (SA)** and **Graph Attention Networks (GAT)**

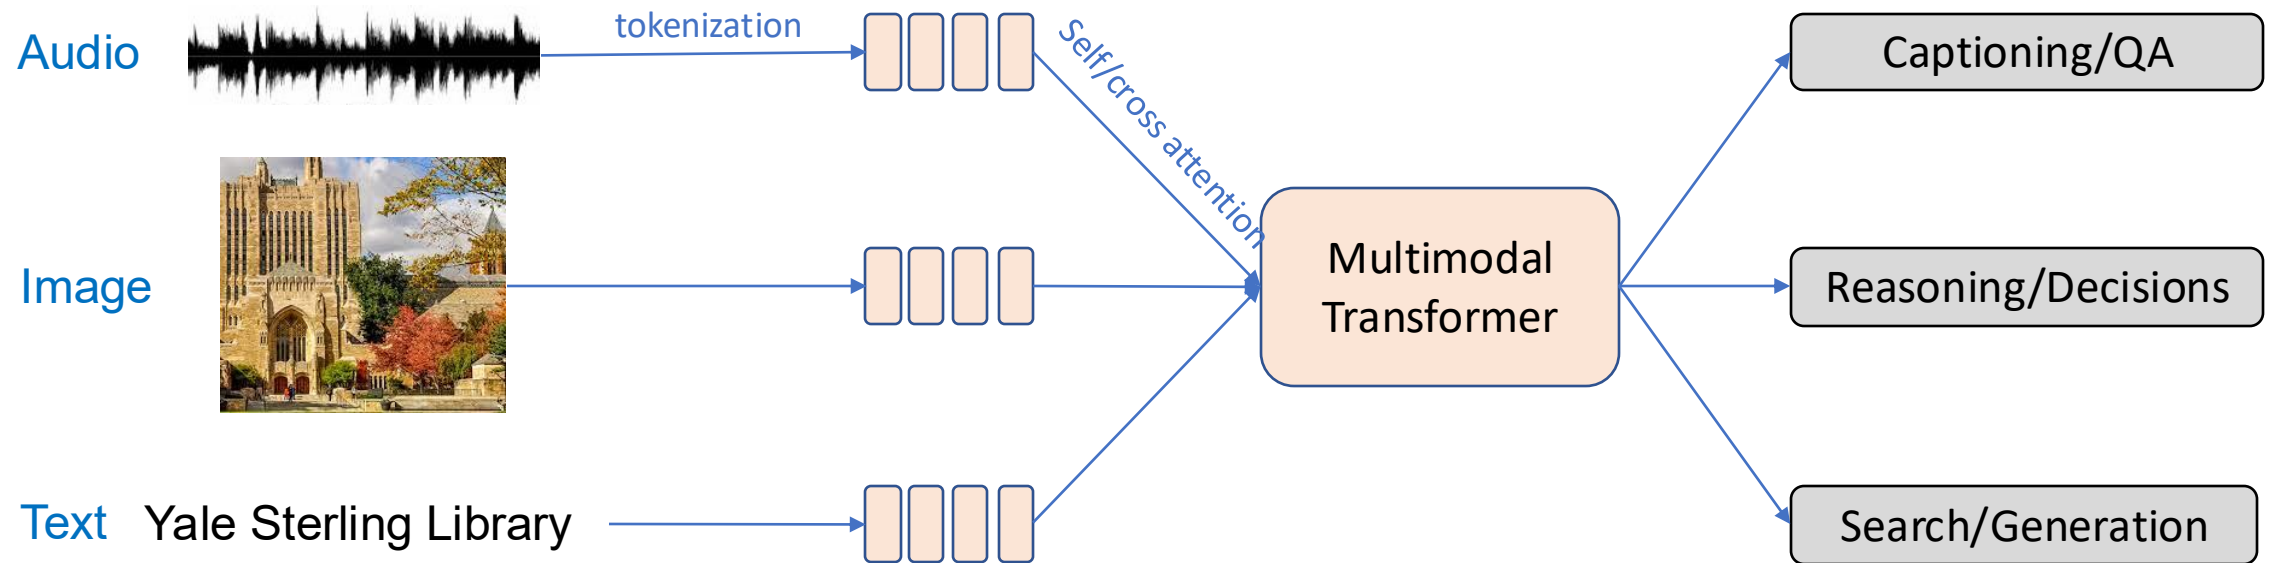
- Step ① Mapping
  - **SA**: different weights for  $q, k, v$ .  $q = w_q x, k = w_k x, v = w_v x$ .
  - **GAT**: shared weights for  $q, k, v$ .  $q = wx, k = wx, v = wx$ .
- Step ② Attention: **SA** uses dot-product attention, while (the original) **GAT** uses concatenation with MLP
  - Dot-product:  $e_{ij} = q_i^T k_j / \sqrt{d}$
  - Concat:  $e_{ij} = \text{act}(W [q_i || k_j])$ , where  $c$  is a weight vector and  $\text{act}$  is the activation function like LeakyReLU

# Graph Attention — in the Language of Transformer

- The above computations do not require the assumption of **the complete graph**.
  - We assume full connectivity, mostly because we do not want to miss any potential token correlations.
- Self-attention can be easily adapted to graph-structured input data where the token correlations are given by the **adjacency matrix**, by replacing the **complete graph** with the **input graph**.
  - $\text{Self-Att}(X) = \text{Softmax} \left( \frac{(\mathbf{W}_k X) (\mathbf{W}_q X)^T}{\sqrt{d}} \odot \mathbf{A}_G \odot \mathbf{W}_E \right) V$ .
  - $\mathbf{A}_G$  is the adjacency matrix of the graph and  $\mathbf{E}$  is the edge weights of the graph if any.
- The complexity is no longer  $O(n^2 d)$  but is linear to the edge number  $O(E)$

# Transformers for Multiple Modalities

- Transformers succeed across domains, enabling **general-purpose AI systems** where one general model can handle multiple input modalities (text, images, audio, etc.)



- E.g., CLIP, DALL-E, Stable Diffusion, GPT-4V, Gemini, Flamingo

# Why is Transformer a Popular Choice

- Resolves various challenges of RNN-based architectures
- Attention makes the architecture **expressive and flexible** for different application scenarios
- It is very amenable to **self-supervised objectives**
  - We can leverage the vast number of **unsupervised examples** to learn a general model
  - Can be fine-tuned for **many downstream tasks**
  - Can out-perform models that are only trained for a specific downstream task

