1. Make sure you are connected to YaleSecure or using Yale VPN.
2. Go to cpsc4710.ycrc.yale.edu and log in. You should see this page:

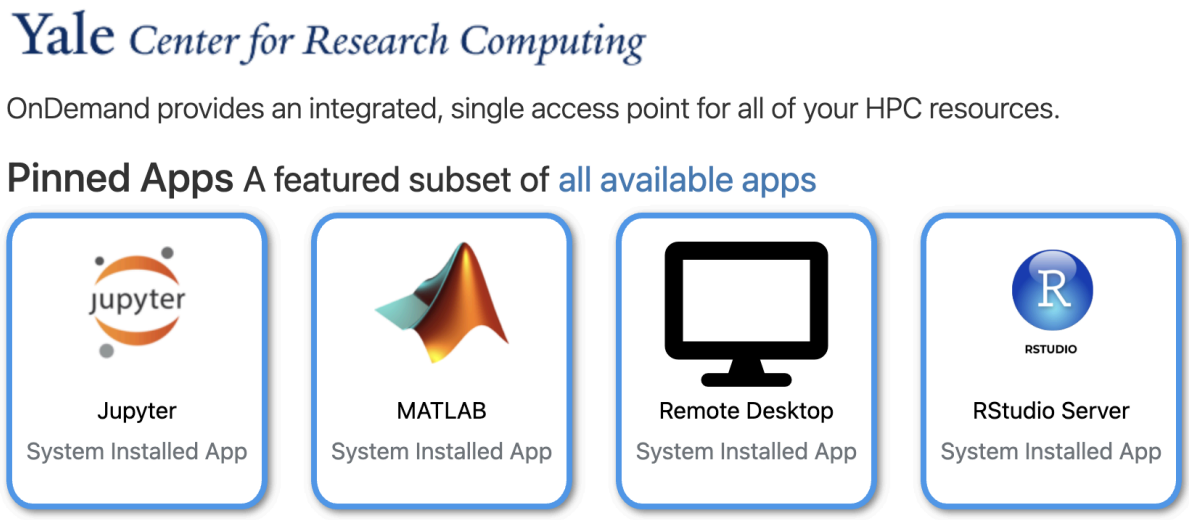

3. Select the “Jupyter” option.
4. In the page that comes up, **select “education_gpu” for the partition field**.
5. Select the desired session length and number of GPUs. Note that longer sessions or more GPUs might increase queueing time. You may leave the other fields blank.
6. Press the “Launch” button.
7. The session will start queueing, which might take some time. Once it is ready to start, click in and you’ll be presented with the Jupyter home page. Navigate to the folder with your NetID and you’ll be able to create files/run scripts as normal.
    a. To ensure everything is working, if you requested, say, 1 GPU, running `nvidia-smi` on terminal should show something along the lines of:

```
+-----------------------------------------------------------------------------------------+
| NVIDIA-SMI 570.133.20             Driver Version: 570.133.20     CUDA Version: 12.8     |
|-----------------------------------------+------------------------+----------------------+
| GPU  Name                 Persistence-M | Bus-Id          Disp.A | Volatile Uncorr. ECC |
| Fan  Temp   Perf          Pwr:Usage/Cap |           Memory-Usage | GPU-Util  Compute M. |
|                                         |                        |               MIG M. |
|=========================================+========================+======================|
|   0  NVIDIA RTX 5000 Ada Gene...    On  |   00000000:AB:00.0 Off |                  Off |
| 30%   32C    P8             14W /  250W |       2MiB /  32760MiB |      0%      Default |
|                                         |                        |                  N/A |
+-----------------------------------------+------------------------+----------------------+

+-----------------------------------------------------------------------------------------+
| Processes:                                                                              |
|  GPU   GI   CI              PID   Type   Process name                      GPU Memory |
|        ID   ID                                                             Usage      |
|=========================================================================================|
|  No running processes found                                                             |
+-----------------------------------------------------------------------------------------+
```

8. The CUDA version is 12.8. Conda is not installed by default and you may want to install it. For more information check out: https://docs.ycrc.yale.edu/applications/modules/.